

INVESTING IN AN AI-DOMINATED ECONOMY

VOLUME I

An Outfinity Venture Studio Guide and Investment Thesis on Research-Driven Startups, Durable Moats, Market Consolidation, and Venture-Scale Returns

*What winning companies will look like when models, software,
and features become commodities*

A research-based framework for venture capital, venture studios, business
angels, and family offices

DRAFT MANUSCRIPT

Version 0.5 | June 2026

OUTFINITY VENTURE STUDIO | INVESTOR EDITION | JULY 2026

INVESTOR EDITION

This monograph is a strategic investment thesis prepared for Outfinity Venture Studio and for investors evaluating artificial intelligence, enterprise software, and research-driven technology. It is not investment, legal, accounting, or regulatory advice. The evidence cutoff is July 2026.

The manuscript separates documented findings, strategic inference, and portfolio priors. A portfolio prior is a decision rule adopted under uncertainty because waiting for complete proof would make investment impossible. The book is deliberately firm where the mechanism is strong and explicit where the empirical record is still immature.

The governing belief is that AI will expand the number of useful products while compressing the number of independent product categories capable of sustaining venture-scale returns. Outfinity should therefore finance the accumulation of knowledge, rights, trust, infrastructure, and operational authority rather than the fastest production of software alone.

Contents

Research and Method Note	4
Executive Investment Thesis	6
Introduction. The Age of Cheap Capability	11
PART I. What Markets Select	15
1. Competition Does Not Select for Truth	16
2. Category Degradation and the Moloch Premium	21
PART II. Recent Market Histories	26
3. Three Histories, Three Mechanisms	27
PART III. The AI Consolidation Investment Thesis	36
4. Capability Abundance and the Death of the Generic Product	37
5. Why AI-Native Products Become Features	42
6. Why Traction No Longer Proves Venture Scale	47
PART IV. What Winning Startups Will Look Like	54
7. The Non-Generic Core: Moats and IP After Code	55
8. Research, Time, and Long-Duration Capital	60
9. How Winning AI-Era Startups Will Be Built	66
PART V. The Outfinity Investment Doctrine	74
10. Building a Portfolio for Consolidation	75
Conclusion. Build Institutions, Not Accelerated Features	84
References	86

Research and Method Note

This book tests a proposition that is stronger than the ordinary claim that most startups fail. It argues that artificial intelligence changes the population of startups and the geometry of their exits. AI lowers the cost of producing competent artifacts, increases entry, and accelerates imitation. At the same time, it strengthens the incentives of model providers, cloud platforms, enterprise incumbents, and customers to absorb narrow functionality into broader systems. The result can be simultaneous entrepreneurial abundance and strategic concentration. More companies may be formed, more useful software may be shipped, and fewer independent application categories may remain capable of producing the extreme outcomes on which venture portfolios depend.

The thesis is addressed to venture capital firms, venture studios, business angels, family offices, and strategic investors. Its unit of analysis is not the product demo and not even the current company. The relevant unit is the future economic position: the set of rights, assets, routines, relationships, permissions, and learning mechanisms that may remain scarce after models improve and visible features become reproducible. A product can be valuable to users while its producer captures little terminal value. A startup can reach product-market fit while remaining structurally designed for a modest acquisition. A technically difficult company can consume enormous capital without creating a defensible institution. The book is concerned with separating these cases before the market does so at the investor's expense.

The research base combines industrial organization, evolutionary economics, strategic management, organizational theory, software engineering, innovation finance, platform economics, institutional governance, and current empirical work on generative AI. Recent market histories are used as bounded evidence. Blockchain is examined only to study the market for narrative capital. Enterprise integration is examined only to study irreducible organizational complexity and vendor consolidation. Mobile and dating platforms are examined only to study finite interface capacity and the divergence between engagement measures and user outcomes. These cases are not universal analogies, and they are not invoked repeatedly as decoration. Each performs a specific explanatory task.

Several claims in the book are well established. AI has improved productivity on bounded tasks, especially for less experienced workers, although gains depend on task structure and the ability to evaluate output (Brynjolfsson, Li, and Raymond 2025; Dell'Acqua et al. 2026). The cost of inference at a given capability level has fallen rapidly, while leading model performance has become more clustered on several public evaluations (Stanford HAI 2025, 2026). Platform owners can envelop adjacent categories, and complementors alter innovation when platform entry becomes plausible (Eisenmann, Parker, and Van Alstyne 2011; Wen and Zhu 2019). Venture outcomes are extremely skewed, so a portfolio cannot be justified by the average useful company (Ouyang, Yu, and Jagannathan 2021; Azevedo et al. 2025). Financial frictions

can bias young firms away from radical innovation, while well-designed research grants can fund prototyping and improve later venture financing (Caggese 2019; Howell 2017).

The book's strongest forward-looking claim is not yet susceptible to definitive cohort analysis because most AI-native companies are young. The claim is therefore presented as a portfolio prior: absent evidence of a non-generic core, Outfinity should assume that a narrow AI application will face feature absorption, price compression, customer insourcing, or a strategic acquisition priced closer to build-versus-buy value than to category leadership. This prior does not declare every application company uninvestable. It changes the burden of proof. A founder seeking venture-scale capital must explain what the company will own after the capability that attracted the first customers becomes common.

The thesis would be weakened by sustained evidence that independent, generic application companies retain pricing power and low churn after model capability converges, platform bundles become available, and customers can reproduce the core workflow internally. It would also be weakened if broad platforms remain persistently narrow, if enterprise buyers prefer vendor proliferation despite security and integration costs, or if application usage automatically creates exclusive outcome data and switching costs without contractual or institutional control. These are observable conditions. Treating the thesis as falsifiable matters because conviction without a revision rule is branding, not investment reasoning.

Finally, the book distinguishes social concern from financial motive. Outfinity may care about market integrity, epistemic quality, and the social consequences of technology. Investors who do not share that concern should still care about category trust because it is economic infrastructure. Inflated claims increase customer acquisition cost, procurement time, regulatory exposure, technical diligence, and the discount applied to every supplier in the category. Responsible market construction is therefore not inserted as philanthropy. It appears where institutional quality affects returns.

Table 1. The status of the central claims

TERM OR ASPECT	DESCRIPTION
Documented premise	A claim supported by peer-reviewed research, official data, or a well-sourced case. It is used within the scope of the original evidence rather than as a universal law.
Strategic inference	A conclusion obtained by connecting documented premises through a stated causal mechanism, such as platform envelopment, vendor rationalization, or asset-stock accumulation.
Portfolio prior	A default underwriting assumption adopted before complete cohort evidence exists. It determines burden of proof, ownership targets, reserves, and valuation.
High-confidence forecast	A directional expectation supported by several mechanisms and early evidence, but still exposed to changes in regulation,

TERM OR ASPECT	DESCRIPTION
	platform strategy, open-source capability, or customer behavior.
Falsification condition	An observable development that would require the thesis to be narrowed, repriced, or abandoned rather than explained away after the fact.

Executive Investment Thesis

Outfinity's base case is that a generic AI application should be underwritten as a feature-in-waiting, not as an independent venture asset. It exploits a gap between a newly available capability and the speed with which platforms, incumbents, and customers can reorganize around it. That gap may be commercially valuable. It can produce revenue, strong user enthusiasm, and an attractive early acquisition. It should not be mistaken for a moat. In a market where the underlying capability is rented, continuously improved, and available to rivals, independence must be proved by the accumulation of something the supplier cannot sell to everyone else.

This is an asymmetric underwriting rule rather than a universal prediction. Outfinity may still finance an application wedge, but only when the wedge is a credible acquisition mechanism for a deeper asset. Validation, early growth, and local product-market fit are necessary evidence; they are no longer sufficient evidence of an independent category. The default question is not whether the product works. It is whether the company will still command a scarce position after the product's visible capability becomes standard.

The economic shift is straightforward. During the classical software era, a polished application was costly evidence. It implied a team capable of writing and maintaining code, designing an interface, integrating a stack, and operating infrastructure. Generative AI lowers the cost of all these signals. A small team can now produce a persuasive product, a tailored sales demonstration, a large content surface, and a plausible technical narrative. This is genuine leverage, but it also reduces the informational content of the artifact. When production becomes cheap, the investor must pay more attention to verification, consequence, and institutional position.

The thesis therefore begins with a migration of scarcity. Capability itself remains important, especially at the research frontier, but broad competence is increasingly available through competing models, open systems, and embedded enterprise products. Value migrates toward the conditions under which competence becomes trusted action: lawful access to data and feedback, authority inside workflows, systems of record, regulatory permission, physical deployment, assurance, distribution, and responsibility for outcomes. The company that controls these conditions can replace a model without replacing itself. The company that does not control them can be replaced when the model changes.

This leads to a consolidation forecast. Model providers have incentives to expand upward into memory, tools, identity, transactions, and application workflows. Enterprise incumbents have incentives to expand outward from the systems that already hold data, permissions, procurement approval, and customer trust. Customers have incentives to reduce supplier count because each AI vendor introduces a new security boundary, legal review, integration surface, monitoring burden, and source of operational risk. Fast imitators can reproduce visible workflows, while internal teams can assemble narrow tools at

declining cost. The generic startup is compressed from four directions before its original market has fully matured.

The likely consequence is not a desert without startups. It is a market with a wide base and a narrow summit. AI should increase company formation because it lowers the minimum efficient scale of experimentation. Emerging research already associates generative-AI exposure with greater startup formation and smaller entrants (Bena, Bian, and Giannetti 2025). Yet the same economics can reduce the number of durable categories because broad platforms can absorb functions and because the fixed investments required for research, assurance, distribution, and infrastructure rise with market size. Entry and concentration are not contradictory forecasts. They are successive stages of the same industrial cycle.

For a large venture fund, this matters because usefulness and venture scale are different propositions. A product may solve a real problem, maintain positive gross margin, and support an excellent small company. It may still have no credible route to an outcome large enough to return a fund. Venture returns remain concentrated in a small minority of investments; removing the best five percent of ventures can eliminate the aggregate risk-adjusted return in historical data (Ouyang, Yu, and Jagannathan 2021). A portfolio of companies likely to be acquired as features can therefore look successful operationally and remain inadequate financially.

Many AI-native exits will be strategically rational but economically modest. An incumbent may buy a startup to obtain a team, remove eighteen months of development delay, acquire a group of customers, or integrate a narrow workflow into a larger suite. The function may become more valuable after acquisition, while the independent company receives a price anchored to replacement cost and negotiating leverage rather than to the total downstream value created. Such an exit can be excellent for a founder, an angel, or a studio with low formation cost. It is a weak foundation for category-leader valuation at entry.

The opposite extreme is also dangerous. Capital intensity does not create defensibility by itself. Frontier models, semiconductors, robotics, energy, biology, and scientific infrastructure can require enormous investment, but money is only productive when it purchases non-compressible learning, assets, or permissions. Outspending rivals is a moat when each dollar advances a research frontier, deepens a manufacturing curve, expands a physical network, improves a proprietary evaluation system, or creates a rights position that cannot be reconstructed later. Capital that merely finances a larger version of a generic product accelerates commoditization rather than preventing it.

Outfinity should therefore search for a non-generic core. This core may be fundamental research, but the category is broader than laboratory science. It includes a difficult method that improves through experimentation, rights to connect decisions with outcomes, authority to execute consequential work, a regulated license, a physical deployment base, an accumulating exception taxonomy, a governed data network, or a brand that functions as a performance bond. The defining property is compounding scarcity: operation should increase the time, cost, or institutional difficulty faced by a competent follower.

Research has a privileged role because it can create knowledge before the commercial surface is obvious and because knowledge accumulation often requires time that cannot be compressed by a late entrant. Yet research alone is not a business. Teece's complementary-assets argument remains decisive: an innovator captures value only when it also possesses or controls manufacturing, distribution, service, regulatory access, or another route to appropriation (Teece 1986). Outfinity's thesis is not "fund scientists and wait." It is to form companies in which research, commercialization, and capital architecture are designed as one system.

The most promising enterprise companies will not sell isolated intelligence. They will control consequential processes. Their work will include data reconciliation, integration, exception handling, evaluation, audit, workflow redesign, and bounded responsibility for a completed outcome. These activities appear laborious because the remaining scarcity is often hidden in organizational complexity. The winning company turns this labor into an asset stock. Each deployment should improve a connector, policy model, decision taxonomy, evaluation set, trust relationship, or contractual permission that makes the next deployment faster and more reliable.

The strongest consumer companies will face a different constraint. People can encounter an unlimited number of applications but cannot maintain unlimited habits, subscriptions, memories, identities, and privacy relationships. A general assistant can absorb many utilities behind one habitual interface. A specialized consumer company must therefore own something besides function: identity, community, content, a transaction network, a longitudinal record, a trusted method, or an experience whose differentiation is itself part of consumption. The comparison is closer to mobile utilities than to games. Games can coexist because creative difference is the product; generic AI utilities converge around task completion.

This thesis implies a different capital model. The standard startup playbook of rapid MVP, broad public validation, aggressive feature expansion, and repeated financing against growth may expose a generic idea before it has accumulated protection. Public experimentation can leak information and attract entry, a problem formalized in recent work on competitive exposure (Gans 2026a). Research-driven and institution-building companies require staged capital tied to uncertainty reduction rather than to motion. A negative experiment can be valuable if it retires a critical path. A pilot is valuable only when it reveals representative implementation economics. A financing round is valuable only when it reaches a more defensible bargaining position.

Outfinity's belief is consequently selective and optimistic. The age of AI does not end venture investing. It raises the minimum intellectual standard for it. The easy company becomes easier to build and harder to own. The hard company becomes more expensive and, when correctly formed, more valuable because the surrounding software layer is being commoditized. The goal is not to predict every winning application. It is to build institutions whose advantage grows when intelligence becomes cheaper.



Figure 1. Value migrates from cheap capability toward scarce economic position

Table 2. The Outfinity thesis in one page

TERM OR ASPECT	DESCRIPTION
Base case	A generic AI application is a perishable arbitrage unless it converts early demand into a compounding, non-generic asset before platforms and customers close the gap.
Market structure	AI increases entry while platform expansion, bundling, vendor rationalization, and endogenous sunk costs increase consolidation.
Return implication	Many useful AI-native companies will produce small or medium strategic exits that are attractive to founders and early capital but insufficient for large-fund return geometry.
Winning core	Durable companies own a research frontier or an institutional scarcity such as outcome rights, workflow authority, assurance, regulated permission, physical assets, or compounding distribution.
Capital implication	Capital must be patient about time but impatient about learning. It should finance proof, asset accumulation, and non-compressible capability rather than narrative speed.
Studio implication	Outfinity should combine research formation, shared technical infrastructure, customer access, governance, and follow-on capital to create companies that ordinary accelerators cannot form.

Table 3. Evidence that would weaken the thesis

TERM OR ASPECT	DESCRIPTION
Persistent generic pricing power	Independent narrow applications retain strong pricing and low churn after comparable functions are bundled into model platforms and incumbent suites.
No platform expansion	Foundation-model and enterprise-platform providers remain structurally narrow despite clear demand for memory, tools, identity, transactions, and workflow automation.
Vendor proliferation at scale	Enterprise customers consistently prefer more narrow AI vendors even after security, integration, monitoring, and contractual costs are fully measured.
Automatic data moat	Ordinary application use reliably creates exclusive, reusable outcome data and switching costs without special rights, workflow control, or customer dependence.
Flat sunk-cost frontier	Research, assurance, brand, infrastructure, and distribution spending do not scale with market size and therefore do not favor organizations capable of repeated reinvestment.

CHAPTER INTRODUCTION

The Age of Cheap Capability

What exactly is an investor buying when software production, synthetic competence, and persuasive demonstration become inexpensive? The introduction defines the object of investment and states why the normal startup playbook is increasingly unable to answer that question.

The history of software is full of functions that survived after the companies built around them disappeared. Browsers absorbed utilities, operating systems absorbed applications, suites absorbed point products, and cloud platforms absorbed infrastructure categories. The function did not fail. It became a feature inside a broader economic relationship. Artificial intelligence increases the speed at which this can happen because the same model can cross product boundaries and because natural language reduces the engineering required to reproduce a visible workflow.

The central distinction of this book is therefore between product value and independent enterprise value. Product value asks whether a user receives a benefit. Independent enterprise value asks whether the company can retain bargaining power after suppliers, rivals, customers, and adjacent platforms respond. Venture value asks a still narrower question: whether the investor can acquire enough ownership in a company whose terminal outcome is large enough to matter to the portfolio. These questions are correlated, but they are not interchangeable.

Generative AI greatly increases the number of products exposed to this fate. The technology is unusually general, its capabilities can be accessed through APIs, and its outputs can be wrapped in interfaces with modest initial engineering. A team can discover a narrow use case, present it convincingly, and reach paying customers before the market has agreed on the category's eventual boundaries. That speed is economically useful. It also creates a period in which temporary scarcity can be mistaken for structural defensibility.

The distinction matters because venture capital is not paid merely for identifying useful things. It is paid for acquiring ownership in organizations whose value can compound faster than dilution and failure consume the portfolio. A product can save customers time and still be unable to retain margin. A company can grow locally and still be impossible to scale without a hidden service organization. A startup can lead a category and still discover that the category was only a feature boundary drawn during an early phase of the technology.

The conventional vocabulary of product-market fit hides these differences. Product-market fit is an observable relationship between a product and a group of users under current conditions. Defensibility concerns the persistence of bargaining power after rivals react, platforms expand, customers learn, regulations change, and input costs decline. Scale concerns the repeatability of delivery when exceptional founders, subsidized inference, unusually clean data, and enthusiastic design partners are removed. Venture return concerns

whether the resulting ownership value is large enough, and arrives soon enough, to matter to a particular fund.

An investor therefore needs several nested units of analysis. The artifact is the visible output: a model, interface, report, workflow, device, or prototype. The product packages the artifact into a repeatable offer. The company coordinates people, capital, contracts, and operations around that product. The system connects the company to data, suppliers, customers, standards, regulators, and complementary assets. The institution persists because it possesses recognized rights, responsibilities, reputation, and governance. AI lowers the cost of the first layers much faster than the last.

The central danger is the artifact illusion. When production is expensive, the existence of a polished product is evidence that the team possesses scarce capability. When production becomes cheap, polish ceases to reveal much about the organization behind it. The same applies to growth. AI can improve outbound sales, content production, onboarding, support, and experimentation. A company may therefore reach impressive early metrics with an operating core that remains shallow. The speed of evidence creation increases faster than the reliability of the evidence.

This is one reason the AI boom can produce both genuine productivity and widespread capital misallocation. Noy and Zhang found that generative AI improved speed and quality on professional writing tasks, while Brynjolfsson, Li, and Raymond documented meaningful productivity gains in customer support, especially among less experienced workers (Noy and Zhang 2023; Brynjolfsson, Li, and Raymond 2025). Dell'Acqua and colleagues described a jagged technological frontier: workers performed better on tasks inside the model's competence but could be misled on tasks outside it (Dell'Acqua et al. 2026). These findings do not imply that AI is superficial. They imply that capability is uneven and that organizational value depends on matching the technology to tasks, controls, and expertise.

The market is simultaneously moving toward lower costs and higher concentration. The Stanford AI Index reported a dramatic decline in the cost of querying models at a given benchmark level between 2022 and 2024, while its 2026 edition described convergence among leading model providers and rapid organizational adoption (Stanford HAI 2025, 2026). Lower cost expands demand and experimentation. Convergence weakens differentiation at the input layer. The combination is favorable to users and to companies that can incorporate interchangeable intelligence into a stronger system. It is unfavorable to companies whose identity depends on privileged access to a capability that is becoming standard.

The word wrapper has become a dismissive shorthand for such companies, but the term is analytically crude. Every software product wraps lower-level components. The relevant distinction is between a replaceable wrapper and an accumulating application institution. A replaceable wrapper translates model output into a convenient interface without acquiring control over data, workflow, distribution, responsibility, or learning. An accumulating institution uses the initial interface to enter a process, observe outcomes, earn permissions,

resolve exceptions, and build assets that remain valuable when the underlying model changes.

Some wrappers can earn substantial revenue during a fast-moving market. Early timing, app-store optimization, paid acquisition, localization, and aggressive subscription design can produce impressive cash flow. The investor should not deny these outcomes. The investor should classify them correctly. A temporary distribution arbitrage may support a bootstrapped business or an attractive small acquisition. It becomes a venture thesis only when the company can explain how temporary attention is converted into a position that a platform cannot cheaply reproduce.

The same distinction applies to deep technology. Scientific depth does not automatically create a moat. A technically brilliant company can fail because it lacks complementary manufacturing, regulatory, distribution, or commercialization assets, a problem emphasized by Teece's work on profiting from innovation (Teece 1986). Yet fundamental research can create a more durable core because it generates lead time, tacit knowledge, patents, specialized infrastructure, and learning curves that are not immediately available through a common API. The strongest companies combine a hard technical core with an institution capable of capturing its value.

The investment thesis is therefore skeptical of two symmetrical errors. The first is application romanticism: the belief that proximity to the customer is sufficient even when the product can be enveloped or bundled. The second is frontier romanticism: the belief that technical difficulty and capital intensity are themselves moats. A research program that cannot commercialize is a laboratory, not necessarily a company. A product that cannot survive capability convergence is a feature, not necessarily a venture asset.

The correct object of investment for Outfinity is a compounding system. Such a system may be centralized or governed through a consortium. It may sell software, services, infrastructure, or outcomes. It may rely on proprietary models, open models, or a portfolio of suppliers. Its defining property is that each period of operation improves its future position through assets that the firm can lawfully retain and use. The improvement may take the form of better data, stronger distribution, lower unit costs, regulatory standing, customer trust, process authority, or scientific knowledge.

The normal startup playbook was designed for a world in which building was expensive and learning from the market was comparatively scarce. A minimum viable product converted technical effort into customer evidence. Public launches created distribution and demonstrated execution. Feature velocity signaled organizational competence. In the AI era, these signals have changed meaning. Building is cheaper, public experimentation is more legible to rivals, and feature velocity may measure access to common models rather than proprietary learning.

This does not make validation obsolete. It makes validation narrower. A successful pilot proves that a customer experienced value under a particular configuration of data, people, attention, and exceptions. It does not prove that the configuration can be reproduced across normal organizations, that the gross margin survives hidden service labor, that the workflow remains

independent after an incumbent bundles the visible function, or that the company can preserve its advantage when the model supplier improves. Validation should reduce uncertainty; it should not be used to declare unrelated uncertainties solved.

The investment object is a compounding system. The system can be centralized, distributed through a governed consortium, or assembled across a network of partners. It can use proprietary models, open models, or several interchangeable suppliers. What matters is that each period of operation creates an asset the firm can lawfully retain and reuse. The asset may be experimental knowledge, outcome-linked data, process authority, physical infrastructure, reputation, or a channel that becomes cheaper with use. Without this accumulation, growth increases visibility faster than defensibility.

The polemical form of the thesis is simple: the feature graveyard will become much larger. The scientifically defensible form is conditional. As model capability converges and application production becomes cheaper, generic product boundaries will have shorter half-lives. Firms can escape by controlling a non-generic core, and the number of escapes will depend on platform conduct, interoperability, regulation, customer preferences, and the ability of investors to finance long learning cycles. Outfinity should underwrite the conditional mechanism, not the slogan.



Figure 2. Markets select the traits rewarded by measures, contracts, and time horizons

PART I

WHAT MARKETS SELECT

Before investors can identify winning AI companies, they need a precise account of what rivalry rewards. Competition is a search process, but its objective function is supplied by institutions, measures, contracts, and time horizons. Poorly specified selection can produce intense activity without proportional progress.

CHAPTER ONE

Competition Does Not Select for Truth

This chapter replaces the moral language of competition with an institutional and evolutionary model. The question is not whether rivalry is good or bad in the abstract, but which traits reproduce under the actual financing and purchasing environment.

The Market as a Selection Environment

Competition is commonly described as a cleansing force. Inferior products disappear, stronger firms expand, and capital is reallocated toward better uses. The description is often correct, but it omits the most important question: better according to which observable criterion, over which time horizon, and with whose costs included? Natural selection does not optimize for beauty, intelligence, or social welfare in the abstract. It preserves traits that increase reproductive success in a given environment. Economic selection works in the same way. Firms survive by satisfying the effective fitness function created by customers, investors, platforms, regulators, employees, and timing.

Evolutionary economics has long treated firms as populations of routines that search, imitate, and adapt under bounded rationality rather than as perfect optimizers of a social objective (Nelson and Winter 1982). This perspective is more useful for technology investing than the simplified belief that competition necessarily produces progress. A market can reward engineering depth, but it can also reward access to distribution, persuasive narratives, regulatory arbitrage, financial leverage, rapid imitation, or the ability to externalize maintenance. The selected trait is whichever combination produces survival and expansion under actual constraints.

Baumol's distinction among productive, unproductive, and destructive entrepreneurship makes the institutional point directly. Entrepreneurial talent is not inherently allocated to socially valuable inventions. The reward structure channels it toward activities that may create value, redistribute value, or destroy value while remaining privately profitable (Baumol 1990). A technology boom with weak measurement can therefore attract capable people into the optimization of financing signals rather than technical substance. The individuals may be rational and the market may be competitive, yet the category may advance slowly.

The relationship between competition and innovation is consequently contingent. Aghion and colleagues found evidence consistent with an inverted-U relationship in which innovation may rise with competition at lower levels and fall when rivalry becomes sufficiently intense (Aghion et al. 2005). The precise shape is not universal, but the underlying mechanism is important. Firms innovate when the prospective gain from escaping competition is large

enough and when they can appropriate part of that gain. If margins are compressed before difficult research pays off, imitation and tactical differentiation can dominate exploration.

This distinction can be expressed as a selection problem. A market environment contains the rules of financing, procurement, platform access, liability, disclosure, and measurement. Firms respond with a portfolio of behaviors that includes research, engineering, sales, bundling, imitation, lobbying, and timing. The population outcome depends on which behaviors are rewarded before their long-term consequences become visible. A market that pays for verified performance selects differently from a market that pays for a credible announcement.

The same technology can therefore develop differently across markets. In medicine, a claim may require trials, regulatory review, post-market monitoring, and professional liability. In consumer software, a claim may be tested by immediate adoption and rapid churn, but long-term harms can remain external. In enterprise software, a product may appear effective during a pilot while implementation cost and exception handling emerge only after broad deployment. The competitive environment defines which portion of reality enters the firm's score.

Investors participate in this environment. Financing milestones, reserve policies, board metrics, and exit expectations are not neutral observations of market quality. They change the traits that founders have reason to develop. When capital is released against user growth without attention to gross-margin normalization, firms optimize acquisition. When financing depends on benchmark performance that is easy to game, firms optimize the benchmark. When valuation rewards category language more than durable evidence, category language proliferates.

Goodhart's law is often summarized as the proposition that a measure ceases to be useful when it becomes a target. Campbell's law extends the point to social decision-making: the more a quantitative indicator is used for consequential decisions, the more pressure there is to corrupt the process it was intended to monitor (Goodhart 1975; Campbell 1976). In organizations, Kerr described the recurring folly of rewarding one behavior while hoping for another (Kerr 1975). Technology markets are especially vulnerable because many desired outcomes are delayed, multidimensional, and difficult to observe, while proxies such as downloads, pilots, model scores, funding rounds, and media visibility are immediate.

The problem is not that metrics are useless. Markets require compression. Investors cannot directly observe every line of code, customer interaction, safety incident, and future adaptation. Metrics are interfaces between complex reality and capital allocation. The danger appears when the proxy is cheaper to improve than the underlying outcome. AI widens this gap because it can produce plausible outputs, synthetic evaluations, persuasive case studies, and tailored narratives at low cost. A market that does not upgrade verification as production becomes cheaper will select for proxy optimization.

The concept of a Red Queen race adds another layer. In evolutionary biology, an organism may need to keep adapting merely to preserve its relative

position. In business, firms introduce features, integrations, announcements, and discounts because rivals do the same, even when the collective escalation produces little customer value. Empirical research has documented Red Queen dynamics in competitive action and imitation (Derfus et al. 2008; Giachetti, Lampel, and Li Pira 2017). The key feature is relative rather than absolute payoff. A company may know that a rushed feature increases maintenance cost, yet omitting it can weaken a sales comparison this quarter.

Such races are not irrational. The individual move can be locally optimal while the population outcome is poor. Faster releases can preserve visibility. Broader feature lists can satisfy procurement checkboxes. Aggressive marketing can reduce the chance that a more careful rival defines the category. The long-term cost is dispersed across technical debt, customer distrust, talent burnout, and a market in which claims become less informative. No participant needs to intend the degradation.

AI increases the importance of this selection model because it lowers the price of surface conformity. A founder can match the visible features of a competitor, generate a persuasive benchmark narrative, and tailor a demonstration to a prospect without possessing the same operational depth. The strategy may be rational when capital and customer attention are allocated before long-run quality is observable. The market is not choosing between truth and falsehood in a philosophical sense. It is choosing among signals with different production costs and verification delays.

This produces a systematic bias toward what can be demonstrated quickly. Fundamental research, safety engineering, process redesign, and institutional permission often have delayed and nonlinear payoffs. Interface polish, pilot count, generated content, and announcement cadence are immediate. When financing milestones reward the immediate variables, investors unintentionally shorten the horizon available to the activities that could create a durable moat. The market later interprets the absence of moats as evidence that the category was weak, although its own selection regime helped prevent them from forming.

When Rivalry Produces Discovery and When It Produces Noise

A useful investment theory must explain not only failure but also the conditions under which competition works well. Rivalry tends to improve a technology market when buyers can distinguish quality, when failure has consequences for the seller, when innovators can capture enough of the value they create, and when entry and switching remain credible. These conditions need not be perfect. They must be strong enough that investing in the underlying outcome is more attractive than investing in its appearance.

Quality observability is the first requirement. Akerlof's market for lemons shows how asymmetric information can reduce the price offered for quality and drive good suppliers from a market (Akerlof 1970). In technology, quality is often revealed only after integration, scale, adversarial use, or regulatory exposure. A weak product can imitate the visible features of a strong one long before

customers observe the difference. If buyers cannot price depth, founders receive a lower return on depth and a higher return on signals.

Reputation can partly solve this problem. Klein and Leffler argued that future rents can function as a performance bond: a firm with a valuable reputation has something to lose by cheating on quality (Klein and Leffler 1981). Brand is therefore not merely emotional differentiation. In markets with high verification cost, it can become an institutional commitment. The mechanism works only when future rents are substantial and when the brand remains accountable for failure. A disposable startup name backed by aggressive customer acquisition does not provide the same bond.

Appropriability is equally important. Teece showed that innovators often fail to capture value when complementary assets such as manufacturing, distribution, service, or regulatory access are controlled by others (Teece 1986). A technically superior startup may create a market that an incumbent monetizes. Patents, trade secrets, lead time, data rights, switching costs, and specialized assets can support appropriation, but each has limits. If the market cannot reward long-cycle work, firms will rationally move toward features with faster proof and faster exit.

The allocation of consequences also shapes selection. A vendor that bears warranty costs, service obligations, liability, or reputational loss has stronger reason to verify. A vendor that can sell a pilot, recognize revenue, and exit before failure is observed has weaker reasons. This does not imply that founders act dishonestly. It means the contract determines which uncertainties remain with the seller. Markets improve when responsibility follows control and when claims are tied to evidence that survives normal use.

Contestability matters because incumbents can otherwise preserve weak outcomes through distribution, standards, or lock-in. Network effects and switching costs may generate efficiency, but they can also make historical position more important than current quality (Arthur 1989; Katz and Shapiro 1985). A market can therefore exhibit intense rivalry among complementors while remaining structurally uncompetitive at the platform layer. Investors who measure only the number of startups may mistake entry abundance for contestability.

Competition becomes destructive when these conditions are reversed. Quality is hard to observe, imitation is cheap, consequences arrive late, platforms control distribution, and investors reward speed through proxies. The resulting market can contain many firms and little meaningful discovery. It may look dynamic because products and narratives change rapidly. Its underlying technical frontier may move slowly.

The practical conclusion is demanding rather than anti-competitive. Competition produces discovery when customers can observe quality, innovators can appropriate a meaningful share of the value, failure remains attached to the seller, and entry remains contestable. It produces noise when the proxy is easier to optimize than the outcome, when reputation can be discarded through rapid entry and exit, and when platforms can appropriate the complement after others have paid to discover it.

Entry alone is not evidence of healthy competition. A market can combine low formation costs with high terminal concentration when scale, data, trust, bundling, or research intensity determine survival. The relevant test is whether entrants can preserve bargaining power after the discovery phase, not whether product demonstrations proliferate during it.

For Outfinity, market analysis must therefore precede product enthusiasm. The diligence question is not merely whether competitors exist or how fast they are moving. It is what the environment requires a company to become in order to survive. If survival requires escalating marketing, copying adjacent features, and accepting unbounded customization, the market is selecting against the very asset accumulation the investment thesis needs. A strong founder can still win tactically, but the portfolio should not confuse tactical fitness with durable value.

Table 4. Conditions under which competition improves or degrades a technology market

TERM OR ASPECT	DESCRIPTION
Quality observability	Rivalry improves outcomes when buyers or credible intermediaries can distinguish durable performance from persuasive imitation before weak suppliers capture most of the value.
Appropriability	Long-cycle innovation is financeable when the innovator can retain value through rights, lead time, complementary assets, accumulated learning, or advantaged distribution.
Consequence allocation	Claims become informative when the seller bears warranty, liability, service, reputation, or outcome-linked costs when the system fails.
Metric integrity	Measures remain useful when they are difficult to game, close to the customer outcome, and revised as strategic behavior adapts.
Contestability	Competition disciplines incumbents only when switching, multihoming, interoperability, and independent distribution remain credible.
Time horizon	Research survives when capital and procurement permit evidence to mature before short-run comparison destroys the expected return.

CHAPTER TWO

Category Degradation and the Moloch Premium

This chapter explains how individually rational behavior can destroy the credibility of an entire market category. It treats trust and semantic clarity as shared economic assets and shows why the prisoner's dilemma is too small to model the resulting system.

Category Trust as Financial Infrastructure

A technology category exists partly as a semantic agreement. Customers, investors, employees, regulators, and partners must believe that the label identifies a coherent problem and a plausible class of solutions. The belief need not be enthusiastic. It must be strong enough that participants are willing to incur the cost of evaluation. When that willingness collapses, even capable firms face a legitimacy discount.

Zuckerman's work on categorical legitimacy showed that organizations can be penalized when evaluators cannot place them within recognized categories (Zuckerman 1999). The inverse problem is equally important. A category can become so broad, inflated, or contaminated that membership conveys little information. The label remains visible but ceases to reduce uncertainty. Buyers respond with heavier diligence, longer procurement, lower willingness to pay, or avoidance.

Category trust can therefore be treated as a common-pool asset. Every supplier benefits when customers believe that demonstrations are meaningful, security claims are credible, and pilots can become production systems. Each supplier also has an incentive to consume that trust by exaggerating capability, hiding service work, using favorable benchmarks, or launching before the system is ready. The private gain arrives quickly; the cost is distributed across competitors and future transactions.

Hardin's tragedy of the commons captures the externality but not the solution (Hardin 1968). Ostrom's research showed that communities can govern shared resources through boundaries, monitoring, graduated sanctions, conflict-resolution mechanisms, and rules adapted to local conditions (Ostrom 1990, 2010). Technology categories rarely recognize trust as a resource requiring governance. They rely on general law, reputation, and customer diligence. These mechanisms may be too slow during a boom in which firms can enter, raise capital, and disappear before consequences mature.

The category-trust mechanism can be written conceptually as an accumulation equation. Trust rises through verified outcomes, transparent standards, credible guarantees, and responsible failure. It falls through disappointment, semantic dilution, security incidents, benchmark gaming, and implementation debt. A category can grow in revenue while the stock of trust

declines because current demand reflects novelty or fear of missing out. The deterioration becomes visible later through procurement skepticism and consolidation.

Category degradation often begins with a useful ambiguity. A new term is broad enough to coordinate investment before the technology is fully understood. The ambiguity attracts researchers, founders, customers, and capital. As attention increases, the label acquires financial value. Firms with weaker technical fit adopt it because the market rewards association. The semantic boundary expands, comparison becomes harder, and marketing intensity rises. Strong firms must either spend more on explanation or imitate the category language to remain visible.

The process resembles adverse selection. When buyers cannot distinguish deep capability from superficial compliance, they discount the entire category. High-quality suppliers face the same skepticism created by low-quality entrants. Some reduce investment, reposition outside the label, or sell to incumbents that can provide reputation. The remaining market becomes more dependent on brand and distribution, which accelerates concentration.

This is why category trust is not merely a public-relations concern. It affects the cost of capital, sales cycles, regulation, talent, and the strategic value of independence. A firm that consumes trust can impose a real economic cost on rivals. An investor who finances such behavior may improve one portfolio company's short-term position while weakening the exit environment for the rest.

The incentives become stronger when firms can exit before consequence. Venture-backed markets separate the time of financing from the time of operational proof. Founders, employees, and early investors may realize partial liquidity through later rounds or acquisitions while customers bear the long tail of implementation. Again, this does not imply intentional deception. It means that the market's settlement period is longer than the participant's holding period.

The same pattern appears when frameworks and standards proliferate. A new abstraction promises interoperability or speed. Vendors extend it to cover more cases, buyers request compatibility, and consultants build practices around it. Exceptions accumulate. The standard becomes a surface of agreement over heterogeneous underlying systems. At some point, the category sells management of the complexity it helped create. Trust may persist because switching is costly, but the original promise loses credibility.

Investor language can accelerate this process. Terms such as platform, ecosystem, agent, infrastructure, proprietary data, and network effect are economically meaningful under specific conditions. During a boom they become status claims. The words are not false in isolation; they are under-specified. A feature can be called a platform because others can connect to it. A collection of customer documents can be called proprietary data even when the company lacks reuse rights. A workflow can be called agentic even when it is a brittle sequence of prompts. Semantic inflation reduces the information content of the term.

A category remains healthy when claims become more precise as capital increases. A degrading category becomes less precise because ambiguity expands the addressable narrative. The investor's duty is to reverse the incentive. Capital should be cheaper for firms that define boundaries, disclose failure modes, and accept testable obligations. Otherwise the market rewards those that privatize attention and socialize doubt.

Category trust matters to an investor even when social welfare is excluded from the objective. A trusted category lowers search cost, improves the quality of inbound customers, shortens procurement, reduces the discount applied to ambitious claims, and expands the set of employees and partners willing to specialize. A degraded category reverses these effects. The strongest companies pay a tax for the weakest because customers cannot cheaply distinguish them before purchase.

The tax can be especially severe in deep technology. Technical quality is difficult to observe, proof cycles are long, and failures can be attributed to implementation, customer readiness, or an unfavorable environment. This creates room for low-quality firms to imitate the language of high-quality firms. Akerlof's lemons mechanism predicts that when buyers cannot distinguish quality, the price paid for depth falls and capable sellers either leave or spend more on costly signals (Akerlof 1970). Brand, certification, warranties, independent evaluation, and long-term contractual responsibility are not cosmetic responses. They are market institutions that make quality legible.

Why the Prisoner's Dilemma Is Too Small

The prisoner's dilemma is often used to explain races in which firms would collectively prefer restraint but individually benefit from defecting. It captures one element of category degradation: each company may prefer a market with honest claims yet fear losing if competitors exaggerate. Repetition, reputation, and credible punishment can support cooperation. The model is useful, but it is too small for the phenomenon described here.

The first limitation is population structure. A technology market contains startups, incumbents, platforms, investors, customers, regulators, consultants, media, and standards bodies. They do not share the same payoff matrix. A platform may benefit from complementor experimentation while retaining the option to enter successful categories. A customer may benefit from low prices in the short term while becoming dependent on a vendor. An investor may prefer a fast acquisition even if the category's independent future weakens. The game is multi-agent and asymmetric.

The second limitation is endogenous rules. In the prisoner's dilemma, the action set and payoffs are given. Technology firms change the game through bundling, acquisitions, API policy, contractual restrictions, standards, lobbying, and control of data. A company with sufficient scale can redefine what counts as a product category. Competition is therefore partly constitutional: actors compete over the rules that determine future competition.

The third limitation is entry and selection. Players in the textbook game remain the same. Real markets replace them. Strategies that grow quickly

attract imitation and capital; strategies that require patience may disappear before their value is visible. The market does not merely choose actions. It changes the population of organizations capable of acting. Evolutionary game theory is therefore more appropriate than a one-shot dilemma.

The fourth limitation is uncertain consequence. Firms often do not know whether an aggressive tactic will damage the category. The effects are delayed, nonlinear, and confounded by technological change. A startup may reasonably believe that a rushed launch creates useful learning. A customer may accept a weak pilot because the option value is high. Category degradation can emerge without common knowledge of the payoff structure.

The fifth limitation is positional competition. Some rewards depend on rank rather than absolute value. Media attention, venture financing, app-store placement, talent, and perceived leadership are scarce. Frank and Cook described winner-take-all markets in which small performance differences can produce large reward differences (Frank and Cook 1995). Under positional pressure, firms may spend resources to preserve rank even when the total market gains little. Restraint is difficult because the benefit of winning depends on others losing.

The sixth limitation is increasing returns. Network effects, learning, data accumulation, and distribution can amplify early advantages (Arthur 1989; Katz and Shapiro 1985). A tactic that wins a temporary lead may create permanent control. This changes incentives dramatically. Firms may rationally overinvest in speed because the market's eventual structure is path dependent. The social cost of redundant entry can coexist with the strategic necessity of reaching scale before a rival.

The seventh limitation is principal-agent conflict. Founders, employees, investors, customers, and acquirers have different horizons. A founder may value a sale that does not return a late-stage fund. A salesperson may close a contract that creates implementation loss for the company. A corporate buyer may sponsor a pilot for internal political reasons. The organization itself does not possess a single payoff function.

Taken together, these models produce what this book calls the Moloch premium: the temporary private return earned by adopting a tactic that degrades the shared environment. Aggressive claims may lower acquisition cost. Rapid imitation may preserve relevance. Underinvestment in maintenance may improve reported margins. A proprietary standard may increase lock-in. Each tactic produces a premium before the external cost is fully priced.

The premium can persist because the market has no immediate mechanism for charging the firm. Later costs appear as regulation, integration labor, customer skepticism, security requirements, higher insurance, longer sales cycles, and consolidation. By the time the bill arrives, the firms that earned the premium may have been acquired or recapitalized. The delayed settlement is central to the problem.

The broader system can be described as a multi-agent evolutionary game with endogenous rules. Firms enter and exit, strategies reproduce, platforms alter interfaces, investors change the payoff to speed, and customers update their willingness to evaluate the category. Increasing returns mean that early tactical

advantages can become structural. Principal-agent conflicts mean that the person receiving the immediate reward may not bear the later cost. Positional competition means that the value of a funding round or launch can depend on rank rather than on absolute customer improvement.

The term Moloch premium names the temporary private return earned by a tactic that degrades the environment on which all participants depend. The premium can arise from overstating a benchmark, disguising service labor as software margin, launching a weak standard before rivals, or promising a broad platform before the narrow product is reliable. The tactic is not necessarily fraudulent. Its defining property is temporal asymmetry: the private benefit is realized before the distributed cost is charged.

Market repair is possible, but it requires institutional design. Ostrom's work on common-pool resources emphasizes boundaries, monitoring, graduated sanctions, conflict resolution, and rules fitted to local conditions (Ostrom 1990, 2010). In technology markets, analogous mechanisms include precise category definitions, benchmark governance, incident reporting, certification, contractual accountability, and investor norms that reward evidence over announcement. These mechanisms are not substitutes for competition. They determine what competition selects.

Table 5. Why no single game-theoretic model is sufficient

TERM OR ASPECT	DESCRIPTION
Common-pool resource	Explains how trust, semantic clarity, security, data quality, and customer attention can be consumed privately while the cost is distributed across the category.
Lemons market	Explains why weak observability lowers the return to technical depth and encourages strong suppliers to imitate signals that weak suppliers can cheaply produce.
Red Queen race	Explains escalation in features, launches, benchmarks, and financing speed that preserves relative position without proportional absolute improvement.
Platform envelopment	Explains how an actor with an adjacent user relationship can bundle a complement and redefine the category in its own favor.
Increasing returns	Explains how temporary differences become durable through learning, data, distribution, switching costs, and network effects.
Evolutionary selection	Explains entry, exit, imitation, and population change when the identity of players and even the rules of competition are not fixed.

PART II

RECENT MARKET HISTORIES

The purpose of history here is not analogy by repetition. Each case isolates one mechanism that will later be generalized without returning to the same names. Narrative capital, accumulated complexity, and finite interface capacity are different problems and should remain analytically separate.

CHAPTER THREE

Three Histories, Three Mechanisms

This chapter uses three recent markets to identify distinct failure modes. Blockchain reveals how category association can be financed before architectural necessity is proved. Enterprise integration reveals how organizational complexity survives new abstractions and rewards vendors that can govern it. Consumer platforms reveal that digital supply is unlimited while human habits and trusted interfaces are not.

Narrative Capital: When an Architecture Becomes a Credential

Blockchain began as a specific family of solutions to difficult coordination problems. Distributed consensus, cryptographic commitments, programmable assets, fault tolerance, and censorship resistance were not marketing inventions. They represented genuine research questions and, in some settings, useful design choices. The category degraded when the term ceased to identify a bounded architecture and began to function as a credential for access to capital, attention, and strategic relevance.

The economic attraction was understandable. A blockchain could be described as a new database, a new ownership layer, a new financial rail, a new governance system, or a new internet. Each description was plausible in a subset of cases. Their combination created an unusually elastic narrative. Founders could attach the term to payments, supply chains, identity, media, healthcare, energy, logistics, social networks, and enterprise data. Investors could treat the category as an option on a systemic transformation without agreeing on the mechanism.

Initial coin offerings intensified the elasticity because financing and product architecture became entangled. Tokens could fund development, coordinate users, reward early adopters, and create liquidity before the underlying network was mature. Howell, Niessner, and Yermack documented that ICOs represented a meaningful financing innovation and that certain disclosure and liquidity characteristics were associated with subsequent outcomes (Howell et al. 2020). The mechanism was not inherently fraudulent. It was unusually permissive of cheap signals because a tradable claim could precede a working system.

In conventional venture financing, an investor receives a claim on a company that owns assets, employs people, and is governed by corporate law. In a token sale, the relationship among the token, the network, the development organization, user rights, and future cash flows could be ambiguous. The ambiguity increased flexibility but weakened the correspondence between technical success and investor return. A project could be appreciated because

the narrative attracted liquidity even when the architecture remained unnecessary or incomplete.

This produced a market for narrative capital. Narrative capital is the capacity of a story to mobilize money, labor, and attention before operational proof. Every emerging technology requires some narrative capital because future systems cannot be financed solely by present evidence. The danger appears when the story becomes transferable across projects at lower cost than the underlying capability. During the blockchain boom, a white paper, token design, advisory board, and prototype could create the appearance of participation in a large technological discontinuity.

Swartz describes the 2017 ICO bubble as a network scam in which social and technical networks helped produce collective credibility without requiring every participant to be individually fraudulent (Swartz 2022). The formulation is useful because it shifts attention from bad actors to system structure. Promoters, exchanges, media, investors, and communities reinforced one another. Rising prices validated the narrative, the narrative attracted projects, and the proliferation of projects expanded apparent adoption.

The result was semantic inflation. Terms such as decentralization, trustlessness, community ownership, and token utility became difficult to evaluate because they described very different systems. A network with a small validator set could be marketed as decentralized. A token required only for payment inside a product could be described as essential utility. A consortium database could be described as blockchain even when a conventional shared system would have been simpler.

Semantic inflation does not merely confuse outsiders. It changes engineering incentives. When the category label is rewarded, teams begin with the label and search for a problem. Architectural choices become identity commitments. A database that should be replaced can persist because abandoning it would invalidate the financing narrative. Technical debt becomes narrative debt: the company must maintain the story that justified its creation.

The Bank for International Settlements warned in 2018 that cryptocurrencies raised questions about scalability, energy use, governance, and the ability to provide the trust functions of money and payment systems (BIS 2018). Some criticisms were too broad when applied to the whole field, but the report identified a central tension. Systems designed to remove trusted intermediaries often reintroduced trust through exchanges, wallet providers, stablecoin issuers, protocol developers, and governance bodies. The architecture did not eliminate institutions; it relocated them.

The investment implication is not that decentralized systems are impossible. It is that decentralization must be specified as a distribution of rights, control, verification, and exit. The word itself is not a moat. A decentralized network can be defensible through liquidity, community, standards, developer adoption, and credible neutrality. It can also be fragile because governance is slow, incentives are unstable, and users concentrate around centralized gateways.

A technically disciplined analysis begins with the trust model. Which parties cannot rely on a common operator? What adversarial behavior must the system tolerate? Which records require shared verification? Who can update the rules?

How are errors reversed? What economic value is created by making the system harder to change? Without explicit answers, decentralization becomes a stylistic preference rather than a functional requirement.

The same discipline applies to token economics. A token can coordinate a network when it allocates scarce resources, rewards verifiable contribution, or provides security through economically meaningful stakes. It becomes ornamental when the product could operate without it and the token exists primarily to create a financing or trading event. The distinction is not ideological. It is a test of whether value accrues to the instrument in proportion to network use.

TradeLens illustrates a different boundary. The shipping platform developed by Maersk and IBM sought to improve document exchange and visibility across a fragmented global ecosystem. The technical problem was real and the platform achieved significant participation, yet the sponsors discontinued it in 2022 because the level of commercial collaboration required for viability had not been achieved (A.P. Moller - Maersk 2022). The case is often used to criticize enterprise blockchain, but its deeper lesson concerns governance and adoption. A shared ledger cannot manufacture incentives for competitors to contribute data, change processes, or accept a common sponsor.

This matters because technological architecture and institutional architecture are complements. A system designed for multiple organizations needs rules for contribution, benefit sharing, liability, dispute resolution, and control. Technical immutability can even make governance more difficult when commercial relationships require correction and exception. Investors who underwrite only the protocol miss the transaction costs that determine participation.

The cycle also revealed how category degradation changes exit opportunities. Once customers associated blockchain with speculative projects, serious enterprise vendors often reduced the visibility of the term or repositioned around digital assets, distributed ledgers, or specific workflow outcomes. The underlying research continued, but the category label lost information. Strong projects paid a legitimacy tax created by weak ones.

The effect was not uniform. Public blockchains developed durable networks, stablecoin markets expanded, and cryptographic research advanced. Some applications created significant value. The mistake is to infer from these successes that the boom's full project population represented technological progress. Market capitalization, token liquidity, and development activity measured different things. The cycle produced infrastructure, speculation, fraud, experimentation, and genuine innovation simultaneously.

For investors, the central distinction is between narrative optionality and architectural commitment. Narrative optionality is valuable early because it attracts resources and keeps strategic paths open. Architectural commitment becomes valuable later when the company proves that a specific design creates an advantage competitors cannot obtain more cheaply. A project that remains at the narrative stage can continue raising money while accumulating little defensibility.

The investor lesson is not that decentralized architectures lack value. It is that technical architecture and financing language can separate. The more valuable

the category label becomes, the greater the incentive to apply it where the trust model does not require it. The resulting semantic inflation makes serious projects harder to price because they inherit the diligence cost created by unrelated claims. A research-driven investor responds by decomposing the label into adversaries, coordination failures, ownership rules, verification requirements, and the institutional complement that makes the architecture operable.

Accumulated Complexity: Integration Is an Organizational Problem Expressed in Software

Enterprise integration is difficult for reasons that no messaging protocol can remove. Organizations accumulate systems across decades, acquisitions, jurisdictions, departments, vendors, and regulatory regimes. The systems encode different definitions of customers, products, identities, risk, time, and authority. Integration therefore requires more than moving data. It requires reconciliation among institutional histories.

Enterprise Integration Patterns provided a rigorous vocabulary for messaging, routing, transformation, orchestration, and error handling (Hohpe and Woolf 2003). Enterprise service buses attempted to operationalize service-oriented architecture by providing common infrastructure for connecting heterogeneous applications (Schmidt et al. 2005). These contributions were technically useful. The market problem emerged when an abstraction for managing complexity was marketed as a mechanism for eliminating complexity.

A bus can standardize transport while leaving semantics unresolved. A canonical data model can reduce pairwise mappings while becoming a contested political object. An orchestration layer can centralize visibility while turning into a bottleneck. Governance can prevent chaos while slowing change. Every solution moves complexity among layers. The customer experiences disappointment when the product promise is framed as removal rather than redistribution.

The economic incentives favor optimistic framing. Integration projects are expensive, and buyers seek a platform that converts bespoke work into reusable configuration. Vendors respond with broader adapters, visual tools, governance modules, transformation languages, monitoring, security, and process engines. The expanding feature set helps win comparisons and increases switching costs. It can also create a product whose internal complexity mirrors the enterprise it was intended to simplify.

Foote and Yoder's "Big Ball of Mud" described software systems whose structure becomes difficult to understand because expedient changes accumulate faster than architectural discipline (Foote and Yoder 2000). Lehman's laws of software evolution similarly emphasize that systems used in changing environments must continually adapt and tend to become more complex unless work is devoted to reducing complexity (Lehman and Belady 1985; Godfrey and German 2014). Enterprise integration products sit directly in

this dynamic. They connect systems that are themselves evolving, so every interface is exposed to change on both sides.

The framework cycle follows a recurring pattern. A new architecture identifies genuine limitations in the previous generation. It offers cleaner boundaries, easier deployment, or greater autonomy. Early adopters benefit because they can avoid historical constraints. As adoption expands, organizations demand compatibility, governance, observability, security, and migration support. The new architecture acquires the operational features of the old one under different names.

Microservices illustrate the mechanism. They can improve team autonomy and deployment flexibility, but they introduce distributed failure, network latency, versioning, observability, and data-consistency problems. Fowler's discussion of microservice trade-offs warned against treating the style as an automatic improvement (Fowler 2015). The market nevertheless rewarded the label, and organizations sometimes decomposed systems before possessing the operational maturity to manage the result.

Machine learning systems add another layer. Sculley and colleagues described hidden technical debt arising from data dependencies, feedback loops, configuration, and changes in the external world (Sculley et al. 2015). Generative AI intensifies the problem because outputs are probabilistic, model providers change, prompts and retrieval systems interact, and evaluation is context dependent. The interface may be simple while the production system requires more monitoring, fallback, and governance than conventional software.

The lesson for investors is that enterprise complexity is not a temporary obstacle waiting for a universal platform. It is a renewable economic resource. Organizational change creates new exceptions faster than architecture can eliminate them. Vendors can build durable businesses by managing this complexity, but they must be honest about where the labor resides and whether each implementation compounds.

The phrase "complexity harvester" describes a company that enters the difficult layer rather than avoiding it. It finds the needle in the haystack: the unusual invoice, conflicting policy, fragmented identity, undocumented dependency, or cross-system exception that prevents automation. The business becomes defensible only if resolving each case improves a reusable taxonomy, data asset, control system, or operating process. A services organization that repeatedly solves the same type of exception can become a software institution. A services organization that merely adds people becomes capacity, not a moat.

Enterprise software markets often move through proliferation and consolidation. A new platform creates technical possibilities, startups specialize in narrow functions, and customers assemble stacks. Over time, integration cost, security review, vendor management, and overlapping features increase. Buyers begin to prefer suites, systems of record, or a smaller set of strategic vendors. Incumbents acquire capabilities and bundle them into broader contracts.

This pattern is not evidence that specialization was useless. Startups discover use cases and create pressure that incumbents would not generate internally. The problem is value capture. A narrow vendor may create an

important workflow improvement while remaining dependent on data, identity, permissions, and distribution controlled by a larger platform. Its product can be valuable to the customer and still be more valuable as a feature inside the platform.

Platform envelopment theory explains part of this movement. An actor with an existing user relationship can enter an adjacent market by bundling the target function with its own platform, leveraging shared users and network effects (Eisenmann, Parker, and Van Alstyne 2011). Enterprise incumbents possess additional advantages: procurement status, security certification, contracts, implementation partners, and access to systems of record. A startup may be technically superior yet impose another vendor relationship.

The result is a paradox. The more serious the workflow, the more valuable the startup's specialization may be, but the more demanding the integration and assurance burden becomes. Customers want innovation without fragmentation. This creates an acquisition path for startups and a consolidation path for the market.

Economics can still support independent companies when specialization produces authority rather than a feature. A vendor that owns the operational record, defines the workflow, carries regulatory responsibility, or coordinates multiple parties is harder to absorb. A vendor that merely improves an interface around another system's data is exposed. The distinction is not vertical versus horizontal in a superficial sense. It is whether the company controls a consequential boundary.

This case contributes to a different lesson from the previous one. The problem is not narrative without need. The need is undeniable. The risk is that every generation sells another layer over unresolved organizational complexity. AI can create a better abstraction, but it cannot abolish the enterprise. The companies that make money will be those that learn to govern the complexity rather than merely rename it.

This history matters because it reverses a common venture instinct. Complexity is not always technical debt waiting to be deleted. In mature organizations, some complexity encodes real differences in law, accountability, timing, ownership, and process. A vendor that merely adds another abstraction may increase the stack. A company that becomes authoritative over exceptions, semantics, and cross-system consequences can turn the same complexity into a compounding asset. The opportunity lies less in promising universal simplicity than in learning why the system resists it.

Finite Interface Capacity: Digital Supply Meets Human Limits

Consumer software appears to offer unlimited shelf space. App stores can distribute millions of products, cloud infrastructure can serve global demand, and digital goods can be replicated at negligible marginal cost. The binding constraint is human. Users possess finite attention, memory, trust, payment

tolerance, and capacity to form habits. An application can be available everywhere and still have no durable place in a person's life.

Large-scale smartphone studies consistently find that a small set of applications dominates habitual use even when users can choose among hundreds of thousands of alternatives. In a longitudinal iOS dataset, Kim and colleagues found that ninety percent of users launched roughly fourteen to eighteen applications in a typical week. Sigg and colleagues documented severe early attrition, with applications losing about sixty-five percent of users during the first week and only a small minority sustaining durable popularity (Kim et al. 2019; Sigg et al. 2019). The exact count varies by population and method. The durable result is concentration of habit, not scarcity of supply.

AI expands the number of functions that can be offered while also expanding the ability of one interface to absorb them. A general assistant can write, summarize, search, translate, plan, analyze, tutor, code, and interact with tools. Specialized applications may perform individual tasks better, but the user pays a cognitive switching cost to remember which system is best, move context, manage identity, and evaluate another privacy relationship. As general assistants improve, convenience can outweigh narrow performance differences.

The consumer startup faces three simultaneous costs. It must acquire attention, earn trust for sensitive input, and preserve relevance after the platform adds similar functionality. Paid acquisition can generate installs, but the company may be renting users from the same platforms that can later compete with it. Virality can reduce cost, but many utility tasks lack social transmission. Subscription revenue can be attractive, but consumers rationalize overlapping tools.

The interface budget becomes more restrictive when AI is embedded into existing products. A writing function inside email, a search function inside a browser, or an analysis function inside a workplace suite does not require a new habit. The incumbent can accept lower standalone revenue because the feature increases retention or supports a broader bundle. A startup that must recover acquisition and support cost through one narrow subscription competes against a different objective function.

This does not eliminate specialized consumer companies. Some categories possess communities, identity, content, or trust that a general assistant cannot easily reproduce. Health, finance, education, creativity, and companionship can support differentiated experiences when the product develops a recognizable method, data advantage, or brand relationship. The startup must become more than a prompt surface. It must own a durable reason for the user to return.

A powerful consumer moat often begins with identity rather than functionality. Users return to products that express status, belonging, taste, progress, or personal history. General assistants are broad and substitutable at the task level. A specialized product can become non-substitutable if it accumulates a meaningful record, community, or practice. The challenge is that such assets take time while AI makes functional imitation immediate.

The economics of early wrappers illustrate the distinction. Some thin applications have generated substantial revenue through speed, distribution, and subscription optimization. Their success demonstrates real entrepreneurial

skill. It does not prove long-term independence. A cash-generating product can be rational for a small team even when the expected terminal value is modest. The same company may be unsuitable for a large venture fund whose return depends on durable category control.

The investor should therefore separate market growth from company defensibility. Billions of hours of AI use can support a huge market while most application providers fail to retain bargaining power. The aggregate opportunity says little about the distribution of returns.

Founders sometimes compare AI applications with games because both markets are hit driven and can produce breakout consumer products. The analogy is incomplete. Games are creative goods consumed partly for their difference. A player may seek a new world, mechanic, story, competition, or community even when another game already satisfies the general need for entertainment. Functional redundancy is not necessarily a weakness because novelty is part of the product.

Caves described creative industries through the "nobody knows" principle: demand is uncertain, quality is difficult to predict, and differentiated projects coexist despite high failure rates (Caves 2000). A successful game can build intellectual property, characters, communities, user-generated content, and mastery. The player does not ask whether one title can summarize the other. Distinctiveness creates consumption.

AI utilities are usually consumed to minimize effort. Users prefer the tool that completes the task with acceptable quality, trust, and friction. When several applications offer comparable summarization, transcription, image cleanup, or planning, the market tends toward aggregation. The best function may survive as a specialized professional tool, but the median consumer feature can be absorbed into a general interface.

The difference affects portfolio logic. A game investor can rationally finance differentiated creative bets because one hit may become a franchise. An investor in generic AI utilities is financing many products that may converge toward the same capability. Correlation is higher than it appears. A major model release, operating-system feature, or pricing change can impair an entire portfolio at once.

Digital dating provides a narrower lesson about metrics. Dating platforms create value by improving discovery and matching, yet their business models can depend on continued engagement. The user's ideal outcome may be to leave the platform after forming a relationship, while the platform measures swipes, messages, subscriptions, and return frequency. This does not mean the platform intentionally prevents successful matches. It means intermediate engagement and final welfare are not identical.

Finkel and colleagues concluded that online dating expanded access but that strong claims about algorithmic matching often exceeded available evidence (Finkel et al. 2012). Pronk and Denissen found that choice overload could produce a rejection mindset in online dating (Pronk and Denissen 2020). Pew's research documented a mixture of positive experiences, frustration, harassment, and skepticism among users (Pew Research Center 2023). The

category shows how more options and more measurable activity can coexist with uncertain improvement in the outcome users ultimately seek.

The analogy matters for consumer AI because many products optimize interaction rather than resolution. A companion, coach, tutor, or productivity assistant may benefit from sustained engagement, yet the user's goal may be autonomy, learning, or completed work. Metrics such as session length and message volume can reward dependence. Goodhart pressure appears when the measurable proxy becomes the monetization target.

A durable consumer brand must therefore define the outcome it is willing to sacrifice engagement to achieve. This is difficult because advertising and subscription models reward use. The strongest companies can align revenue with progress, transaction, or verified result. Others may rely on reputation: users stay because the product is trusted not to exploit them. Brand becomes a governance promise.

The constructive path for a consumer startup is narrower but not empty. It can own a community, a trusted vertical identity, a proprietary content universe, a regulated relationship, a transaction, or a longitudinal data record. It can build a method that users recognize and a brand willing to make testable promises. It can integrate into a device or physical practice. What it cannot safely assume is that a convenient interface around general intelligence will remain a separate category.

The market will contain many profitable small companies, particularly those with disciplined acquisition, localization, and niche distribution. Their fragility should be reflected in valuation and financing structure. A lifestyle business can be attractive when it maintains low fixed cost and harvests cash. It becomes fragile when growth capital assumes a durable moat that does not exist or when the business depends on one model API, one app store, and one acquisition channel.

The three histories now separate cleanly. The first concerns financing claims before necessity. The second concerns the survival of real complexity beneath new technical layers. The third concerns demand concentration when supply is cheap. Together they show why more entry and more features do not imply more durable companies. The common investor error is to price the visible product while ignoring the market institution that will ultimately govern it.

Table 6. The three historical mechanisms and their investment use

TERM OR ASPECT	DESCRIPTION
Narrative capital	Use when category association can attract financing before the architecture is shown to be necessary under normal technical, economic, and governance conditions.
Accumulated complexity	Use when the customer problem consists of heterogeneous systems, exceptions, semantics, and authority that cannot be removed by a new interface alone.
Finite interface capacity	Use when digital supply is abundant but user attention, habit, trust, identity, and subscription tolerance remain bounded.
Do not overgeneralize	The mechanisms are evidence about market behavior, not claims that every future technology will replay the same

TERM OR ASPECT	DESCRIPTION
	sequence or that every company in the historical category was weak.

PART III

THE AI CONSOLIDATION INVESTMENT THESIS

AI changes what is scarce and destabilizes the boundaries of software products. The central industrial forecast is that capability will spread broadly while bargaining power concentrates around platforms, systems of record, research frontiers, and organizations able to own consequences.

CHAPTER FOUR

Capability Abundance and the Death of the Generic Product

This chapter explains why falling production cost does not distribute value evenly. It develops the production-verification asymmetry, distinguishes useful wrappers from durable applications, and treats model convergence as an underwriting assumption rather than a technological slogan.

Abundant Intelligence and the Production-Verification Asymmetry

General-purpose technologies create value because they spread across many sectors and stimulate complementary invention (Bresnahan and Trajtenberg 1995). Foundation models possess this property in an unusually direct form. The same underlying system can support writing, coding, analysis, search, image generation, customer service, scientific work, and tool use. Natural-language prompting makes the capability programmable by a much broader population than previous software platforms.

Burkhardt and Rieder argue that foundation models should be understood as platform models because prompting turns general capability into an interface through which others build applications (Burkhardt and Rieder 2024). The platform interpretation matters economically. Application developers do not merely buy an input. They build on an actor that can observe demand, change prices, improve the base product, modify policies, and enter adjacent functions.

The cost trend strengthens the platform effect. Stanford's 2025 AI Index reported that the cost of querying a model at roughly GPT-3.5 benchmark performance fell more than two hundredfold between late 2022 and late 2024, while smaller models reached performance levels that previously required much larger systems (Stanford HAI 2025). The 2026 report described continued capability gains, convergence among leading providers, and widespread organizational adoption (Stanford HAI 2026). These trends make intelligence cheaper, more substitutable, and easier to embed.

Lower cost is good for customers and for economic growth. It expands the number of tasks worth automating and enables startups to test ideas with little capital. It also removes scarcity from the component around which many early companies were formed. A product whose advantage was access to capable text generation loses differentiation when several providers offer comparable output at declining price.

This process resembles commoditization but should not be oversimplified. Models differ in reliability, latency, modalities, tool use, privacy, and domain performance. Frontier capability remains expensive to create, and the best system for one task may not be best for another. The investment question is not

whether all models become identical. It is whether the customer's willingness to pay for the application depends on a difference that remains visible after model substitution.

The strongest economic effect is a production-verification asymmetry. AI makes candidate generation fast. It can produce many designs, code paths, analyses, messages, or diagnoses. The cost of determining which candidate is correct in context may decline more slowly. Verification requires access to ground truth, domain expertise, testing, permissions, and responsibility. In consequential settings, the last mile is not a cosmetic layer. It is the product.

The asymmetry is visible in empirical work. AI tools have improved productivity on bounded tasks, yet performance depends on whether the task lies within the model's competence and whether users can evaluate output (Noy and Zhang 2023; Dell'Acqua et al. 2026). Shumailov and colleagues showed that recursive training on generated data can degrade model quality, illustrating how abundant synthetic production can pollute future information environments when provenance and verification are weak (Shumailov et al. 2024). AI can accelerate knowledge production and epistemic noise at the same time.

This changes the locus of value. When answers are scarce, the answer generator earns rent. When answers are abundant, the organization that can authorize one answer, connect it to action, and bear the consequence becomes scarce. A hospital does not need only diagnostic suggestions. It needs validated workflow, clinical responsibility, privacy, audit, and escalation. A bank does not need only document classification. It needs policy alignment, evidence, controls, and accountability. A factory does not need only predictions. It needs integration with physical operations and safe response.

The asymmetry also affects marketing. A startup can generate polished demonstrations, benchmark narratives, and customer-specific prototypes rapidly. Investors receive more evidence-like material but not necessarily more evidence. The cost of producing a claim declines, while the cost of verifying deployment remains. Diligence that does not change becomes less informative. The empirical record supports both optimism and caution. Generative AI has increased output and quality in bounded tasks, with especially large gains for less experienced workers in customer support and selected knowledge work (Brynjolfsson, Li, and Raymond 2025; Dell'Acqua et al. 2026). Other field evidence finds meaningful time savings without immediate changes in task composition, and economy-wide effects remain difficult to identify after only a few years of adoption (Dillon et al. 2025; Humlum and Vestergaard 2025). The correct conclusion is not that AI is weak. It is that capability, organizational change, and captured profit move on different clocks.

The Stanford AI Index documented a decline from roughly twenty dollars to seven cents per million tokens for systems reaching a GPT-3.5-like MMLU score between late 2022 and late 2024, a reduction of more than two orders of magnitude (Stanford HAI 2025). The 2026 report also described clustering among leading models on several public evaluations and continued rapid company formation and investment (Stanford HAI 2026). Such convergence does not mean every model is identical. It means an investor should assume

that many visible differences will narrow faster than the company's financing horizon.

The production-verification asymmetry is the decisive consequence. AI can create more candidate code, documents, designs, decisions, and experiments than an organization can responsibly validate. Where the consequence is trivial, the asymmetry creates consumer surplus. Where the consequence is material, it shifts value toward ground truth, evaluation, permission, escalation, and liability. The product is no longer the answer. It is the institution capable of deciding when the answer can act.

The Wrapper Test

A wrapper is often defined as an application that calls a foundation-model API and adds an interface. The definition is technically unhelpful because most modern software depends on external services. A payment company wraps banking rails, a marketplace wraps logistics and identity, and a cloud application wraps databases and operating systems. The strategic issue is not dependence. It is whether the company adds assets that change bargaining power.

A replaceable wrapper has three properties. Its output is largely determined by a model available to competitors. Its customer context can be moved with low cost. Its distribution does not improve through operation. Such a product can be useful and profitable, but it has little protection when the supplier improves, a rival copies the interface, or an incumbent bundles the function.

An accumulating application uses the model as a component inside a system that gains strength. It may acquire lawful outcome data, proprietary evaluation, workflow authority, customer trust, integration depth, or a network of participants. The model can be replaced without replacing the company. This reversibility is one of the clearest tests of a durable AI business.

Generic application startups face a four-sided squeeze. Upstream model providers can expand into memory, retrieval, tools, agents, identity, and transactions. Downstream enterprise incumbents can embed AI into systems that already hold data and permissions. Lateral startups can reproduce visible features quickly. Customers can assemble internal solutions as tools become easier to use. The startup must create value faster than all four forms of substitution.

The upstream risk is not limited to feature copying. A model provider can change price, rate limits, safety policy, context handling, data terms, or product architecture. Even favorable changes can destroy differentiation. If the provider adds a capability that the startup spent a year building around, the customer may receive better performance without changing vendors.

The downstream risk is distribution. Enterprise systems possess identity, procurement, security, and workflow access. Consumer platforms possess habitual interfaces and app-store control. A startup may outperform the incumbent's feature but still lose because the buyer prefers one contract and one data boundary. Bundling changes the comparison from product quality to total relationship cost.

The lateral risk is semantic imitation. Generative tools reduce the time required to copy visible workflows, landing pages, onboarding, and integrations. A company can still possess deep operations behind the surface, but the market may not observe the difference. Customer acquisition becomes more expensive because every claim has close substitutes.

The customer-insourcing risk is uneven. Large enterprises can combine internal data, cloud services, open-source models, and consultants. Smaller firms may prefer packaged products. The direction nevertheless matters. When model APIs and orchestration tools become easier, the threshold for building a narrow internal application falls. A startup must offer more than development convenience.

Schrepel and Pentland emphasize that competition among foundation models is shaped by compute, data, cloud relationships, and ecosystem effects (Schrepel and Pentland 2024). The market can be competitive at the model level and concentrated at the infrastructure level. Application companies should therefore avoid confusing provider multiplicity with supplier independence. Several APIs may rely on the same clouds, chips, and distribution channels. Commoditization is not a single event. It is a moving boundary. Today's rare capability becomes tomorrow's API, and tomorrow's API becomes an operating-system function. Startups survive by moving the boundary of what they own. A company may begin as a wrapper because speed is rational. It becomes investable when the wrapper is a wedge into a harder position.

This creates a strategic clock. The company has a limited period in which temporary differentiation can be converted into assets. Growth that does not improve defensibility may shorten the clock by attracting platforms and imitators. The investor should therefore ask not only how fast the company is growing but what each unit of growth accumulates.

The market can be large and the independent opportunity small. This is the central analytical error in many AI investment theses. Total economic value created by generative intelligence may reach extraordinary scale. Most of that value can accrue to users, incumbent complements, infrastructure, and a few broad platforms. Market size does not guarantee a large population of durable application companies.




Figure 3. The four-sided squeeze on generic AI applications

The word wrapper should therefore be used as an economic diagnosis, not an insult. Every technology company wraps lower layers. The question is whether the higher layer accumulates independent bargaining power. A payment company built on banking rails may own risk models, regulatory permission, merchant distribution, and transaction data. An AI application built on a model API can be equally substantive if it owns outcomes, authority, and learning rights. Dependence is normal; replaceability is the problem.

Outfinity's burden-of-proof rule is strict. Assume that the current model supplier can be replaced by a competitor, that the visible feature can be copied, and that the customer can reproduce the demonstration. The company must remain necessary because it controls the harder system around the model. If the investment case fails under these counterfactuals, the company is a product opportunity but not yet a venture-scale institution.

Table 7. The economic wrapper test

TERM OR ASPECT	DESCRIPTION
Model substitution	The company remains valuable if the current model is replaced without losing the customer relationship, operating knowledge, rights, or evaluation system.
Context portability	A rival cannot reproduce the outcome merely by importing the same documents, prompts, and configuration into another interface.
Workflow authority	The company can recommend, approve, route, execute, or govern consequential action rather than only generate content.
Learning rights	The company can lawfully connect decisions to outcomes and reuse what it learns across time, customers, or deployments.
Distribution compounding	Usage lowers future acquisition cost or deepens an installed relationship rather than relying permanently on purchased traffic.

TERM OR ASPECT	DESCRIPTION
Exception ownership	Difficult cases improve reusable systems instead of returning as unpriced manual work or being transferred to the customer.

CHAPTER FIVE

Why AI-Native Products Become Features

This chapter moves from capability economics to industry structure. It explains why easier entry can coexist with fewer durable categories, why platform envelopment and vendor rationalization favor broader firms, and why many apparently successful exits will be priced as feature acquisitions.

Entry, Envelopment, and the Consolidation Machine

The intuitive model of competition predicts that lower entry costs produce more firms and more persistent variety. Digital markets complicate the prediction. Lower entry costs create more experiments, but network effects, bundling, data, distribution, and fixed investment can concentrate value. The population of startups can expand while the number of independent strategic positions contracts.

AI intensifies both forces. Generative tools lower the cost of building and selling a product. Recent research suggests that exposure to generative AI is associated with increased startup formation and changes in the scale and composition of entrants, although the evidence remains early and context dependent (Bena, Bian, and Giannetti 2025; Gupta et al. 2026). Entry is real. The question is whether the entrants become enduring firms or discovery mechanisms for larger platforms.

John Sutton's theory of endogenous sunk costs is highly relevant. In some markets, firms can escalate spending on research, advertising, quality, or distribution as the market grows. Because the expenditure raises perceived or actual quality, concentration need not decline with market size (Sutton 1991, 1998). AI creates exactly such escalation opportunities. Frontier research, compute, safety, brand, ecosystem development, and global distribution can absorb enormous capital. A larger market can support a more expensive race rather than more equal competitors.

The implication is a barbell, but not a simplistic one. At one pole are organizations able to finance frontier research, infrastructure, chips, energy, robotics, and broad platforms. At the other are focused companies that own a narrow but consequential institution. The exposed middle contains generic products whose capability can be supplied by either pole. They are too narrow to command platform power and too shallow to control an outcome.

Platform envelopment provides the direct mechanism. A platform can enter an adjacent category by bundling the target function with its existing relationship, turning the complement's standalone product into one feature of a broader offer (Eisenmann, Parker, and Van Alstyne 2011). Foundation-model providers can envelop applications because the model is already the technical substrate.

Enterprise incumbents can envelop them because the workflow, data, identity, and contract are already present.

Wen and Zhu found that app developers changed innovation and pricing in response to the threat of platform-owner entry, shifting effort away from affected categories (Wen and Zhu 2019). The welfare effect can be ambiguous because platform entry may reduce redundant development while also weakening complementor incentives. This ambiguity is central to AI. Many standalone features will be more efficiently supplied inside a general assistant or suite. The social value of consolidation can coexist with poor returns for the startups that discovered the feature.

Bundling strengthens the effect because the platform need not earn the same standalone margin. A broad provider can price an AI function near zero if it increases retention, cloud consumption, advertising, or suite adoption. The startup must cover acquisition, support, and model cost directly. Even a superior product can lose when the comparison is between a narrow subscription and an included capability.

Customers contribute through vendor rationalization. Every AI supplier creates a new data relationship, security review, legal process, integration surface, and operational dependency. As experimentation becomes production, enterprises prefer fewer accountable vendors. Consumer users similarly rationalize subscriptions and habits. Consolidation is therefore pulled by demand as well as pushed by supply.

The mechanism is accelerated by category overlap. A summarizer becomes a research assistant, which becomes a document workspace, which becomes an agent platform. Product boundaries are unstable because the underlying model is general. Companies that appear to occupy separate markets can become substitutes after one capability release. Investors who construct portfolios by current labels may underestimate common exposure.

Increasing returns turn small differences into structural advantages. A broad platform gathers more usage, can distribute new features immediately, and can cross-subsidize experimentation. A vertical operator gathers more domain outcomes and earns deeper trust. The generic middle accumulates less because its customer relationship is thin. It becomes an acquisition target or a temporary layer.

The industrial pattern is a classic shakeout accelerated by general-purpose technology. Beraja and Buera model concentrated industries as moving from intense entry to exit and eventual concentration when innovation creates large scale differences (Beraja and Buera 2026). AI supplies unusually strong conditions for this path. The input capability is broad, the product boundaries overlap, the distribution assets are concentrated, and the cost of offering one more function inside an existing platform is lower than the cost of sustaining a separate company around it.

Platform envelopment is not the same as superior product innovation. A provider enters an adjacent category by bundling the target function with an existing user relationship and by using network effects or switching costs that were created elsewhere (Eisenmann, Parker, and Van Alstyne 2011). The startup may remain technically better. The platform can still win because the

comparison moves from feature quality to total relationship cost, data boundary, procurement simplicity, and convenience.

Enterprise vendor rationalization strengthens the pressure. Every narrow vendor creates recurring costs that do not appear in the product's subscription price: security review, access control, integration maintenance, contract management, monitoring, incident response, and employee training. Once experimentation ends, chief information officers and procurement teams have rational incentives to reduce the number of suppliers. An incumbent suite can offer a weaker feature and still deliver a lower total coordination cost.

Customer insourcing is the fourth mechanism. Generative systems do not eliminate the need for internal technical talent, but they reduce the cost of assembling a narrow workflow from models, cloud services, open-source components, and consulting. The startup must therefore be better than internal development on the difficult parts: continuous maintenance, cross-customer learning, assurance, exception handling, and responsibility. Development convenience alone has a shrinking price.

Exit Geometry: When the Technology Survives but the Company Does Not

The claim that most startups fail is trivial. The more important claim is that many successful AI-native startups may still produce exits too small for conventional venture underwriting. They can have customers, revenue, and capable teams while lacking the strategic independence required for a large terminal value.

A tuck-in acquisition can be rational for all parties. The startup's feature becomes more valuable inside the acquirer's distribution. The acquirer avoids development time and obtains talent, integrations, and customer learning. Founders receive liquidity, and customers gain a better-capitalized supplier. The transaction can nevertheless be mediocre for a large fund if the exit value is small relative to ownership dilution and portfolio requirements.

Venture returns are highly concentrated. Jackson and Strebulaev's 2026 study of more than one hundred thousand venture professionals reported that a small share of investors generated most investment profits, reinforcing earlier evidence of power-law outcomes and persistent differences in access and skill (Jackson and Strebulaev 2026). A fund cannot treat every positive exit as equivalent. The same company can be an excellent angel investment and an irrelevant late-stage outcome.

The consolidation thesis therefore changes financing fit. A company with a credible path to a fifty-million-dollar acquisition may be attractive to founders, angels, or a small fund. It may be structurally wrong for a billion-dollar venture vehicle. Misalignment appears when the financing plan assumes platform-scale value while the product architecture points toward feature acquisition.

The kill-zone literature adds a darker possibility. Kamepalli, Rajan, and Zingales model how the prospect of acquisition by an incumbent platform can reduce complementor adaptation and financing for a superior entrant, weakening the entrant's standalone value (Kamepalli, Rajan, and Zingales

2022). Cunningham, Ederer, and Ma document killer acquisitions in pharmaceuticals, where some acquired projects were discontinued to preempt competition (Cunningham, Ederer, and Ma 2021). These mechanisms should not be mechanically transferred to AI, but they show that acquisition markets can alter innovation incentives before any transaction occurs.

In AI, the more common risk may be not deliberate suppression but absorption. The startup's technology survives as a feature, while its independent category disappears. The buyer may genuinely expand the product. The investor still loses the standalone upside.

The market will also contain a large number of acqui-hires. AI increases the value of teams that understand model behavior, product design, and domain integration. A company may be purchased for human capital after its product moat weakens. This can preserve some capital but rarely supports a power-law thesis unless the acquisition price reflects exceptional strategic urgency.

There are important counterexamples. Some application companies will become platforms because they own the customer relationship and expand faster than model providers can move downstream. Some will develop proprietary data and research. Some will become systems of record. Some will build trusted brands in categories where accountability matters. The consolidation forecast is probabilistic, not deterministic.

The investor should therefore map the likely acquirers before investing, not as a substitute for independence but as a test of category structure. If every plausible buyer can reproduce the feature internally and none needs the customer relationship, exit optionality is weak. If the company controls a scarce asset that several strategic buyers need, acquisition can be valuable. If the company can remain independent because customers depend on its authority, the venture case is stronger.

A direct forecast follows. The AI application market is likely to produce many more products than enduring companies, and many more useful companies than venture-scale outcomes. Consolidation will occur both through acquisition and through category expansion by large providers. Early revenue will not protect a company whose function becomes a default capability.

The constructive response is not to avoid the application layer. It is to finance companies that use applications to acquire a deeper position. The practical obstacle is conversion: local traction often depends on founder labor, favorable data, and exceptional customers, while scale requires those conditions to become a repeatable operating system before the consolidation clock expires.

The firm can fail as an independent category while its technology succeeds. This is the most important distinction for AI-native underwriting. A broad acquirer may distribute the function to millions of users, improve it with proprietary context, and create far more total value than the startup could have created alone. Yet the acquisition price may reflect only the acquirer's alternative cost, the startup's bargaining position, and the scarcity of the team. The investor captures a fraction of the downstream value because the complement, not the feature, owns distribution.

This makes acquisition strategy vehicle-specific. A venture studio can deliberately create narrow companies for strategic buyers when formation cost, ownership, and timing are disciplined. A business angel can earn an excellent multiple from a modest outcome. A large fund cannot use the existence of possible buyers as evidence of venture scale. Buyer concentration and roadmap volatility often reduce price precisely when the startup needs the exit most.

The book is firm here because the alternative assumption is costly. Outfinity should underwrite consolidation as the base case and require evidence for independence. It should not underwrite independence as the base case merely because no platform has copied the feature yet. Absence of current absorption is a window, not a moat.

Table 8. Consolidation mechanisms in AI markets

TERM OR ASPECT	DESCRIPTION
Platform envelopment	An adjacent platform bundles the startup function with an existing relationship and changes the comparison from standalone quality to total relationship cost.
Vendor rationalization	Customers reduce supplier count when security, integration, monitoring, and contractual burdens exceed the benefit of narrow specialization.
Category overlap	General models and suites blur product boundaries, enabling one interface to absorb functions that previously supported separate applications.
Endogenous sunk costs	Research, brand, assurance, compute, and distribution spending rise with market size, favoring organizations capable of repeated reinvestment.
Customer insourcing	Lower development cost allows capable customers to reproduce narrow workflows, leaving external vendors to justify their role through deeper complements.
Feature survival	The function continues inside a larger product even when the independent supplier captures little terminal value.

CHAPTER SIX

Why Traction No Longer Proves Venture Scale

This chapter separates evidence of local value from evidence of repeatable institutional scale. It examines pilots, hidden services, exception economics, and the return geometry that determines whether a successful company is appropriate for venture capital.

The Pilot Economy and the Hidden Service Layer

AI products are unusually easy to demonstrate because language and images make capability visible. A model can summarize a contract, answer a question, generate code, or classify a document within minutes. Enterprise adoption is much harder because the demonstration must become a controlled process connected to data, authority, and consequence. The gap between the two creates a pilot economy.

Pilots are useful experiments. They reveal demand, technical limits, and organizational resistance. They become misleading when the conditions of the experiment are treated as normal operations. A design partner may provide clean data, motivated experts, executive attention, permissive security, and a narrow scope. Founders may monitor every output and resolve failures personally. Inference may be subsidized and support cost omitted. The result is genuine but not representative.

Local traction is evidence that value exists under a particular boundary. Scale requires the boundary to widen without destroying economics or quality. The investor must identify which conditions were essential and whether they can be reproduced across ordinary customers.

Founder dependence is the first hidden variable. Early founders possess product knowledge, customer trust, and authority to improvise. They can recognize when a model is wrong, rewrite a workflow, negotiate with stakeholders, and work outside the contract. The customer's experience appears automated because exceptional humans absorb uncertainty. Hiring more people may preserve delivery but convert software margin into services margin.

Champion dependence is similar. An internal executive can force access to data, align departments, and protect a pilot from normal procurement. The company may interpret adoption as product pull when it is partly political sponsorship. Scale requires success in accounts without such a champion or a repeatable method for creating one.

Data subsidy appears when evaluation uses curated inputs, historical labels, or a small number of document types. Production data contain missing fields, changing formats, adversarial behavior, and organizational ambiguity. The

model's accuracy can remain high while the cost of data preparation makes the deployment uneconomic.

Exception invisibility is often the decisive issue. Demonstrations emphasize the common path because it is frequent and visually compelling. Rare cases determine loss, liability, and labor. A system that automates ninety percent of transactions may be valuable, but the remaining ten percent can require more expertise than the original process because the easy cases have been removed. Automation changes the composition of human work.

Gans and Goldfarb's O-ring automation model formalizes this complementarity. When production quality depends on multiple tasks whose performance multiplies, automating one task changes the value of improving the others. Adoption can be discrete and bundled rather than a sequence of independent substitutions (Gans and Goldfarb 2026). The implication for AI startups is that a strong component may create little value until surrounding tasks, controls, and skills are redesigned.

This helps explain why individual productivity gains do not automatically become firm-level profit. Dillon and colleagues found time savings when AI was integrated into tools used by workers, yet did not detect broad changes in task composition during the study period (Dillon et al. 2025). Brynjolfsson, Rock, and Syverson's productivity J-curve argues that general-purpose technologies often require complementary intangible investment before aggregate gains appear (Brynjolfsson, Rock, and Syverson 2021). AI access is not the same as organizational transformation.

The startup can create value by supplying those complements, but the business model changes. It must design workflow, train users, integrate systems, monitor outcomes, and manage risk. The company is no longer selling a model-assisted feature. It is becoming an operator or institutional partner. AI makes the pilot economy unusually persuasive because language, images, and code expose capability immediately. A founder can demonstrate value before the organization has solved data governance, role redesign, exception handling, or liability. The pilot is not false. It is a bounded experiment supported by subsidies that often disappear in production: founder intervention, an internal champion, curated inputs, permissive security, manual review, and unusual customer attention.

The diligence error is to treat the pilot boundary as representative. Scale requires the boundary to expand while preserving economics, reliability, and service quality. A useful test is to rebuild the revenue under ordinary staff, representative data, production security, steady-state support, and the expected distribution of rare cases. If the gross margin depends on invisible expert labor, the company has either discovered a service business or found the raw material for a learning system. The investment outcome depends on which one it becomes.

Scaling Through Exceptions Rather Than Around Them

The conventional software instinct is to avoid services because services scale with labor. In AI, avoiding services too early can prevent learning. The difficult cases contain the information required to build a durable system. The strategic question is whether human intervention compounds.

A non-compounding service resolves the customer's problem and leaves little reusable assets. The next engagement begins again. A compounding service records the case, normalizes the data, improves evaluation, creates a rule, strengthens a connector, and clarifies a contractual boundary. The next engagement becomes faster or more reliable. Over time, the service organization turns into a decision system.

The process resembles learning by doing. Dierickx and Cool emphasize that some strategic assets must be accumulated over time and cannot be purchased instantly because of time compression diseconomies, causal ambiguity, and asset interconnectedness (Dierickx and Cool 1989). Exception knowledge has this property when it is embedded in people, taxonomies, data, and customer relationships. A rival can copy the interface but not instantly reproduce years of adjudication.

The company must possess rights to learn. Many enterprise contracts restrict reuse of customer data and derived knowledge. Privacy, confidentiality, and regulation may prevent pooling. A startup that performs valuable work without negotiating learning rights builds the customer's asset rather than its own. The moat therefore begins in contract design.

Evaluation is another compounding asset. Generic benchmarks measure broad capability, but production systems require domain-specific tests, adversarial cases, calibration, and monitoring. A company that accumulates a trusted evaluation suite can route among models, detect regressions, and prove performance. The model becomes replaceable while the assurance system remains.

Operational authority can also compound. A product initially recommends actions. After repeated success, the customer permits it to route, approve, or execute. Authority increases switching cost because removal requires redesign of controls and responsibility. The company should not seek authority faster than its evidence supports. Premature autonomy converts technical error into institutional harm.

A scalable AI company therefore often moves through a sequence from assistance to bounded automation to managed outcome. The sequence is not a hidden product roadmap; it is a transfer of responsibility. Revenue quality improves when the company can price an outcome and control the variables required to deliver it. Liability also increases. The moat and the risk are the same boundary.

The diligence method follows directly. Investors should reconstruct economics under representative data, ordinary staff, steady-state model cost, explicit monitoring, and the full exception path. They should ask what happens when the customer champion leaves, the model changes, a security review tightens,

and the founders stop joining every call. The result may remain attractive. If so, the company has stronger evidence than most.

Scale also depends on sales structure. A startup can grow revenue by accepting heterogeneous custom work, but each contract expands the product surface. The pipeline can hide increasing implementation debt. A high backlog may signal demand or the inability to standardize. Investors should distinguish repeatable configuration from bespoke development.

The best founders use early revenue as research funding. They choose customers whose complexity reveals a reusable market, not merely customers willing to pay. They refuse custom work that does not strengthen the core. They measure the decline in intervention per unit of outcome. They treat service margin as an investment in asset accumulation rather than as evidence that the product is already scalable.

This chapter turns the consolidation thesis into an operating requirement. A generic startup has limited time to build a harder position. It cannot do so by hiding the labor required to reach the customer outcome. It must enter the labor, learn from it, and convert it into rights, process, evidence, and authority. The phrase "finding a needle in a haystack" is useful only if the company learns why the needle is difficult to find. A complexity harvester converts repeated investigation into a taxonomy of exceptions, a reusable evidence model, a rights position, and an operating method. Each new case should reduce uncertainty for the next case. If every engagement begins from zero, labor scales linearly and the company remains consultancy-like no matter how sophisticated the interface appears.

This is one reason services can be strategically positive in AI. Early service work exposes the constraints that a generic product cannot see. The danger is not human involvement; it is non-compounding involvement. The company should measure intervention per outcome, time to deployment, exception recurrence, cross-customer reuse, and the share of decisions that can move from recommendation to bounded execution. These measures reveal whether operations are becoming an asset or merely supporting revenue.

Return Geometry and Investor Fit

The phrase venture scale often obscures a simple arithmetic fact. A company can be successful in ordinary economic terms and unsuccessful for a particular fund. The relationship depends on entry valuation, ownership, dilution, follow-on capital, exit value, time, and the size of the vehicle. When AI lowers company-formation cost, the number of plausible investments rises. The distribution of outcomes need not become more favorable. More experiments can coexist with stronger concentration in the few companies that control infrastructure, distribution, or a consequential workflow.

Venture returns have long been highly skewed. Research on fund performance and entrepreneurial outcomes has repeatedly found substantial dispersion and persistence, while recent work using data on more than 100,000 United States venture professionals estimates that 90 percent of investment profits are generated by 5 percent of venture capitalists (Kaplan and Schoar

2005; Hall and Woodward 2010; Jackson and Strebulaev 2026). This concentration does not prove that only large companies should be financed. It does mean that a strategy relying on moderate exits must be reconciled with fund mathematics rather than narrated as though every acquisition were equivalent.

AI-native application companies create a particular return ambiguity. Their early development can be capital efficient, their growth can be rapid, and strategic buyers may value teams, customers, or integration knowledge. These features can produce frequent exits. Yet the same low construction cost and strategic substitutability can cap acquisition values. A buyer that can recreate the product may pay for speed, talent, or customer access but not for a durable monopoly rent. The founder may experience a life-changing outcome while a late-stage investor receives a mediocre multiple.

The distinction between a small exit and a failed thesis is therefore vehicle specific. A business angel entering before product-market fit can earn an attractive return from an acquisition that would not move a multibillion-dollar fund. A venture studio can create value by repeatedly forming companies around validated problems, shared infrastructure, and disciplined sale points. A family office can combine direct equity with strategic assets, patient ownership, real estate, industrial operations, or regulated market access. A traditional venture fund generally needs a credible path to a small number of outcomes large enough to return the fund.

This diversity should change diligence. The question is not whether the company can exit. Almost any capable team with customers may attract some strategic interest. The question is what kind of asset the buyer is purchasing and how many buyers can rationally pay. A feature tuck-in has a narrow buyer set and a price anchored to build-versus-buy economics. A workflow institution has a broader strategic value because it controls customer authorization, data flows, and operating responsibility. A research platform may attract buyers seeking a technical frontier, but its value depends on reproducibility, freedom to operate, key-person risk, and the cost of continuing the program.

Local traction creates another source of confusion. A company can dominate a country, language, profession, or customer segment through relationships and contextual knowledge. That position can be durable and valuable. It becomes venture-scalable only when the source of local advantage transfers or when the local market is itself large enough. Relationships that cannot be delegated, integrations that must be rebuilt, and regulatory knowledge that changes by jurisdiction may produce a strong regional business rather than a global software multiple. The correct conclusion is not to dismiss the company. It is to choose a financing model consistent with the replication burden.

Lifestyle businesses deserve equally careful treatment. AI permits very small teams to serve global niches and produce substantial cash flow. This can be an attractive form of entrepreneurship. The fragility arises when the company depends simultaneously on one model supplier, one acquisition platform, one app store, and one narrow function. Low headcount reduces fixed cost but does not create bargaining power. A durable small business needs the same strategic disciplines as a venture company, though at a different scale:

diversified routes to market, owned customer relationships, a recognized method, operational expertise, cash reserves, and the ability to switch suppliers.

The most demanding cases sit at the other end of the spectrum. Frontier models, chips, robotics, energy systems, scientific platforms, and regulated infrastructure may possess deeper technical moats but require sustained capital. Their risk is not “easy imitation”. It is that the company cannot finance the sequence of experiments and complementary assets required to reach a self-supporting position. Deep technology can fail with excellent science because the financing architecture is too shallow, because manufacturing or deployment is underestimated, or because the company reaches a bargaining point with insufficient runway.

An investor should consequently distinguish capital intensity from capital defensibility. Spending billions does not by itself create a moat. Capital becomes defensive when it purchases cumulative learning, scarce capacity, regulatory evidence, a production curve, or an installed base that later entrants cannot cheaply replicate. Conversely, a capital-light company can be deeply defensible when it controls rights, trust, and workflow authority. The relevant variable is not burn but the irreversibility and compounding value of what burn creates.

Portfolio construction must reflect these different geometries. A collection of apparently diverse AI applications may be exposed to the same model release, distribution policy, privacy rule, or enterprise bundling decision. The portfolio is correlated at the dependency layer even when end markets differ. True diversification comes from distinct control points, customer budgets, regulatory regimes, physical assets, research programs, and routes to value. An education assistant, a legal assistant, and a sales assistant built on the same external models and sold through the same acquisition channels may represent one platform bet disguised as three vertical bets.

The investor's comparative advantage should also be explicit. Access to capital is less differentiating when strong companies can choose among investors. Specialized diligence, customer access, scientific judgment, governance, follow-on capacity, and the ability to coordinate complementary actors can change a company's probability of survival. Jackson and Strebulaev findings are consistent with persistent investor-specific skill and with the self-reinforcing access of recognized investors (Jackson and Strebulaev 2026). In an AI market rich in persuasive artifacts, the scarce investor capability is not deal discovery alone. It is discriminating between acceleration and accumulation.

The practical conclusion is deliberately unsentimental. Many AI companies will deserve to exist without deserving venture-scale financing. Some will generate useful products, high margins, and respectable exits while remaining structurally transient. Others will look slow because they are building the rights, evidence, physical systems, or scientific lead time required for a larger outcome. Capital allocation improves when these paths are named early.

The practical implication is unsentimental. A small SaaS company can be profitable, resilient, and valuable to its owners while remaining a poor venture

asset. AI may increase this class of company because tiny teams can serve narrow markets efficiently. It may also make the class more fragile because suppliers, channels, and features change faster. Valuation should reflect expected cash generation and strategic optionality, not the terminal multiple of a category leader that the architecture cannot become.

Outfinity's investment committee should therefore separate four questions. Is there real customer value? Can the company become economically repeatable? Can it remain independent after consolidation pressure? Can the resulting outcome return the relevant capital vehicle? A positive answer to the first question is necessary and insufficient. The book's polemic is directed at the habit of allowing early traction to answer all four.

Table 9. What early traction proves and what it does not prove

TERM OR ASPECT	DESCRIPTION
Customer enthusiasm	Proves that a user experienced value under a particular boundary; does not prove representative implementation cost, authority, or independence.
Pilot revenue	Proves willingness to pay for an experiment; does not prove steady-state gross margin after founder labor, exception handling, and security requirements are normalized.
Fast growth	Proves that demand and distribution currently exist; does not prove that growth accumulates a moat faster than it attracts imitation and platform attention.
High retention	Proves continued use under the present product architecture; does not prove resilience to bundling, model convergence, or customer insourcing.
Local dominance	Proves advantage in a bounded geography, language, or network; does not prove transferability without losing the institutional conditions that produced the advantage.
Strategic acquisition interest	Proves that the asset has build-versus-buy value; does not prove category-leader economics or a price capable of returning a large fund.

PART IV

WHAT WINNING STARTUPS WILL LOOK LIKE

The preceding analysis is not an argument for retreat. It is a filter for identifying the small set of companies that can become more valuable as models and software become cheaper. Winners will own a non-generic core and will pair it with capital designed for long learning cycles.

CHAPTER SEVEN

The Non-Generic Core: Moats and IP After Code

This chapter defines defensibility when code, interfaces, and many model capabilities are reproducible. It treats a moat as a causal system of scarcity, compounding, appropriability, and customer consequence rather than as a collection of reassuring labels.

From Feature Advantage to System Defense

The word moat is used so broadly in venture markets that it often functions as reassurance rather than analysis. Founders describe superior design, faster execution, proprietary prompts, integrations, data, community, patents, and brand as moats. Each can matter. None is a moat simply because it exists.

A useful moat performs three economic functions. It excludes a rival from reproducing the customer outcome at comparable cost. It retains the customer when alternatives improve. It compounds through operation so that time strengthens the position rather than merely aging the product. A feature advantage may satisfy the first function for a few months. A durable company needs the system.

Barney's resource-based view argues that sustained advantage depends on resources that are valuable, rare, difficult to imitate, and not easily substituted (Barney 1991). Dierickx and Cool add that important strategic assets often accumulate through time-dependent processes that cannot be compressed by spending alone (Dierickx and Cool 1989). These ideas become more important when AI accelerates ordinary imitation. The scarce asset is increasingly the history that cannot be regenerated from the current interface.

The first analytical tool is the "moat half-life". Every advantage should be evaluated by the time and cost required for a competent rival, platform, incumbent, customer, or standard to neutralize it. A prompt library may have a half-life measured in weeks. A specialized evaluation suite may last longer. A regulatory license, longitudinal outcome dataset, or installed physical network can persist for years. The half-life is not fixed; it shortens when tools improve and lengthens when assets interact.

Interaction matters because moats are systems. Data without rights may be unusable. Patents without manufacturing or distribution may create little value. Workflow integration without trust may increase switching cost but also customer resentment. Brand without performance becomes vulnerable. The strongest position combines several imperfect assets so that a rival must reproduce the whole system.

This is why "proprietary data" is often a weak claim. Data can be abundant, stale, legally restricted, biased, or easily recreated. The economically relevant questions concern exclusivity, outcome linkage, freshness, coverage of difficult

cases, permission to learn, and whether the data improves a customer's value. A collection of customer documents is not automatically an asset the company can reuse.

Workflow authority is stronger because it creates institutional position. A system that merely drafts a recommendation can be replaced more easily than one authorized to approve, route, reconcile, or execute. Authority embeds the product in controls and responsibility. It also creates risk: the company must prove reliability and accept bounded consequences.

Distribution remains a moat when it compounds rather than when it is purchased. Paid acquisition is a channel expense. A brand that lowers diligence, a community that produces referrals, a standard that attracts participants, or an installed system that opens adjacent workflows can reduce future customer-acquisition cost. The distinction is between rented attention and owned demand access.

Switching costs require similar discipline. Friction can retain customers, but artificial lock-in invites regulation and motivates replacement. Productive switching costs arise because the system contains accumulated context, trusted controls, trained users, and operational history. The customer stays because replacement would destroy value, not because data is held hostage.

The investor should map the moat as a causal system. What asset is created by each transaction? Who owns it? How does it improve performance, margin, retention, or distribution? What event would cause it to decay? This map is more informative than asking whether the company has a moat in the abstract.

A useful formulation is that moat strength rises with scarcity, compounding, appropriability, and outcome relevance, and falls with substitutability. This is not a precise financial equation. It is a discipline for causal analysis. A scarce asset that does not matter to the outcome creates little pricing power. A valuable asset that the company cannot lawfully retain creates little shareholder value. An advantage that does not improve with use has a short half-life. A powerful system combines several layers so that the weakness of one does not collapse the whole.

The non-generic core is the part of this system that cannot be purchased as a common input. It may be a scientific method, a time-dependent learning curve, an exclusive permission, a physical network, a high-integrity outcome dataset, an authoritative workflow, or an accountable brand. The core need not be technically exotic. A company can own a difficult institutional position that a better model cannot recreate because the position depends on contracts, history, trust, and accepted responsibility.

Intellectual Property, Rights, and Freedom to Operate

AI complicates intellectual property because it changes both invention and imitation. Generative systems can assist research, code creation, design, and prior-art discovery. Patent activity in generative AI has accelerated sharply; WIPO reported in July 2026 that published generative-AI patent families rose substantially in 2024 and 2025 (WIPO 2026). More patents do not necessarily

mean stronger moats. They may indicate technological activity, defensive filing, and an emerging thicket.

Patents can protect a specific technical method when the claims are valid, enforceable, and commercially relevant. They are most useful when a rival must practice the protected method to achieve the outcome and when infringement can be detected. They are less useful for rapidly changing software implementations that can be designed around or for companies without resources to enforce them.

The USPTO's revised 2025 guidance reaffirmed that only natural persons can be inventors and that the ordinary inventorship standard applies when AI is used as a tool (U.S. Patent and Trademark Office 2025). The guidance provides legal clarity but does not solve the business problem. A patent establishes a right to exclude; it does not establish customer demand, freedom to operate, or the ability to implement at scale.

Trade secrets may be more valuable for data pipelines, evaluation methods, operating procedures, model routing, and manufacturing know-how. Their strength depends on secrecy, employee practices, contracts, and the difficulty of independent discovery. AI can weaken secrecy by accelerating reverse engineering and employee productivity. It can also strengthen a trade-secret strategy when the valuable knowledge is tacit and distributed across operations.

Copyright is relevant to code, content, and datasets but rarely protects functionality itself. Contract rights often matter more. Customer agreements determine whether the company can use inputs, outputs, feedback, and derived statistics. Supplier agreements determine access, pricing, and continuity. Data licenses, exclusivity, audit rights, and assignment clauses can create or destroy the moat before any technical issue arises.

Freedom to operate is the neglected side of IP. A company may own patents and still require licenses to use underlying technologies, data, standards, or open-source components. Shapiro's analysis of patent thickets and Heller and Eisenberg's anticommons framework show how overlapping rights can impede innovation even when each right is individually justified (Shapiro 2001; Heller and Eisenberg 1998). In AI, model licenses, training data, software dependencies, and sector regulation create a similar stack of permissions.

Open source changes the location of advantage rather than eliminating it. A company can build on open models and still own distribution, operations, data rights, evaluation, and trust. Open source can reduce supplier dependency and accelerate adoption. It can also eliminate licensing revenue and allow cloud or platform providers to capture complementary value. Lerner and Tirole's work on open-source participation and patent pools shows that governance and complementary assets determine who benefits from shared technology (Lerner and Tirole 2002, 2004).

A credible open strategy answers two questions. What is intentionally commoditized to increase adoption or reduce ecosystem friction? What remains scarce enough to support investment? Publishing code can create a standard, attract developers, and improve trust. It is not a business model unless the company controls a complementary asset.

The hierarchy of moats is contextual. A consumer creative tool may rely on brand and community. A medical company may rely on evidence, regulatory approval, and clinical integration. A semiconductor company may rely on process knowledge, design IP, supply relationships, and capital. The investor should not seek the same asset everywhere.

The most robust question is counterfactual. Assume that a competent rival receives the same foundation model, similar funding, and access to modern coding tools. What prevents it from reproducing the customer outcome within eighteen months? If the answer is only execution speed, the company has a lead. If the answer is a connected system of rights, learning, authority, and distribution, it may have a moat.

AI makes intellectual-property rhetoric easier and IP strategy harder. More patents can be generated, more code can be produced, and more technical alternatives can be explored. The investor should not count documents. The investor should identify the bottleneck that the documents protect, the ease with which infringement can be detected, the cost of enforcement, the availability of design-arounds, and the complementary assets required to commercialize the claim.

Contract rights often dominate formal IP in application companies. The right to retain derived insights, connect decisions to outcomes, reuse evaluation artifacts, and operate across customer boundaries may determine whether deployment creates an asset. A startup can possess proprietary data and still have no moat if the data are stale, unlinked to outcomes, legally non-reusable, or easily regenerated by the customer. Conversely, a carefully governed feedback loop can be highly defensible even when the underlying model is open.

Freedom to operate belongs inside the investment thesis from inception. A research company that discovers a superior method but depends on blocked manufacturing, exclusive datasets, or a patent thicket may create knowledge without capturing value. The strongest companies align invention, rights, and complementary assets before scaling makes redesign expensive.




Figure 4. The non-generic core combines several forms of compounding scarcity

Table 10. The AI-era moat stack and its reality test

TERM OR ASPECT	DESCRIPTION
Research frontier	Repeatedly solves problems unavailable through common models; later entrants cannot compress the learning history.
Rights and outcome data	Lawfully links actions to outcomes, preserves provenance, and reuses knowledge that customers and rivals cannot readily reproduce.
Workflow authority	Authorized to initiate, approve, route, or govern consequential actions inside the customer's operating structure.
Assurance and trust	Evidence, guarantees, continuity, and incident handling put future rents at risk when performance fails.
Physical or regulated assets	Devices, manufacturing, logistics, laboratories, licenses, or infrastructure create deployment knowledge and permissions beyond software.
Compounding distribution	The channel, community, or system position becomes cheaper and more valuable with use rather than rented through paid acquisition.

CHAPTER EIGHT

Research, Time, and Long-Duration Capital

This chapter makes the positive case for research-driven formation and long-horizon finance. It explains when capital creates a moat, why scientific uncertainty needs different milestones, and why patience must be conditional on learning rather than confused with tolerance for drift.

The Right to Outlearn

A non-generic company possesses a core that cannot be reconstructed from common tools and ordinary execution. The core may be a scientific discovery, a specialized model, a manufacturing process, a proprietary instrument, a unique dataset, a regulated operating system, or a trusted network of institutions. It does not need to be completely isolated from the market. It must create a learning trajectory that competitors cannot enter at the same point.

Fundamental research can supply such a trajectory because it produces knowledge before the commercial application is fully specified. The company learns what is possible, not merely how to package a known capability. This is particularly relevant in materials, biology, energy, robotics, semiconductors, cryptography, and scientific AI. The work is slow, uncertain, and often capital intensive. Its value lies in changing the feasible set.

Deep technology is therefore not defined by technical difficulty alone. Many enterprise integrations are difficult but not deep tech. A deep-tech core has uncertainty about whether a new technical principle or performance frontier can be achieved, and success creates an asset with broader commercial consequences. The company must still solve commercialization, regulation, manufacturing, and adoption.

Gans and Stern distinguish strategies for commercializing technological innovation through cooperation with incumbents or competition in product markets (Gans and Stern 2003). The choice depends on intellectual-property strength and access to complementary assets. A startup with strong patents but weak manufacturing may license. A startup with weak formal IP but unique capabilities may need to build the full product to prevent appropriation. The financing plan should follow the commercialization regime.

Capital becomes a moat when it purchases time-dependent learning that cannot be compressed later. Training frontier models, building fabs, running clinical trials, deploying robots, constructing energy infrastructure, or collecting longitudinal evidence can create barriers. Cottier and colleagues documented the rising cost of training frontier AI systems, illustrating how the technical frontier can become accessible only to a small number of organizations (Cottier et al. 2024). Stanford's 2026 AI Index similarly showed the concentration of

notable model production in industry and the strategic importance of data-center and chip infrastructure (Stanford HAI 2026).

The ability to outspend competitors is not automatically defensible. Waste can scale as easily as learning. Capital supports a moat when each dollar improves a connected asset: performance, cost, evidence, supply, distribution, or regulatory standing. A company that burns more to reproduce a commoditizing capability is financing delay.

The strongest frontier companies possess the right to outlearn. They can run more experiments, observe more consequential outcomes, attract scarce researchers, and integrate feedback into the system. Their organizational advantage may exceed any individual patent. Research culture, infrastructure, and tacit knowledge become an asset stock.

This is also why small startups can remain competitive in selected domains. They may not outspend the largest firms in general, but they can concentrate learning on a narrow frontier where the incumbent's broad platform is inefficient. A specialized robotics company can own a difficult environment. A scientific software company can build around a unique instrument. A regulated operator can accumulate evidence that a general model provider cannot access.

The investor should distinguish capital requirement from capital strategy. A company that requires billions must have a financing ecosystem, milestones, syndication plan, and governance capable of surviving technical delay. A company that requires tens of millions should not imitate frontier spending merely to appear ambitious. The objective is not maximum capital. It is sufficient capital to preserve the learning trajectory.

The phrase "research-driven" should be used narrowly enough to retain meaning. It describes a company whose future economics depend on producing new, testable knowledge rather than mainly assembling known components. The research may concern algorithms, materials, biology, hardware, robotics, energy, cryptography, or a specialized scientific model. The commercial product is not separate from the research program; deployments create evidence, constraints, and data that improve the frontier.

The investment conviction is deliberately stronger than the claim that research is merely desirable. For large-return venture capital, the absence of a non-generic core should be treated as a disqualifying condition unless the entry price and exit plan explicitly assume a tactical acquisition. Fundamental or applied research is the preferred core when it creates a frontier that compounds through experiments and cannot be leased from the same supplier by every rival. Where research is not the core, an equally durable institutional scarcity must perform the same economic function.

Research is especially valuable in an AI-dominated economy because common capability compresses the value of ordinary implementation. When everyone can build faster, the remaining source of exceptional return is often the ability to discover something competitors cannot yet ask an API to provide. This is not a claim that every venture winner must train a foundation model. It is a claim that extraordinary terminal value requires an extraordinary scarcity. A frontier of knowledge is one of the few scarcities that can expand rather than erode as software production becomes cheaper.

The ability to outspend is defensible only when it creates the right to outlearn. Frontier training costs, fabrication facilities, laboratories, field deployments, and regulatory programs can require sustained capital. Cottier and colleagues document rapidly rising training costs at the frontier, while Sutton's work explains how firms can escalate endogenous sunk costs as markets expand (Cottier et al. 2024; Sutton 1991, 1998). The investor should ask whether spending produces an asset stock and whether a late entrant can reproduce that stock simply by spending more later. Non-compressible time is often the real moat.

Planning capital provides a useful extension. Caplin describes frontier investment as staged discovery in which costly inquiry creates diagnostic understanding that improves future inquiry (Caplin 2026). The asset is not only a result; it is the organization's ability to choose better questions, interpret failures, and navigate related problems. A research company with planning capital can move into adjacent opportunities without returning to zero. This is a stronger foundation for a platform than a generic product roadmap.

Financing a Proof Architecture

Venture capital is commonly described as a mechanism for accepting uncertainty that conventional finance cannot price. The description is incomplete. Venture capital also determines which uncertainties a company is allowed to resolve and which it must conceal in order to raise the next round. A financing market that rewards visible growth before technical, legal, or operational proof can make weak architectures look rational. A financing market that recognizes staged evidence can turn patient learning into a competitive advantage.

The design problem begins with milestone selection. A milestone is not merely a target. It is a measurement regime, and measurement changes behavior. When the next financing depends on pilot count, founders optimize pilot count. When it depends on contracted annual value, they may discount implementation and pull revenue forward. When it depends on benchmark performance, they optimize the benchmark. Goodhart's and Campbell's laws do not require fraudulent actors. They describe predictable adaptation to consequential measures (Goodhart 1975; Campbell 1976).

The alternative is not to abandon metrics but to finance a proof architecture. A proof architecture links each stage of capital to the uncertainty that must be retired before the next scale commitment. In a research-centered company, early capital may establish reproducibility, mechanism, and a credible frontier rather than revenue. In an enterprise operator, it may establish outcome improvement, implementation reuse, and bounded liability. In a physical system, it may establish reliability under field conditions, manufacturing yield, and regulatory evidence. The milestone should make the next dollar more informative, not merely make the chart steeper.

This principle changes the meaning of speed. Speed is valuable when it shortens the learning cycle. It is destructive when it advances the organization before the evidence on which later scale depends. A company that signs ten

heterogeneous pilots may appear faster than one that deeply instruments three deployments. The second may be learning more. Investors need enough domain knowledge to distinguish commercial motion from epistemic progress.

AI makes the distinction harder because persuasive evidence is cheaper to manufacture. Synthetic demonstrations, generated collateral, benchmark selection, and polished customer narratives can create the surface of maturity. At the same time, AI can improve real verification by generating tests, exploring failure modes, monitoring systems, and compressing expert review. The investor's response should not be generalized skepticism. It should be a higher ratio of verifiable claims to presentation.

Governance can reinforce this ratio. Boards should receive information about exception rates, incident classes, human override, implementation labor, supplier dependencies, data permissions, and technical debt alongside revenue and pipeline. These variables are leading indicators of whether the company is becoming an institution or a services layer hidden behind software metrics. They also allow the board to authorize deliberate investment in validation before a failure converts into category distrust.

Financing terms shape the horizon over which such investment is rational. Extremely high early valuations can force a company to pursue growth paths that preserve the next mark rather than the best long-run architecture. Insufficient ownership can reduce founder incentives after repeated capital-intensive rounds. Short runways can make a technically necessary redesign appear impossible. Conversely, unlimited capital without hard evidence can protect weak assumptions and delay correction. The objective is not patient capital in the abstract. It is the capital whose patience is conditional on learning.

Research and development require a special form of patience because their output distribution differs from feature development. Scientific work may produce a negative result that sharply improves the investment decision. A financing system that treats all negative results as failure encourages selective disclosure and incremental programs. Public grants, corporate partnerships, milestone-based venture rounds, and patient family capital can be combined so that the party best able to bear each uncertainty finances it. Howell and colleagues found that early research grants can have persistent effects on innovation and venture funding, illustrating how staged public support can alter the path of young technology firms (Howell 2017).

Consortia create another financing problem. Shared data, interoperability layers, safety evidence, or sector infrastructure may generate value for many participants without offering one company an easy monopoly. Conventional venture capital can struggle because the asset's social value exceeds the appropriable value. The solution is not necessarily a nonprofit. A governed entity can charge for access, certification, operations, or premium services while preserving rules that prevent capture. Ostrom's institutional analysis and the literature on open-source governance show that credible boundaries, contribution rules, monitoring, sanctions, and conflict resolution are compatible with economic activity (Ostrom 1990; Lerner and Tirole 2002, 2004).

Such arrangements require investors to think constitutionally. Who can change the rules? Which data can be reused? Can a dominant member discriminate against competitors? What happens when a participant exits? How are improvements licensed? Who pays for verification? These questions determine whether the consortium becomes a moat shared by its members, a neutral utility, or an asset gradually appropriated by the strongest participant.

Category trust also belongs in the proof architecture. Investors can require precise claims, independent evaluation where practical, transparent customer references, documented failure handling, and marketing language proportionate to evidence. These practices may appear conservative during a boom. They create an option when the category enters its inevitable period of skepticism. A company with an accountable record can acquire customers, talent, and competitors while others spend capital rebuilding credibility.


The best financing plans are therefore narratives of uncertainty reduction. They explain what the company must learn, what asset the learning creates, why the asset will remain valuable after the surrounding capability becomes common, and how much capital is required to reach that state. Growth remains important. It becomes evidence of a compounding system rather than a substitute for one. Research finance fails when the capital clock is shorter than the proof clock. Financing risk can push investors toward less novel projects because a technically ambitious company may need future rounds before its quality is legible (Nanda and Rhodes-Kropf 2017). Financial constraints can also suppress radical innovation relative to incremental work (Caggese 2019). These findings support a practical conclusion: the capital structure is part of the technology strategy.

Well-designed public or non-dilutive capital can improve the sequence. Howell's quasi-experimental study of Department of Energy grants found that early awards approximately doubled the probability of subsequent venture financing and increased patenting and revenue, with evidence consistent with funding prototyping rather than mere certification (Howell 2017). Such capital is valuable when it retires uncertainty before private investors must price commercial scale. It becomes harmful when it sustains research without a commercialization path or when grants replace customer and technical evidence.

The appropriate model is patient about time and strict about knowledge. Each stage should state the uncertainty being retired, the experiment that can retire it, the asset created by either result, and the decision that follows. A failed experiment can be a successful milestone if it closes a path cheaply. A positive demonstration can be a failed milestone if it produces no transferable knowledge and encourages premature scale.

Outfinity's studio model can improve this architecture by sharing expensive formation capabilities across companies. Technical diligence, experimental design, IP strategy, security, evaluation, procurement access, and research recruiting can be institutional capabilities rather than rebuilt by every founding team. The studio earns an advantage when these shared capabilities improve the quality and speed of learning without forcing different companies into the same product template.

LONG-HORIZON CAPITAL SHOULD FINANCE A LEARNING LOOP



Capital is defensible only when each cycle makes the next cycle faster, cheaper, or more informative.

Figure 5. Long-horizon capital should finance a learning loop, not a sequence of narratives

Table 11. Capital structures matched to proof cycles

TERM OR ASPECT	DESCRIPTION
Milestone venture rounds	Appropriate when a defined technical or commercial uncertainty can be retired before the next scale commitment and investors can reserve for successful learning.
Non-dilutive research capital	Appropriate for prototyping and knowledge creation with spillovers, provided the company retains a credible route to complementary assets and commercial capture.
Family-office capital	Appropriate for long cycles, physical assets, regulated infrastructure, and concentrated positions where conventional fund duration creates premature pressure.
Strategic corporate capital	Useful for data, manufacturing, validation, or distribution, but dangerous when dependence narrows the company's future bargaining set before technical proof.
Consortium capital	Appropriate for shared data, standards, or assurance that no participant can efficiently build alone and whose governance prevents unilateral capture.
Founder or cash-flow finance	Appropriate when the company can reach profitability before a durable moat is proven and when accepting venture-scale expectations would distort the business.

CHAPTER NINE

How Winning AI-Era Startups Will Be Built

This chapter turns the moat analysis into company archetypes. Winning startups fall into two broad families: frontier builders that own a difficult technical trajectory, and consequence owners that control the institutional system through which AI produces a verified result.

Frontier Builders

The first family owns a frontier. Its competitive advantage is generated by research, engineering, physical deployment, or specialized infrastructure that ordinary software teams cannot reproduce with common models. Frontier builders include scientific platforms, domain-specific models where unique data and evaluation matter, semiconductors, robotics, energy systems, advanced cryptography, industrial hardware, and tools that materially improve the economics of frontier development.

A frontier is a trajectory, not a benchmark snapshot. A company does not become defensible because it leads one public evaluation or demonstrates a result that no competitor has announced. The investor must understand why the organization can continue moving the boundary after the current result becomes known. The relevant assets are experimental systems, tacit knowledge, instrumentation, failure histories, specialized teams, data provenance, manufacturing feedback, and the ability to choose the next useful question. A result can be copied; a superior research process is harder to reconstruct.

The investment case is therefore not technical difficulty by itself. Difficulty can be expensive without being valuable, and novelty can remain commercially sterile. The company must connect the frontier to a bottleneck that customers will pay to remove, possess a route to the complementary assets required for commercialization, and show that each experiment improves a reusable capability. The strongest frontier builder creates an expanding option set: one research program can support several products because the underlying method, infrastructure, or knowledge base remains scarce.

This option set distinguishes a platform of discovery from a one-product science project. A research program should make adjacent problems cheaper to investigate even when the first commercial hypothesis fails. Shared instruments, representations, protocols, simulation environments, manufacturing methods, and evaluation systems can preserve value after a particular market route is rejected. The company is then accumulating planning capital rather than merely consuming runway in search of one binary result.

The relationship between research and customers must also be designed carefully. Customers provide constraints that keep the frontier economically

relevant, but a single customer can capture the roadmap through bespoke requests and confidential data restrictions. A frontier builder needs a commercial wedge that exposes important problems without turning the company into a contract laboratory. Deployment should generate knowledge the firm can lawfully generalize, while the research program should remain coherent enough to produce advantages across accounts.

A physical AI company can be especially defensible because the model is only one component of an embodied system. Sensors, manufacturing, calibration, field reliability, maintenance, safety, supply chains, and customer operations create learning curves that software imitation cannot erase. The capital burden is real, but so is the opportunity to build non-compressible operational knowledge. A rival can license a comparable model and still lack the field evidence, installation base, quality system, or service organization required to deliver the same outcome.

A specialized infrastructure company can likewise win without owning a frontier foundation model. It may provide privacy, latency, reliability, evaluation, data lineage, domain performance, or hardware-software co-design under constraints that broad platforms do not optimize. The test is not whether the problem sounds technical. The test is whether specialization creates a superior learning loop, whether customers face a persistent constraint, and whether the market can finance continued advantage after the platform providers notice the opportunity.

Capital absorption is the decisive diligence question. Every major expenditure should purchase an asset that a later entrant cannot fully recreate on demand: a validated process, a yield improvement, a deployed network, a regulatory dossier, a difficult dataset, a safer architecture, or a better diagnostic model of the problem. If spending merely funds more engineers to assemble familiar components, the company is capital intensive without being capital defensible. Large capital requirements become attractive only when they increase the distance a competent follower must travel.

The team is part of the moat because frontier companies fail at disciplinary boundaries. Research without product discipline can optimize elegant but irrelevant objectives. Commercial pressure without research governance can consume the frontier in custom delivery. Manufacturing, regulation, security, and field operations may determine whether an invention can be appropriated. Outfinity should form teams and boards capable of protecting the research trajectory while forcing it through the complementary assets required for market power.

The evidence standard is correspondingly demanding. A frontier builder should be able to show reproducible experiments, credible baselines, explicit uncertainty, freedom-to-operate analysis, a staged route from proof to production, and an account of what each deployment adds to the next. The investment thesis should survive the removal of the current model supplier and the failure of the first product configuration. When it does, the company is not merely using AI to enter a market. It is building a technical institution that can keep creating scarce capability as the surrounding software becomes common.

Consequence Owners

Enterprise AI spending is growing because the technology is useful, not merely fashionable. Stanford's 2026 AI Index reported broad organizational adoption, while deployment of agents in business functions remained much earlier than general use (Stanford HAI 2026). The gap between access and scaled autonomy is the location of the commercial opportunity.

Model access will become a smaller share of total value as capability diffuses. Enterprises can buy APIs, deploy open models, or use AI embedded in existing software. They still need to decide which data the system may access, how outputs are evaluated, which actions are permitted, who approves exceptions, how incidents are investigated, and how economic impact is measured. These are not peripheral services. They define whether the technology can enter a consequential workflow.

The highest-value enterprise positions are likely to control one of four boundaries. A system of record controls authoritative data. A system of action controls execution. A system of assurance controls evidence, safety, audit, and compliance. A system of coordination aligns multiple parties, tools, and incentives. A company may own more than one, but it must be indispensable at least somewhere.

Systems of record possess structural advantage because AI needs context. Customer relationship, enterprise resource planning, clinical records, source code, financial ledgers, identity, and document repositories already contain data and permissions. Incumbents can add intelligence without asking customers to create another data plane. Startups can displace them only by becoming the new authoritative record or by owning a workflow the record does not serve well.

Systems of action turn recommendations into operations. They create value by shortening the distance between insight and consequence. The moat emerges when the product receives permission to act and when its controls become part of the customer's operating model. An agent that drafts an email is convenient. A system that resolves a claim, reorders inventory, or routes a payment is strategically embedded.

Systems of assurance will become more important as model capabilities and providers change. Enterprises need evaluation suites, provenance, access controls, monitoring, incident response, and evidence for regulators and boards. The product may not generate the most visible output, but it determines whether output can be trusted. Assurance can support multiple models and therefore benefit from commoditization upstream.

Systems of coordination are valuable in fragmented ecosystems. Healthcare, construction, logistics, energy, and financial services involve multiple organizations with incompatible data and incentives. AI can reduce information-processing cost, but a coordinating institution must allocate rights and responsibility. The company that becomes the neutral operating layer may capture value beyond software licensing.

Weber and colleagues' study of AI platforms in medical imaging identifies orchestration logics involving resources, data-centric collaboration, distributed

refinement, and application brokering (Weber et al. 2025). The finding supports a broader point: vertical AI platforms succeed by organizing actors and data, not only by offering a model. The platform is an institution for co-production.

The profit pool also includes implementation and change. Enterprise buyers do not purchase technology in isolation; they purchase a reduction in operational risk. Teams that go to the customer, map the process, integrate data, and redesign roles can create more value than the model. Large AI brands may combine proprietary models with such teams and capture a substantial share of the market. Specialized startups can compete by knowing a domain deeply enough to deliver a better outcome.

The conventional software subscription prices access. Consequential AI may require different models because value and cost depend on use, outcome, and responsibility. Usage pricing aligns with model cost but exposes the vendor to commoditization. Seat pricing fits productivity tools but becomes weak when agents replace seats. Outcome pricing captures value but requires control over confounding variables. Managed-service pricing reflects operational work but can obscure scalability.

The strongest business model often evolves. A company begins with project or service revenue because the workflow is not standardized. It moves to platform and usage revenue as repeatable components emerge. It may eventually price outcomes where evidence and control are sufficient. The transition should be visible in declining human intervention, increasing reuse, and stronger contractual authority.

A labor-intensive business is not automatically unattractive. AI can increase the productivity of expert teams and allow a company to profitably address problems that were previously too fragmented. The relevant margin is not the gross margin in the first year. It is the trajectory of contribution margin as learning accumulates. Investors should penalize disguised services, not services themselves.

The company can defend enterprise margin through proprietary models, but model ownership must be economically justified. Training or fine-tuning a model is valuable when domain performance, latency, privacy, cost, or control creates a customer outcome unavailable from general providers. Owning a model merely to claim vertical integration can consume capital without strengthening the business.

Data and knowledge become a moat when they are operational. A taxonomy of exceptions, a history of adjudications, a graph of dependencies, or a corpus of verified outcomes can improve deployment. Knowledge embedded only in consultants does not compound unless the organization captures it. The operating system must convert tacit expertise into reusable assets without eliminating the judgment that makes it valuable.

Brand matters differently in enterprise markets. A trusted vendor reduces the perceived career risk of the buyer. This can favor incumbents and large AI providers. Startups overcome the disadvantage through evidence, domain authority, references, guarantees, or partnership with established institutions. The founder's personal reputation is useful but must become organizational.

The market will likely contain both centralized operators and governed consortia. A centralized provider can move quickly and assume responsibility. A consortium can pool data and establish standards where participants fear vendor capture. Hybrid models are plausible: a commercial operator runs infrastructure under rules set by members, or a neutral foundation governs standards while companies compete on implementation.

The money in enterprise AI will not be distributed evenly. Generic generation will create enormous value but face price pressure. Integration, assurance, and operation will appear less glamorous but support persistent budgets. The company that makes intelligence safe to authorize may capture more value than the company that produces the most impressive demo.

This forecast also explains why incumbents are formidable. They already own records, workflows, channels, and contracts. Startups need a wedge that creates a new source of authority or solves a problem the incumbent structurally avoids. Merely being more innovative is insufficient when the customer must add another vendor.

The independent opportunity remains large in verticals where data are fragmented, workflows are poorly served, and outcomes are valuable enough to support embedded work. The startup must accept that the product is partly institutional. It needs domain experts, implementation, governance, and customer-specific trust. These are not distractions from software. They are the source of the moat.

The most credible escape route is to own a consequential vertical outcome. A vertical AI operator does not merely tailor a model to an industry. It combines domain knowledge, data rights, workflow integration, evaluation, and responsibility around a result that customers value. The company may process claims, manage compliance, optimize industrial operations, support clinical decisions, or coordinate logistics. Its verticality is institutional rather than cosmetic.

A narrow market can support a large company when the outcome is valuable, recurring, and connected to adjacent workflows. The operator begins with one use case but accumulates authority and data that allow expansion. The expansion follows the customer's process rather than a generic feature roadmap. This reduces exposure to platform envelopment because the product's boundaries are defined by consequence.

The main risk is hidden services. Vertical expertise often resides in people, and early deployments require substantial customization. The company must design for compounding from the beginning. Contracts should permit lawful learning. Operations should record exceptions. Product teams should convert repeated interventions into controls and tools. Sales should select customers whose complexity is representative.

The complexity harvester is a related form. It targets problems that are economically important but unattractive to pure software because they contain fragmented information, exceptions, and coordination. Paul Graham's "schlep blindness" describes the tendency of founders to avoid tedious work even when the work protects an opportunity (Graham 2012). AI lowers the cost of entering these neglected domains by making expert labor more productive.

The company finds the needle in the haystack, but the metaphor should be extended. It does not merely locate the rare item. It builds a map of why the haystack was difficult, which signals identify future needles, and which permissions are required to act. The service engagement becomes a sensor for a reusable system.

Examples include reconciling complex accounts, investigating compliance cases, resolving industrial maintenance exceptions, preparing regulated submissions, or integrating records across institutions. The market may appear small when measured as software spend but large when measured as labor, delay, loss, and risk. AI converts part of that hidden cost into addressable revenue.

The escape route fails when the company becomes a general consultancy with an AI narrative. The distinction is the rate at which intervention per outcome declines and the extent to which each engagement strengthens a reusable asset. Investors should track learning velocity rather than celebrate services backlog.

A second escape route is the research-centered company. It owns a technical frontier and uses applications to commercialize it. The application is not the moat; it is the route by which research becomes customer value. The company must prevent services and product requests from fragmenting the research program. It also must avoid the opposite failure of pursuing science without a credible market.

A third route combines AI with physical assets. Robotics, sensors, laboratories, manufacturing, logistics, and energy systems create deployment constraints that software cannot bypass. The company learns from the physical world and may develop data unavailable to pure model providers. Capital and operational risk are higher, but the position is harder to absorb as a feature.

A fourth route is the trust brand. In categories with high uncertainty, a startup can become an authority through transparent evidence, certification, guarantees, and consistent handling of failure. The brand must be attached to a method and operating system. Marketing alone cannot create the performance bond.

Some moats cannot be built by one company because the necessary data, standards, or participation are distributed among actors that do not trust a dominant sponsor. A governed consortium can create a shared asset while allowing members to exploit it commercially under rules. This is decentralization as institutional design rather than as branding.

The consortium must define its boundary. Who can participate, what must they contribute, which uses are permitted, and how can they exit? Ambiguous membership encourages free riding and disputes. Ostrom's design principles emphasize clear boundaries, monitoring, graduated sanctions, conflict resolution, and rule-making rights adapted to the community (Ostrom 1990; Cox, Arnold, and Villamayor-Tomas 2010).

Rights are the central issue. Members need to know whether raw data, derived models, benchmarks, and commercial outputs are shared, licensed, or exclusive. A rule that is too restrictive prevents investment. A rule that is too permissive allows one participant to appropriate the commons. Bounded

exploitation means members can build businesses while preserving the shared resource.

Governance should separate contribution from control. A large contributor may deserve influence but should not be able to rewrite rules unilaterally. Independent monitoring and transparent decision procedures can preserve credible neutrality. Dispute resolution must be faster than ordinary litigation because technical markets change quickly.

Consortia can support standards, data pools, evaluation, identity, and procurement. They are especially valuable in regulated and fragmented sectors. Their weakness is speed. Startups and investors must decide whether the shared asset is worth slower coordination. A consortium should not govern product decisions that competitive firms can make independently.

Strategic openness follows the same principle. A company can open a model, protocol, or interface to increase adoption while retaining control over assurance, hosting, data, or specialized operation. Open source can prevent supplier lock-in and attract an ecosystem. It can also enable larger companies to capture the market. The firm needs a clear complementary asset before commoditizing the core.

Farrell and Saloner showed that standardization can improve compatibility while affecting innovation incentives (Farrell and Saloner 1985). Shapiro and Varian emphasized that information businesses often compete through standards, switching costs, and complements (Shapiro and Varian 1999). The AI startup should treat openness as market design. What should become common infrastructure so that the company's scarce layer becomes more valuable?

These forms can be combined. A vertical operator may use open models, contribute to a consortium, and build a trust brand. A research company may deploy physical systems and harvest complexity through services. The combination must remain coherent. A company that claims every moat usually possesses none deeply.

The investor should look for a transition equation. What temporary advantage gets the company into the market? What hard asset is created through use? What institutional position is gained? How does that position expand the addressable market? The answer should be visible in contracts, product architecture, hiring, and unit economics.

The startup also needs a stopping rule. Not every useful product should become a platform. A disciplined founder may choose to run a profitable niche business, sell early, or avoid venture capital. The problem arises when financing structure forces a fragile feature business to pretend it will become an institution. Correct classification is itself a strategy.

The constructive message is that AI does not eliminate startup opportunity. It changes the path. The startup must enter through speed and differentiate through accumulation. It should borrow intelligence while building ownership over the conditions of action. The final part translates this principle into capital allocation, diligence, portfolio design, and governance for different classes of investor.

These companies do not merely sell intelligence. They own a completed consequence. A vertical outcome operator accepts responsibility for a bounded process and combines models with data, workflow, assurance, and human escalation. A complexity harvester begins with difficult service work and converts repeated cases into reusable decision systems. An assurance company becomes the trusted layer for evaluation, monitoring, audit, and incident response across changing models. A governed network coordinates data or action among parties that cannot rely on a single participant.

The distinction between the two families is analytical, not absolute. A research company may become an outcome operator, and a complexity harvester may develop a research frontier from its accumulated cases. The crucial point is that the company owns either the creation of scarce capability or the institution that makes capability consequential. The exposed middle consists of firms that own neither.

Consumer winners fit the same framework. A pure utility is vulnerable to aggregation, but a company can own identity, community, intellectual property, a transaction network, a longitudinal record, or a trusted method. These assets turn the product from a replaceable tool into a relationship. The opportunity is real, but the burden of proof is high because a general assistant already occupies the most valuable interface: the place where the user begins.

TWO FAMILIES OF VENTURE-SCALE WINNERS




Figure 6. Two families of venture-scale winners in an AI-dominated economy

Table 12. Winning company archetypes and their failure mode

TERM OR ASPECT	DESCRIPTION
Research-centered company	Owns a technical frontier and uses products to commercialize it; fails when research lacks complementary assets or product demands consume the program.
Physical AI system	Combines models with devices, manufacturing, deployment, or infrastructure; fails when capital and operational complexity do not create a superior learning curve.

TERM OR ASPECT	DESCRIPTION
Specialized infrastructure	Optimizes privacy, latency, reliability, evaluation, or domain constraints; fails when a broad platform can justify equivalent investment and bundle it.
Vertical outcome operator	Controls a consequential workflow and accepts bounded responsibility; fails when "vertical" means only industry-specific prompts and branding.
Complexity harvester	Converts exception-rich service work into reusable assets; fails when labor grows linearly and learning remains tacit or contractually trapped.
Trust and assurance institution	Lowers verification cost through evidence, monitoring, guarantees, and accountable failure; fails when reputation is marketing without operational backing.
Governed network or consortium	Creates data, standards, or coordination no member can build alone; fails when rights are ambiguous or a dominant member can rewrite the rules.
Identity or community company	Owns a durable consumer relationship beyond function; fails when acquisition is rented and the user can move history and value to a general assistant.

PART V

THE OUTFINITY INVESTMENT DOCTRINE

A thesis becomes useful only when it changes selection, ownership, governance, and portfolio construction. The final part states what Outfinity should reject, what it should be willing to finance, and how a venture studio can create an advantage in markets that ordinary startup formation underserves.

CHAPTER TEN

Building a Portfolio for Consolidation

This chapter converts the argument into an Outfinity underwriting system. It treats consolidation as the default scenario, research and institutional scarcity as the source of exceptional outcomes, and the venture studio as a mechanism for financing and forming companies whose proof cycles are too demanding for generic acceleration.

The Outfinity Selection Rule

Outfinity should begin every investment with a counterfactual rather than a market-size slide. Assume that model quality converges, inference cost falls, the visible feature is bundled into an adjacent platform, and the customer can reproduce the demonstration. What remains necessary? The answer defines the investable company. If the answer is merely better execution, speed, or branding around a common capability, the opportunity may support a cash-flow business or a tactical exit but should not receive category-leader underwriting.

This counterfactual makes consolidation the base case rather than a remote risk. It reverses the usual burden of proof. A founder does not establish independence by showing that no incumbent has copied the feature today; the founder must show why copying, bundling, or internal reconstruction will not eliminate the company's bargaining position over the investment horizon. The passage of time is not neutral. Every financing round spent on a generic surface increases exposure to platforms whose cost of entry is falling.

The selection rule is not "research or reject." It is "non-generic core or reprice." Fundamental or applied research is the preferred route where technical discovery can create a durable option set. In other markets, institutional scarcity is equally valid: outcome rights, workflow authority, regulated permission, physical assets, trusted assurance, or a governed network. The common requirement is that use strengthens the company after the initial feature becomes common. The startup must be able to replace its models more easily than the customer can replace the startup.

The first analytical gate is commoditization exposure. Outfinity should identify every capability that can be rented, copied, bundled, or rebuilt and assign it no durable value merely because it is currently impressive. This includes prompt methods, generic retrieval, ordinary agent orchestration, standard connectors, common interfaces, and access to broadly available models. These components can be commercially useful, but their contribution to terminal value should be discounted unless they create a controlled downstream asset.

The second gate is accumulation. The investor should trace how one additional experiment, customer, deployment, or month of operation changes

the competitive position. Revenue that creates no reusable knowledge is cash flow, not compounding. Data that cannot be linked to outcomes or reused across customers are operational exhaust, not a moat. Services that solve each case anew are labor, not a learning system. By contrast, a validated method, exception taxonomy, outcome right, installed base, evaluation corpus, permission, or trusted role can make the next cycle faster and the company harder to displace.

The third gate is capital fit. The company must reach its durable position before imitation and consolidation remove its negotiating leverage. This requires a financing sequence matched to the proof clock, enough ownership and reserves to support the likely outcome, and milestones that retire technical or institutional uncertainty rather than merely increase public visibility. Long-duration capital is justified when time creates knowledge or assets. It is not justified when patience only extends the life of an undifferentiated product.

Outfinity should therefore reject businesses whose primary moat is prompt quality, early access to one model, a polished interface, paid acquisition, an uncontracted claim to proprietary data, or a roadmap that assumes platforms will remain still. It should also reject capital-intensive companies whose spending does not create a time-dependent learning advantage. The studio's purpose is not to finance the most fashionable layer of the stack. It is to form companies that can survive the stack's reorganization.

The rule also permits tactical opportunities, but it requires honest pricing. A narrow SaaS product with strong validation can be a rational studio formation when capital needs are low, strategic buyers are identifiable, ownership is concentrated, and the expected exit is modeled as a feature acquisition. The same company becomes a poor venture investment when it raises at a category-leader valuation, builds a cost structure for global independence, and treats buyer interest as proof of a moat. The economic error is not building the product; it is mismatching product architecture and capital architecture.

A promising company should present a transition equation. The temporary wedge explains how it enters the market. The compounding asset explains what each customer, experiment, or deployment creates. The institutional position explains why the asset remains controlled by the company. The financing plan explains how the company reaches that position before commoditization closes the window. Without all four elements, the pitch describes activity rather than an investment thesis.

The final selection principle is intentionally severe. For venture-scale underwriting, a generic company is presumed non-durable until evidence establishes the contrary. Outfinity is not paid to discover that a useful feature has customers; AI will make such discoveries increasingly common. It is paid to identify or create the rare organization whose value increases when the feature itself becomes cheap.

Underwriting, Ownership, and Return Geometry

Conventional startup diligence begins with a customer problem, a product, a market, a team, and evidence of demand. Those elements remain necessary.

They are no longer sufficient when the productive capability behind the product can be rented from several suppliers, reproduced by a small team, or bundled into a platform already authorized to serve the customer. The AI-era investment question is not merely whether a company can grow under present conditions. It is whether the company becomes stronger when the enabling technology becomes cheaper and more widely available.

This reversal is important because falling input cost has two opposing effects. It improves the startup's gross margin and expands the set of tasks that can be automated. It also lowers the cost of imitation, customer insourcing, and platform entry. The Stanford AI Index documented a decline of more than 280 times in the inference price required to reach roughly GPT-3.5-level benchmark performance between November 2022 and October 2024, while its 2026 edition reported a cluster of leading models with increasingly similar top-line performance (Stanford HAI 2025, 2026). Cheaper intelligence enlarges the market, but it does not assign the enlarged surplus to an application vendor. In many settings it shifts bargaining power toward whoever controls the workflow, customer relationship, data rights, or distribution bundle.

A useful diligence process therefore begins with a consolidation test. The investor asks what remains valuable if model quality converges, inference becomes nearly free, the customer's incumbent software vendor introduces an acceptable substitute, and a technically competent customer can assemble an internal version. This is not an extreme downside case. Each assumption describes an incentive held by a powerful actor. Model providers seek broader usage. incumbents seek higher retention and wallet share. customers seek fewer vendors and lower dependency. imitators seek profitable niches. A company that survives only while these actors remain slow is making a timing trade, not necessarily building a moat.

The test should be applied to a specific unit of value rather than to the company narrative. Consider a system that drafts regulatory documents. Its visible output may be reproducible, but the company may still control a validated corpus, rights to feedback, a review workflow, an audit trail, a network of accountable experts, and contractual responsibility for defined classes of error. The draft is the feature; the governed production system is the company. By contrast, a product that sends documents to a general model and returns polished text may be useful and profitable without possessing an independent strategic position. The distinction cannot be inferred from the interface.

Revenue quality must be examined with the same care. Early AI revenue often combines software, founder attention, manual exception handling, customization, and customer willingness to experiment. Reported recurring revenue can conceal a sequence of bespoke projects whose renewals depend on continued labor. That is not automatically a defect. Services can be the instrument through which a company learns the real workflow. The diligence question is whether each engagement creates an asset available to the next engagement. If the company accumulates reusable process models, evaluation sets, integration components, outcome data, permissions, and reference credibility, service work may finance a compounding system. If every contract begins from zero, growth increases organizational entropy.

Cohort analysis should therefore be extended beyond retention. The investor should compare implementation hours, exception rates, gross margin after fully loaded expert labor, time to verified outcome, deployment reuse, and the share of customer-specific work that becomes a general asset. A company can show excellent logo retention because founders absorb complexity at uneconomic cost. It can also show modest initial margins because it is deliberately converting hard cases into a repeatable operating advantage. The numbers require a theory of learning.

The same principle applies to data. The phrase proprietary data is too weak for investment purposes. Data can be proprietary yet irrelevant, legally unusable for training, non-exclusive, poorly labelled, detached from outcomes, or easy for a customer to recreate. Valuable data has an economic chain of custody. The company has permission to collect it, a reason to receive it continuously, a mechanism for linking it to results, and the ability to use the resulting signal across customers without violating contract or trust. The moat lies as much in the rights and feedback process as in the stored records.

Intellectual property requires similarly disciplined interpretation. A growing patent landscape indicates strategic activity, not automatic defensibility. WIPO reported that published generative-AI patent families rose from roughly 14,000 in 2023 to more than 37,800 in 2025, with large enterprises prominent among leading applicants (WIPO 2026). A crowded patent field can increase both exclusion opportunities and freedom-to-operate risk. The relevant inquiry is whether the company's claims protect an economically important bottleneck, whether infringement can be detected, whether enforcement is financeable, and whether alternative implementations are easy. Under current United States guidance, AI remains a tool and only natural persons can be named as inventors, which also makes documentation of human conception and contribution operationally important (U.S. Patent and Trademark Office 2025).

Team quality should be judged against the company's next constraint rather than against generic founder archetypes. A fast product team may be ideal for discovering a wedge and poorly configured to build a regulated operating system. A distinguished research team may generate real technical progress and still lack the commercial architecture needed to capture it. The investable team is not the one with the maximum credentials. It is the team whose capabilities match the sequence of bottlenecks and whose governance can absorb the specialists required as the company moves from artifact to institution.

The consolidation test ends with a buyer map, but not only in the usual exit sense. The investor identifies every actor able to destroy, appropriate, or accelerate the company's position. A platform may become a supplier, competitor, distributor, or acquirer. An incumbent may bundle the feature while partnering for high-liability cases. A customer may insource the common path but outsource exceptions. A regulator may convert a capability into a controlled function. The company is defensible when these changes increase the value of its controlled assets or at least leave it with a credible role. It is exposed when every plausible development transfers value to someone else.

The result of this analysis is not a binary verdict. A thin product can be an excellent angel investment, a disciplined cash-generating company, or a strategically timed acquisition candidate. The error is financing it at a valuation and burn rate that require category ownership. Underwriting should connect the company's actual durability to the return path and capital structure rather than forcing every useful product into the same venture template.

Return geometry should be explicit at formation. A studio can create value through a portfolio of disciplined strategic exits, but it must not finance each company as though it were a future independent platform. Ownership, burn, option pool, follow-on reserves, and valuation should fit the likely outcome. Conversely, a company with a credible research or institutional moat should receive enough capital and governance support to pursue independence rather than being forced into an early sale because the initial vehicle was too small.

The distinction between a good business and a venture asset is not moral. It is arithmetic. A profitable niche product may be an excellent founder outcome and a poor use of a large fund's time. A research-intensive platform may be an extraordinary venture opportunity and unsuitable for an angel who cannot follow on. Outfinity's comparative advantage should lie in matching company architecture to capital architecture before financing incentives distort the company.

Venture outcomes are skewed not only across startups but across investors. Recent evidence finds that a small minority of venture professionals account for a very large share of investment profits (Jackson and Strebulaev 2026). This concentration suggests that access alone is insufficient. Persistent return requires a differentiated selection and formation capability. Outfinity's thesis is valuable only if it produces companies that other capital providers systematically underform, misprice, or cannot patiently construct.

The Venture Studio as an Institutional Advantage

Investors influence markets not only through individual boards but through portfolios. A portfolio aggregates dependencies, incentives, and reputational effects. In an AI cycle, companies presented as different verticals can share the same foundation models, cloud providers, vector databases, acquisition channels, benchmark assumptions, and enterprise buyers. Their risks are correlated by architecture. Diversification based on application labels alone is therefore illusory.

A stronger portfolio map begins with control points. One company may control a scientific method, another a regulated workflow, another a physical network, another an assurance layer, and another a consumer community. Their revenues can still be correlated with macroeconomic conditions, but their strategic survival does not depend on the same platform decision. The map should also show where the portfolio is a complement to dominant platforms and where it competes with them. Complementors can grow rapidly while remaining exposed to rule changes; competitors can possess more upside but require capital and distribution adequate to survive retaliation.

Platform dependency is not inherently negative. Cloud infrastructure, external models, app stores, and enterprise ecosystems allow startups to avoid rebuilding the economy. The risk appears when dependency is asymmetric and non-portable. A supplier can change price, functionality, access, or terms while the startup cannot switch without losing quality, data, or customers. Investors should treat portability as an option with a measurable cost. The relevant question is not whether the company is model agnostic in presentation, but whether it has actually tested alternative suppliers, preserved evaluation data, isolated integration layers, and negotiated rights consistent with switching.

Portfolio construction also needs a theory of consolidation timing. Early in a capability wave, numerous narrow products can grow because incumbents are slow and customers are experimenting. Later, platforms bundle common functions and buyers standardize procurement. The exposed middle contracts. An investor can still finance application companies, but the strategy should specify whether value will be realized before consolidation, through consolidation, or after the company crosses into a defended position. Ambiguity about timing is often disguised as optionality.

Acquisition is one legitimate path through consolidation. The problem is not the exit itself but the price and strategic meaning. A studio or seed fund may deliberately create companies that discover workflows and become attractive to incumbents. Such a strategy can be coherent when formation cost is controlled, ownership is high, and sale values are realistic. It becomes incoherent when companies are financed and valued as independent category leaders while the expected outcome is a feature acquisition.

Investors can sometimes improve exit geometry by coordinating assets before the buyer does. Complementary portfolio companies can share evaluation infrastructure, distribution, regulatory work, or data standards without being forced into an artificial merger. A holding company can combine several narrow products into a broader operating platform when customers prefer one accountable vendor. A consortium can preserve independent innovation while pooling a common layer. These structures must be designed around real economies of scope rather than around financial engineering. The test is whether coordination reduces customer verification and integration cost.

Market repair becomes relevant when category degradation threatens all participants. Investors with several positions in a field can support common benchmarks, disclosure norms, incident taxonomies, certification, interoperability, or shared legal infrastructure. The private return may arise through lower diligence cost, faster procurement, and stronger category legitimacy rather than through direct ownership of the standard. This resembles the role that industry institutions have played in other complex markets. It also reduces the incentive for each startup to invent a proprietary vocabulary that increases sales friction for everyone.

There is a tension between standardization and differentiation. Premature standards can freeze weak architectures, favor incumbents, and narrow experimentation. Absent standards can make every customer repeat basic diligence and integration. The useful boundary is usually to standardize interfaces, evidence, and safety obligations while preserving competition in

implementation and outcome. A consortium should govern what participants need to trust together, not determine every technical choice.

The same logic applies to strategic openness. Open-source components can commoditize a supplier layer, accelerate adoption, attract contributors, and create a de facto standard. They can also destroy a startup's only differentiator. Lerner and Tirole's work on open source emphasizes heterogeneous motivations and complementary business models rather than a simple opposition between open and proprietary software (Lerner and Tirole 2002). A company should open what enlarges the market or weakens a bottleneck controlled by others, while retaining the assets through which it learns, assures, distributes, or assumes responsibility.

Capital providers differ in their capacity to perform market repair. Large venture firms can convene customers and later-stage capital. Specialist funds can define evidence standards because they understand the science. Venture studios can build shared infrastructure and avoid repeating undifferentiated work. Family offices can support longer-lived institutions and consortia, particularly when they own operating assets. Corporate investors can provide deployment and standards access, though governance must prevent them from turning a common layer into exclusive control.

The investor's own incentives remain the final constraint. Fundraising cycles, markups, media status, and competitive access can generate the same Red Queen dynamics described earlier. A firm may finance an inflated round because rivals are doing so, encourage a portfolio company to announce an immature category because the next fund is being raised, or avoid a necessary down round because it damages reported performance. The behavior can be individually rational and collectively corrosive. Capital cannot govern Moloch dynamics while exempting itself from them.

A credible investment institution therefore needs internal mechanisms for epistemic discipline. Decision memoranda should preserve falsifiable assumptions. Post-investment reviews should compare the original moat thesis with what actually accumulated. Failed investments should be analyzed for selection error rather than converted into generic lessons about execution. Valuation changes should be separated from technical progress. Partner incentives should reward realized judgment over durable horizons, not only deployment volume or interim marks.

Institution building is not philanthropy. It is a strategy for markets in which trust, interoperability, and evidence are complements to private value. An investor who helps create these conditions may improve the quality of selection across a portfolio. The work is especially important when AI makes entry cheap enough that the category cannot rely on scarcity to screen participants.

The studio model can create four advantages that a conventional fund often lacks. It can invest before the company exists by forming the research question and recruiting the team around it. It can reuse technical and institutional infrastructure across ventures without forcing them into identical products. It can preserve a longer proof horizon because formation cost and governance are controlled. It can coordinate later capital, customers, and strategic partners around evidence rather than around a public fundraising race.

These advantages disappear if the studio becomes a factory for similar wrappers. Shared access to models, design, sales, and cloud infrastructure can increase speed while creating a highly correlated portfolio. The shared layer should instead be the difficult layer: research operations, evaluation, IP and rights architecture, security, regulated market access, enterprise implementation, and disciplined post-investment learning.

Outfinity should also preserve the option to create consortia or shared institutions when the moat cannot belong to one startup. A data pool, evaluation standard, identity layer, or assurance network may generate more value as governed infrastructure than as a proprietary point product. The investment design can still capture returns through access, certification, premium operations, or equity in businesses built on the common asset. The key is constitutional clarity about contribution, control, exploitation, and exit.

The social chapter of the thesis belongs here in economic form. A studio that supports honest benchmarks, precise claims, incident learning, and credible assurance improves the category in which its portfolio competes. This lowers verification cost and may expand demand for high-quality suppliers. Outfinity should support these institutions where they protect private value and align with its broader preference for technology that improves rather than corrupts the environments on which it depends.

The final portfolio belief is a barbell with a selective middle. One side contains frontier builders capable of absorbing large capital productively. The other contains consequence owners that begin with focused markets but accumulate authority and knowledge. The excluded middle is not all application software. It is application software that remains generic after use. Outfinity's task is to identify the conversion early enough to finance it and to decline the company when no conversion is plausible.

Table 13. The Outfinity underwriting doctrine

TERM OR ASPECT	DESCRIPTION
Consolidation first	Model the company after likely bundling, platform expansion, vendor rationalization, and customer insourcing rather than assuming present product boundaries remain stable.
Non-generic core	Require a research frontier or institutional scarcity that cannot be purchased as a common input and that becomes stronger through operation.
Transition equation	Identify the temporary wedge, the asset created through use, the institutional position that retains it, and the financing path that reaches it in time.
Representative economics	Normalize founder intervention, service labor, exceptions, data preparation, security, assurance, and steady-state support before treating pilot revenue as scalable software.
Rights before data rhetoric	Verify collection, provenance, outcome linkage, reuse, exclusivity where relevant, retention, and revocation rather than accepting the phrase proprietary data.

TERM OR ASPECT	DESCRIPTION
Capital-product fit	Match fund size, ownership, reserves, and duration to the likely outcome; do not finance a feature exit with category-leader cost structure.
Learning governance	Boards should monitor uncertainty retired, asset accumulation, model and platform dependencies, intervention per outcome, and evidence that can falsify the thesis.
Acquisition discipline	Separate strategic-system value from talent, time-to-market, customer access, and defensive build-versus-buy value when modeling exits.

Table 14. What Outfinity should finance and what it should reprice

TERM OR ASPECT	DESCRIPTION
Finance for venture scale	Research-centered companies, physical AI, specialized infrastructure, vertical outcome operators, complexity harvesters, assurance institutions, and governed networks with compounding control.
Finance selectively	Application wedges that have explicit rights, workflow, data, distribution, or research conversion milestones and can reach them before platform convergence.
Price as tactical	Useful products with real cash flow but short moat half-life, supplier concentration, rented distribution, and likely strategic acquisition rather than independent category control.
Reject at venture terms	Thin wrappers, undifferentiated micro-SaaS, prompt advantages, ungoverned data claims, capital intensity without learning, and products whose independence depends on platform inaction.

CONCLUSION

Build Institutions, Not Accelerated Features

The conclusion states the conviction beneath the analysis: AI will create immense value, but exceptional venture returns will accrue to the organizations that own scarce knowledge or the institutional conditions of consequence.

Artificial intelligence does not make intelligence free, and it does not make companies easy. It makes many cognitive operations cheap enough to stop distinguishing the organizations that perform them. Code generation, drafting, classification, search, translation, design, and analysis can create enormous user surplus while reducing the bargaining power of a company whose only asset is access to those operations. The economic value of the capability and the terminal value of the supplier can move in opposite directions.

This is why the normal startup is in a weaker position than it appears. The artifact is easier to build, the market is easier to enter, and the first revenue can arrive with a small team. The same conditions make imitation faster, product boundaries less stable, and customer insourcing more credible. The company can be locally successful while its category is disappearing into a larger platform. Its technology can become ubiquitous while its exit is priced as a feature.

Outfinity's investment belief is that this outcome should be assumed, not feared abstractly. Consolidation is the default scenario against which the company must be designed. The assumption is not universal monopoly and not the disappearance of entrepreneurial opportunity. Open models, regulation, specialization, new interfaces, and customer heterogeneity will preserve contestability in important markets. The burden nevertheless belongs on the startup to show why it remains necessary after the capability becomes common.

The companies capable of doing so will own a non-generic core. Some will produce knowledge through research that cannot be compressed by late capital. Some will control physical or regulated assets. Some will acquire lawful rights to connect decisions with outcomes. Some will become authoritative inside workflows, resolve exceptions, and accept bounded responsibility. Some will govern networks or build brands that reduce the cost of verification. Their common property is not technological purity. It is compounding scarcity.

This is also the case for long-term capital. Patience is not a virtue when it merely postpones recognition of failure. It is productive when it allows a company to complete a learning cycle that competitors cannot cheaply reproduce. The right investor finances research, deployment, and institutional formation as one proof architecture. It values negative evidence, designs milestones around uncertainty, and refuses to force premature scale in order to preserve a financing narrative.

The result is a narrower and stronger venture strategy. Outfinity should not attempt to own every useful AI product. It should own a meaningful share of the few institutions that become more powerful as useful AI products multiply. It should prefer a difficult company with an expanding frontier to an easy company with a shrinking boundary. It should finance the accumulation of knowledge, rights, trust, and authority rather than the acceleration of a feature that a platform will eventually distribute more cheaply.

The belief can be stated without qualification at the level of action. When models, software, and interfaces become commodities, invest in what cannot be commoditized on the same clock. Build the research organization, the outcome system, the governed network, the physical platform, or the accountable brand. Build an institution, not an accelerated feature.

References

- A.P. Moller - Maersk and IBM. 2022. 'A.P. Moller - Maersk and IBM to Discontinue TradeLens, a Blockchain-Enabled Global Trade Platform.' Corporate announcement, 29 November.
- Aghion, Philippe, Nick Bloom, Richard Blundell, Rachel Griffith, and Peter Howitt. 2005. 'Competition and Innovation: An Inverted-U Relationship.' *Quarterly Journal of Economics* 120 (2): 701-728.
- Akerlof, George A. 1970. 'The Market for Lemons: Quality Uncertainty and the Market Mechanism.' *Quarterly Journal of Economics* 84 (3): 488-500.
- Alexander, Scott. 2014. 'Meditations on Moloch.' *Slate Star Codex*, 30 July.
- Arthur, W. Brian. 1989. 'Competing Technologies, Increasing Returns, and Lock-In by Historical Events.' *Economic Journal* 99 (394): 116-131.
- Autor, David, David Dorn, Lawrence F. Katz, Christina Patterson, and John Van Reenen. 2020. 'The Fall of the Labor Share and the Rise of Superstar Firms.' *Quarterly Journal of Economics* 135 (2): 645-709.
- Azevedo, Eduardo M., Florian Scheuer, Kent Smetters, and Min Yang. 2025. 'Dilution vs. Risk Taking: Capital Gains Taxes and Entrepreneurship.' NBER Working Paper 34512. DOI: 10.3386/w34512.
- Bank for International Settlements. 2018. 'Cryptocurrencies: Looking Beyond the Hype.' *Annual Economic Report 2018*, chapter V. Basel: BIS.
- Barney, Jay. 1991. 'Firm Resources and Sustained Competitive Advantage.' *Journal of Management* 17 (1): 99-120.
- Baumol, William J. 1990. 'Entrepreneurship: Productive, Unproductive, and Destructive.' *Journal of Political Economy* 98 (5): 893-921.
- Bena, Jan, Bo Bian, and Mariassunta Giannetti. 2025. 'Prompted to Start: How Generative AI Is Transforming Entrepreneurship.' HKU Jockey Club Enterprise Sustainability Global Research Institute Paper 2025/182. DOI: 10.2139/ssrn.5749564.
- Beraja, Martin, and Francisco J. Buera. 2026. 'The Life-cycle of Concentrated Industries.' NBER Working Paper 34770. DOI: 10.3386/w34770.
- Bloom, Nicholas, John Van Reenen, and Heidi Williams. 2019. 'A Toolkit of Policies to Promote Innovation.' *Journal of Economic Perspectives* 33 (3): 163-184.
- Bresnahan, Timothy F., and Manuel Trajtenberg. 1995. 'General Purpose Technologies: Engines of Growth?' *Journal of Econometrics* 65 (1): 83-108.
- Brynjolfsson, Erik, Danielle Li, and Lindsey R. Raymond. 2025. 'Generative AI at Work.' *Quarterly Journal of Economics* 140 (2): 889-942. DOI: 10.1093/qje/qjae044.
- Brynjolfsson, Erik, Daniel Rock, and Chad Syverson. 2021. 'The Productivity J-Curve: How Intangibles Complement General Purpose Technologies.' *American Economic Journal: Macroeconomics* 13 (1): 333-372.
- Burkhardt, Marcus, and Bernhard Rieder. 2024. 'Foundation Models Are Platform Models: Prompting and the Political Economy of AI.' *Big Data & Society* 11 (2). DOI: 10.1177/20539517241247839.
- Caggese, Andrea. 2019. 'Financing Constraints, Radical versus Incremental Innovation, and Aggregate Productivity.' *American Economic Journal: Macroeconomics* 11 (2): 275-309.
- Campbell, Donald T. 1976. 'Assessing the Impact of Planned Social Change.' *Occasional Paper 8*. Kalamazoo, MI: Public Affairs Center, Western Michigan University.
- Caplin, Andrew. 2026. 'Planning Capital and Discovery-Based Learning-by-Doing: Investment as Staged Discovery in the Hyperscale Era.' NBER Working Paper 35349. DOI: 10.3386/w35349.

- Caves, Richard E. 2000. *Creative Industries: Contracts Between Art and Commerce*. Cambridge, MA: Harvard University Press.
- Cottier, Ben, Robi Rahman, Loredana Fattorini, Nestor Maslej, Tamay Besiroglu, and David Owen. 2024. 'The Rising Costs of Training Frontier AI Models.' arXiv:2405.21015.
- Cox, Michael, Gwen Arnold, and Sergio Villamayor-Tomas. 2010. 'A Review of Design Principles for Community-Based Natural Resource Management.' *Ecology and Society* 15 (4): 38.
- Cunningham, Colleen, Florian Ederer, and Song Ma. 2021. 'Killer Acquisitions.' *Journal of Political Economy* 129 (3): 649-702.
- Dell'Acqua, Fabrizio, Edward McFowland III, Ethan R. Mollick, Hila Lifshitz-Assaf, Katherine C. Kellogg, Saran Rajendran, Lisa Kraye, Francois Candelon, and Karim R. Lakhani. 2026. 'Navigating the Jagged Technological Frontier.' *Organization Science* 37 (2): 403-423.
- Derfus, Pamela J., Patrick G. Maggitti, Curtis M. Grimm, and Ken G. Smith. 2008. 'The Red Queen Effect: Competitive Actions and Firm Performance.' *Academy of Management Journal* 51 (1): 61-80.
- Dierickx, Ingemar, and Karel Cool. 1989. 'Asset Stock Accumulation and Sustainability of Competitive Advantage.' *Management Science* 35 (12): 1504-1511.
- Dillon, Eleanor W., Sonia Jaffe, Nicole Immorlica, and Christopher Stanton. 2025. 'Shifting Work Patterns with Generative AI.' NBER Working Paper 33795. DOI: 10.3386/w33795.
- Eisenmann, Thomas, Geoffrey Parker, and Marshall Van Alstyne. 2011. 'Platform Envelopment.' *Strategic Management Journal* 32 (12): 1270-1285.
- Erdem, Tulin, and Joffre Swait. 1998. 'Brand Equity as a Signaling Phenomenon.' *Journal of Consumer Psychology* 7 (2): 131-157.
- Farrell, Joseph, and Garth Saloner. 1985. 'Standardization, Compatibility, and Innovation.' *RAND Journal of Economics* 16 (1): 70-83.
- Finkel, Eli J., Paul W. Eastwick, Benjamin R. Karney, Harry T. Reis, and Susan Sprecher. 2012. 'Online Dating: A Critical Analysis from the Perspective of Psychological Science.' *Psychological Science in the Public Interest* 13 (1): 3-66.
- Foote, Brian, and Joseph Yoder. 2000. 'Big Ball of Mud.' In *Pattern Languages of Program Design 4*. Boston: Addison-Wesley.
- Fowler, Martin. 2015. 'Microservice Trade-Offs.' MartinFowler.com, 1 July.
- Frank, Robert H., and Philip J. Cook. 1995. *The Winner-Take-All Society*. New York: Free Press.
- Gans, Joshua S. 2026a. 'Competitive Exposure and Entrepreneurial Experimentation.' NBER Working Paper 35172. DOI: 10.3386/w35172.
- Gans, Joshua S. 2026b. 'Belief Dispersion and Entrepreneurial Entry.' NBER Working Paper 35091. DOI: 10.3386/w35091.
- Gans, Joshua S., and Avi Goldfarb. 2026. 'O-Ring Automation.' NBER Working Paper 34639. DOI: 10.3386/w34639.
- Gans, Joshua S., and Scott Stern. 2003. 'The Product Market and the Market for Ideas: Commercialization Strategies for Technology Entrepreneurs.' *Research Policy* 32 (2): 333-350.
- Giachetti, Claudio, Joseph Lampel, and Stefano Li Pira. 2017. 'Red Queen Competitive Imitation in the U.K. Mobile Phone Industry.' *Academy of Management Journal* 60 (5): 1882-1914.
- Godfrey, Michael W., and Daniel M. German. 2014. 'On the Evolution of Lehman's Laws.' *Journal of Software: Evolution and Process* 26 (7): 613-619. DOI: 10.1002/smr.1636.
- Goodhart, Charles A. E. 1975. 'Problems of Monetary Management: The U.K. Experience.' In *Papers in Monetary Economics*, volume 1. Sydney: Reserve Bank of Australia.

- Graham, Paul. 2012. 'Schlep Blindness.' Essay, January.
- Gupta, Abhinav, Franklin Qian, Elena Simintzi, and Yifan Sun. 2026. 'Generative AI and Entrepreneurship.' Paper presented at the NBER Productivity, Innovation, and Entrepreneurship Conference, Spring 2026.
- Hall, Robert E., and Susan E. Woodward. 2010. 'The Burden of the Nondiversifiable Risk of Entrepreneurship.' *American Economic Review* 100 (3): 1163-1194.
- Hardin, Garrett. 1968. 'The Tragedy of the Commons.' *Science* 162 (3859): 1243-1248.
- Heller, Michael A., and Rebecca S. Eisenberg. 1998. 'Can Patents Deter Innovation? The Anticommons in Biomedical Research.' *Science* 280 (5364): 698-701.
- Hohpe, Gregor, and Bobby Woolf. 2003. *Enterprise Integration Patterns: Designing, Building, and Deploying Messaging Solutions*. Boston: Addison-Wesley.
- Howell, Sabrina T. 2017. 'Financing Innovation: Evidence from R&D Grants.' *American Economic Review* 107 (4): 1136-1164.
- Howell, Sabrina T., Marina Niessner, and David Yermack. 2020. 'Initial Coin Offerings: Financing Growth with Cryptocurrency Token Sales.' *Review of Financial Studies* 33 (9): 3925-3974.
- Humlum, Anders, and Emilie Vestergaard. 2025. 'Still Waters, Rapid Currents: Early Labor Market Transformation under Generative AI.' NBER Working Paper 33777, revised March 2026.
- Jackson, Blake, and Ilya A. Strebulaev. 2026. 'Human Capital in Venture Capital: Evidence from 100,000 Venture Capitalists.' NBER Working Paper 35501. DOI: 10.3386/w35501.
- Kamepalli, Sai Krishna, Raghuram G. Rajan, and Luigi Zingales. 2022. 'Kill Zone.' NBER Working Paper 27146, revised June 2022. DOI: 10.3386/w27146.
- Kaplan, Steven N., and Antoinette Schoar. 2005. 'Private Equity Performance: Returns, Persistence, and Capital Flows.' *Journal of Finance* 60 (4): 1791-1823.
- Katz, Michael L., and Carl Shapiro. 1985. 'Network Externalities, Competition, and Compatibility.' *American Economic Review* 75 (3): 424-440.
- Kerr, Steven. 1975. 'On the Folly of Rewarding A, While Hoping for B.' *Academy of Management Journal* 18 (4): 769-783.
- Kim, Dohyun, Joshua Gluck, Malcolm Hall, and Yuvraj Agarwal. 2019. 'Real World Longitudinal iOS App Usage Study at Scale.' CoRR abs/1912.12526.
- Klein, Benjamin, and Keith B. Leffler. 1981. 'The Role of Market Forces in Assuring Contractual Performance.' *Journal of Political Economy* 89 (4): 615-641.
- Lehman, Meir M., and Laszlo A. Belady, eds. 1985. *Program Evolution: Processes of Software Change*. London: Academic Press.
- Lerner, Josh, and Ramana Nanda. 2020. 'Venture Capital's Role in Financing Innovation: What We Know and How Much We Still Need to Learn.' *Journal of Economic Perspectives* 34 (3): 237-261.
- Lerner, Josh, and Jean Tirole. 2002. 'Some Simple Economics of Open Source.' *Journal of Industrial Economics* 50 (2): 197-234.
- Lerner, Josh, and Jean Tirole. 2004. 'Efficient Patent Pools.' *American Economic Review* 94 (3): 691-711.
- Myers, Kyle R., and Lauren Lanahan. 2022. 'Estimating Spillovers from Publicly Funded R&D: Evidence from the U.S. Department of Energy.' *American Economic Review* 112 (7): 2393-2423.
- Nanda, Ramana, and Matthew Rhodes-Kropf. 2017. 'Financing Risk and Innovation.' *Management Science* 63 (4): 901-918.
- Nelson, Richard R., and Sidney G. Winter. 1982. *An Evolutionary Theory of Economic Change*. Cambridge, MA: Harvard University Press.
- Noy, Shakked, and Whitney Zhang. 2023. 'Experimental Evidence on the Productivity Effects of Generative Artificial Intelligence.' *Science* 381 (6654): 187-192.

- Ostrom, Elinor. 1990. *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge: Cambridge University Press.
- Ostrom, Elinor. 2010. 'Beyond Markets and States: Polycentric Governance of Complex Economic Systems.' *American Economic Review* 100 (3): 641-672.
- Ouyang, Shumiao, Jiaheng Yu, and Ravi Jagannathan. 2021. 'Return to Venture Capital in the Aggregate.' NBER Working Paper 27690, revised October 2021. DOI: 10.3386/w27690.
- Pew Research Center. 2023. 'The Experiences of U.S. Online Daters.' 2 February.
- Pronk, Tila M., and Jaap J. A. Denissen. 2020. 'A Rejection Mind-Set: Choice Overload in Online Dating.' *Social Psychological and Personality Science* 11 (3): 388-396.
- Schmidt, Marc-Thomas, Beth Hutchison, Peter Lambros, and Rob Phippen. 2005. 'The Enterprise Service Bus: Making Service-Oriented Architecture Real.' *IBM Systems Journal* 44 (4): 781-797.
- Schrepel, Thibault, and Alex Pentland. 2025. 'Competition between AI Foundation Models: Dynamics and Policy Recommendations.' *Industrial and Corporate Change* 34 (5): 1085-1103. DOI: 10.1093/icc/dtaf004.
- Sculley, D., Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips, Dietmar Ebner, Vinay Chaudhary, Michael Young, Jean-Francois Crespo, and Dan Dennison. 2015. 'Hidden Technical Debt in Machine Learning Systems.' *Advances in Neural Information Processing Systems* 28.
- Shapiro, Carl. 2001. 'Navigating the Patent Thicket: Cross Licenses, Patent Pools, and Standard Setting.' *Innovation Policy and the Economy* 1: 119-150.
- Shapiro, Carl, and Hal R. Varian. 1999. *Information Rules: A Strategic Guide to the Network Economy*. Boston: Harvard Business School Press.
- Shumailov, Iliia, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. 2024. 'AI Models Collapse When Trained on Recursively Generated Data.' *Nature* 631: 755-759.
- Sigg, Stephan, Emil Lagerspetz, Ella Peltonen, Petteri Nurmi, and Sasu Tarkoma. 2019. 'Exploiting Usage to Predict Instantaneous App Popularity: Trend Filters and Retention Rates.' *ACM Transactions on the Web* 13 (2), Article 13. DOI: 10.1145/3199677.
- Simon, Herbert A. 1971. 'Designing Organizations for an Information-Rich World.' In *Computers, Communications, and the Public Interest*, edited by Martin Greenberger. Baltimore: Johns Hopkins University Press.
- Stanford Institute for Human-Centered Artificial Intelligence. 2025. *AI Index Report 2025*. Stanford University.
- Stanford Institute for Human-Centered Artificial Intelligence. 2026. *AI Index Report 2026*. Stanford University.
- Sutton, John. 1991. *Sunk Costs and Market Structure*. Cambridge, MA: MIT Press.
- Sutton, John. 1998. *Technology and Market Structure: Theory and History*. Cambridge, MA: MIT Press.
- Swartz, Lana. 2022. 'Theorizing the 2017 Blockchain ICO Bubble as a Network Scam.' *New Media & Society* 24 (7): 1695-1713.
- Teece, David J. 1986. 'Profiting from Technological Innovation: Implications for Integration, Collaboration, Licensing and Public Policy.' *Research Policy* 15 (6): 285-305.
- U.S. Patent and Trademark Office. 2025. 'Revised Inventorship Guidance for AI-Assisted Inventions.' *Federal Register guidance*, 26 November.
- Weber, Michael, Andreas Hein, Jorg Weking, and Helmut Krcmar. 2025. 'Orchestration Logics for Artificial Intelligence Platforms: From Raw Data to Industry-Specific Applications.' *Information Systems Journal* 35 (3): 1015-1043. DOI: 10.1111/isj.12567.

- Wen, Wen, and Feng Zhu. 2019. 'Threat of Platform-Owner Entry and Complementor Responses: Evidence from the Mobile App Market.' *Strategic Management Journal* 40 (9): 1336-1367.
- World Intellectual Property Organization. 2026. 'GenAI Innovation Soaring, With Patent Activity Nearly Tripling in Two Years.' *Technology SPARK*, 14 July.
- Yudkowsky, Eliezer. 2017. *Inadequate Equilibria: Where and How Civilizations Get Stuck*. Berkeley, CA: Machine Intelligence Research Institute.
- Zuckerman, Ezra W. 1999. 'The Categorical Imperative: Securities Analysts and the Illegitimacy Discount.' *American Journal of Sociology* 104 (5): 1398-1438.