

INVESTING IN AN AI-DOMINATED ECONOMY

VOLUME II

THE NEXT COMPUTING PLATFORM

**An Outfinity Guide and Investment Thesis for
Long-Horizon Investors**

*How research-driven ventures can occupy the strategic gaps created by
AI, platform renewal, and the institutional limits of software*

DRAFT MANUSCRIPT

Version 0.1 | July 2026

OUTFINITY VENTURE STUDIO | RESEARCH EDITION | JULY 2026

INVESTOR NOTE

This volume is not a catalogue of companies in the Outfinity pipeline. It is a map of the market that the pipeline is intended to explore. The names of ventures may change, individual prototypes may fail, and some research directions may never become independent companies. The investment thesis should survive those changes because it concerns the direction of the computing platform rather than the fate of any one product.

Volume I argued that cheap access to models and software weakens the independent economics of generic AI applications. When a capability can be rented, copied, bundled, or rebuilt inside a larger suite, traction may prove demand without proving strategic independence. Volume II asks the constructive question that follows: if the obvious application layer is crowded and structurally exposed, where can companies with venture-scale durability still be built?

The answer developed here is that artificial intelligence is renewing the computing platform. Every such renewal solves an old scarcity, absorbs a generation of point solutions, and creates a new set of missing institutions. AI makes generation, prediction, and increasingly action abundant. It therefore raises the value of what abundance does not provide by itself: governed execution, durable context, verification, rights over data, legitimacy, scientific reproducibility, and social environments capable of sustaining meaning rather than merely maximizing exposure.

Outfinity's conviction is that these gaps are not peripheral ethics features. They are candidate markets. They become more valuable as AI becomes more capable, because the cost of acting on an unverified result, delegating to an unidentified agent, losing control over sensitive context, or organizing social life through synthetic abundance also rises. The studio is therefore oriented toward research that can discover and occupy these layers before they are recognized as conventional categories.

Contents

Investor Note	2
Preface. A Thesis, Not a Catalogue	4
1. Computing History Is a Sequence of Platform Renewals	5
2. AI Is the Next Platform, Not the Next Feature	8
3. The Obvious AI Market Will Consolidate	11
4. The Industries Created by AI's Own Deficits	14
5. Research Is Venture Infrastructure	19
6. The Outfinity Investment Thesis	22
Conclusion. Invest in the Missing Institutions	25
References	26

Preface. A Thesis, Not a Catalogue

The easiest way to misunderstand the AI market is to look only at the products currently visible. The interface of the moment attracts imitation because it is legible. A chat box, a coding assistant, a document generator, or an autonomous workflow can be demonstrated quickly, described in familiar metrics, and financed with a recognizable SaaS narrative. The resulting activity is real. It is also concentrated on the surface of a much larger transition.

Technological history rarely rewards investors who confuse the first popular interface with the final control point. The mainframe was not merely a faster calculator. The personal computer was not merely a smaller mainframe. The web was not merely a cheaper publishing channel, and cloud computing was not merely rented servers. Each wave reorganized what counted as a product, which complements became necessary, where bargaining power accumulated, and which previously independent applications became ordinary functions of a platform.

Artificial intelligence is now producing a comparable renewal. The models are the most visible component, but they are not the whole platform. A model does not know which organizational facts are authoritative, which files it is allowed to read, which action can bind a company, how a scientific claim depends on an experiment, or whether a digital identity deserves trust. It can produce an answer without owning the consequences. Once AI moves from generating text to participating in work, the missing layer becomes institutional rather than merely computational.

This book advances a strong portfolio prior. The obvious AI application market will remain prolific, but much of it will consolidate. Many useful products will survive as features, internal tools, acqui-hires, or modest independent businesses. The largest durable outcomes will tend to appear where a venture owns a non-generic core and becomes part of the infrastructure through which intelligence is authorized, verified, remembered, coordinated, or socially legitimated. Research is central to this thesis because the relevant categories are not fully specified. Their architecture, standards, failure modes, and commercial boundaries must be discovered before they can be scaled.

This is not a prediction that every research lab will become a valuable company, nor an argument that capital intensity is a moat. Research can become theater, and ambitious infrastructure can become an excuse for avoiding the customer. The discipline is stricter: every research cycle must reduce a consequential uncertainty and leave behind an asset that compounds. That asset may be a representation, protocol, benchmark, corpus, runtime, validated method, permission, or body of tacit knowledge. A venture earns the right to raise more capital when those assets create a better position in the next experiment and a stronger claim on the future market.

CHAPTER 1

Computing History Is a Sequence of Platform Renewals

The first investment question is not which AI application is growing fastest. It is which layer of the computing system is being rewritten, and which permanent functions will migrate into that layer.

The history of information technology can be read as a sequence of platform renewals. A platform renewal occurs when a new technical substrate becomes sufficiently general that many industries reorganize around it. Bresnahan and Trajtenberg described such general-purpose technologies as pervasive, continuously improvable, and productive of complementary innovation. Their most important observation for investors is that the value of the core technology depends on a large body of downstream invention, while the incentives of the core provider and its users do not automatically align. A decentralized market can therefore produce both frantic activity and "too little, too late" investment in the complements that make the technology genuinely productive (Bresnahan and Trajtenberg, 1995).

The mainframe era concentrated computation, expertise, and authority. Computing was scarce, expensive, and administered by institutions that decided who could use it. The personal computer did not merely reduce the size of the machine. It redistributed control and created the operating system as the durable intermediary between changing hardware and applications. A large class of utilities that had once been purchased separately became functions of the operating system, while new categories appeared above it because the standardized environment lowered the cost of invention.

The web renewed the platform again. It made documents, services, and later transactions addressable across networks. Browsers absorbed protocols and rendering complexity; search engines made an expanding public corpus navigable; identity and payment systems became essential complements. Cloud computing then standardized the acquisition of servers, storage, and databases, while mobile platforms joined software to location, sensors, persistent identity, and app-store distribution. Each wave preserved much of the previous infrastructure, but changed its economic status. What had been a differentiated product often became a commodity input or a bundled feature. What had been an implementation detail became a strategic control point.

This pattern matters because markets are organized around permanent human and organizational problems, not around permanent product forms. People continually need to compute, remember, communicate, coordinate, prove identity, evaluate claims, and belong to groups. The form through which those needs are met changes with the platform. A successful word processor in the desktop era and a successful collaborative document system in the cloud era both address the production of text, but they occupy different architectures,

distribution channels, data models, and sources of lock-in. The investor who identifies only the visible product category is late to the deeper change.

The AI transition begins by making a new class of operations cheap: interpretation, generation, prediction, and increasingly the selection and execution of actions. These operations were previously embedded in human labor or in narrow, manually specified software. Once they become reusable components, the surrounding system must be redesigned. The new platform must decide which model is appropriate, what context it may use, how an intention becomes a process, when an action requires approval, how an output is tested, and which record survives after the model is replaced.

The platform renewal is therefore not the replacement of all software by one universal model. It is the insertion of programmable cognition into the software stack. The No Free Lunch results in optimization do not prove that artificial general intelligence is impossible or remote; they show only that no optimizer dominates over all arbitrary problem distributions. Real environments are structured, and general models can exploit that structure. The investment implication is narrower and more useful: even extremely capable models will operate within domain distributions, legal boundaries, data rights, physical constraints, and institutional objectives. Specialization and composition remain economically important because the world is not an unconstrained benchmark.

Industry history also warns that the first phase of a platform renewal is deceptive. Entry expands rapidly because technical possibility outruns stable category definitions. Firms experiment with interfaces, pricing, and distribution. Later, dominant designs emerge, platforms bundle adjacent functions, and the number of independent firms contracts. Klepper's studies of industry life cycles and shakeouts show that an early expansion in firm formation can coexist with later concentration; the two are stages of the same evolutionary process rather than contradictory observations (Klepper, 1997; Klepper and Simons, 1997).

For investors, the practical distinction is between the permanent problem and its temporary implementation. A startup may appear to own a market because it is the first convenient interface to a new capability. Yet if the capability is destined to become native to the next platform, the company may be occupying a transitional product form. The more durable opportunity lies in the layer that will still be necessary after the interface, model, and deployment pattern have changed.

Table 1. Permanent functions and their likely AI-era platform form

TERM OR ASPECT	DESCRIPTION
Access to computation	Model routing across local, edge, private, and public infrastructure under explicit cost and data policies.
Production of artifacts	Specification-driven processes in which models, tools, deterministic code, and human judgment create versioned outputs.
Organizational memory	Persistent project state, provenance, decisions, exceptions, and evaluation history that survive model replacement.

TERM OR ASPECT	DESCRIPTION
Coordination	Governed workflows among people, agents, tools, and organizations, with declared roles and bounded authority.
Identity and trust	Verifiable identity, authorization, performance evidence, audit, and remedies rather than reputation inferred from presentation alone.
Belonging and status	Bounded social environments with rules, roles, memory, progression, and consequences rather than a single undifferentiated feed.




Figure 1. Each computing wave commoditizes part of the previous stack and moves scarcity toward a new control point.

CHAPTER 2

AI Is the Next Platform, Not the Next Feature

The platform thesis becomes investable only when the migration of scarcity is explicit: models become more capable and cheaper, while the system surrounding their actions becomes more valuable.

The strongest evidence for a platform transition is not that models can perform impressive demonstrations. It is that capability is becoming simultaneously better, cheaper, and more widely available. The Stanford AI Index reported that the cost of querying a model with roughly GPT-3.5-level benchmark performance fell from about \$20 per million tokens in November 2022 to \$0.07 by October 2024, a reduction of more than 280 times. By March 2026, several frontier providers were clustered closely in human preference rankings, shifting competition toward cost, reliability, deployment, and domain performance rather than simple access to an obviously superior model (Stanford HAI, 2025; 2026a).

Abundance does not imply completion. The same 2026 report found that agents still failed roughly one in three attempts on structured computer-use benchmarks, while scientific agents remained far below experts on end-to-end research and paper-scale replication. Responsible-AI measurement and governance lagged capability, and documented incidents continued to rise. The market is therefore defined by a productive tension: intelligence is becoming an inexpensive input before it has become a dependable institution (Stanford HAI, 2026a; 2026b; 2026c).

This tension explains why a chatbot is too small a metaphor. A model can generate an answer; a platform must govern a consequence. Consequences require persistent context, identity, authorization, tools, state, evaluation, and accountability. The system must know not only what the model inferred, but what source was authoritative, which version of a policy applied, whether the action was reversible, who approved it, and what should happen when the result is disputed. These requirements are not optional additions for regulated industries. As agents acquire access to files, terminals, communication channels, and financial or operational systems, they become the conditions under which autonomy can be purchased at all.

The next control point is therefore likely to be an operating layer for intelligence. Such a layer does not need to replace Linux, Windows, or cloud infrastructure. It operates above them, coordinating models, agents, tools, data, permissions, and project memory. Its purpose is analogous to an operating system in the economic sense: it stabilizes interfaces among components that evolve at different speeds, hides irrelevant implementation complexity, and exposes the decisions that must remain under human or institutional control.

The AssistOS Book from outfinity.ch site illustrates this category with clarity. It describes a local-first and cloud-optional environment in which projects, artifacts, agents, permissions, execution traces, and model routing belong to a stable workspace rather than to a single chatbot or API. The visible product is a workspace; the proposed long-term architecture progresses toward a runtime, package ecosystem, and control layer. Components such as Ploinky, MRP-VM, Wormhole, and OpenDSU explore sandboxed execution, executable processes, inter-organizational communication, provenance, and rights. The manuscript explicitly separates operational components, prototypes, and research hypotheses, which is the correct way to present an architectural thesis at this stage (Axiologic Research, 2026).

AssistOS is useful here as an example, not as proof that one implementation will win. The strategic point is that the same missing functions appear across enterprise adoption. Organizations are experimenting with many models and agents, but their durable assets are the projects, policies, data relationships, and validated procedures that must survive provider changes. A neutral operating layer can benefit from model commoditization because every new model increases the value of portability, routing, and common governance. The more capable the components become, the more expensive it is to administer them separately.

This is also why local and sovereign executions matter without requiring an anti-cloud ideology. Training the largest models will remain centralized and capital intensive. Many workloads will continue to use elastic cloud infrastructure. Yet inference, retrieval, private memory, and narrow models are increasingly feasible on controlled hardware. A hybrid system can keep sensitive context near the user, escalate only when necessary, and preserve a functional core when a provider, connection, price, or jurisdiction changes. Sovereignty, in this sense, is continuity and optionality rather than isolation.

The economics of the transition can be expressed as a migration from capability to consequence. At the capability layer, a company rents intelligence available to many competitors. At the process layer, it combines that capability with domain representations, tools, and evaluation. At the consequence layer, it controls an outcome that matters to a customer. At the institutional layer, it holds the rights, trust, standards, and memory that allow the outcome to be repeated and defended. Value does not automatically increase at each step, but durability usually does because imitation requires reproducing more than an interface.

A platform thesis is credible only if it begins with a wedge. The mistake of infrastructure founders is to build a universal substrate before a single workflow makes the substrate necessary. The correct sequence runs in the opposite direction. A difficult and consequential process exposes the missing abstractions; repeated implementations turn those abstractions into a runtime; shared components create a platform; and only demonstrated interoperability can produce an ecosystem. The operating layer is not declared. It is earned by becoming the best place to complete work that weaker architectures cannot govern.




Figure 2. As model capability becomes rentable, durable value migrates toward governed processes, consequential outcomes, and institutional control.

CHAPTER 3

The Obvious AI Market Will Consolidate

Cheap experimentation creates more companies, not necessarily more durable categories. The investor must distinguish a wave of entry from the structure that remains after platforms absorb the obvious functions.

Paul Graham's observation that the space of convenient startup ideas has been "stripped pretty clean" is particularly relevant in an AI cycle. His argument was not that the opportunity had ended. It was that founders systematically avoid ideas involving tedious integration, regulation, physical operations, difficult partnerships, or ambiguous categories. The result is schlep blindness: too many people compete in the tractable space, while valuable problems remain in the messy territory they prefer not to see (Graham, 2012a; 2012b).

AI intensifies this asymmetry. GitHub reported more than 180 million developers, 630 million repositories, 4.3 million AI-related projects, and more than 1.1 million public repositories importing an LLM SDK in 2025. These figures do not count viable companies, but they measure the extraordinary capacity of the market to explore an obvious software idea. A working paper on entrepreneurship found that industries with greater exposure to generative AI experienced about 20 percent more startup formation after diffusion of the technology. Lower experimentation costs create genuine discovery, but they also increase the number of teams able to reach the same visible opportunity (GitHub, 2025; Bena, Bian, and Giannetti, 2025).

The resulting paradox is central to this volume: AI can generate more startups in the short run and fewer independent categories in the long run. Industrial evolution frequently begins with broad entry and ends with a shakeout. Dominant designs reduce uncertainty, scale advantages accumulate, and platforms use distribution, data, installed relationships, and bundling to move into adjacent markets. Eisenmann, Parker, and Van Alstyne describe platform envelopment as entry through a multi-product bundle that leverages shared users and network effects. A narrow AI product is unusually exposed because a model provider, operating system, productivity suite, CRM, design platform, or cloud vendor can add its function to an existing relationship and price it at an incremental amount that an independent company cannot match (Eisenmann, Parker, and Van Alstyne, 2011).

Teece's analysis of technological innovation explains why technical originality alone does not prevent this outcome. When imitation is easy or intellectual property is weak, returns migrate to owners of complementary assets such as manufacturing, distribution, service, regulatory position, or customer access. In AI applications, the code and even much of the visible capability may be reproducible. The complementary assets are therefore decisive: proprietary rights over outcome-linked data, integration into a

consequential workflow, trusted distribution, licenses, physical infrastructure, domain assurance, or a research trajectory that competitors cannot rent from the same API (Teece, 1986).

Capital structure reinforces the pattern. The NVCA 2026 Yearbook reports that AI captured 65.4 percent of US venture deal value in 2025, while the five largest companies raised almost \$60 billion collectively. This is not a uniform boom. It is a barbell in which enormous amounts of capital flow to model and infrastructure leaders while a very large population of smaller application companies competes for attention, distribution, and acquisition. An application company can be useful, fast-growing, and even profitable without possessing the economic independence required to return a large fund (NVCA, 2026).

The appropriate underwriting rule is therefore firm: in the absence of a demonstrable non-generic core, a generic AI application should be valued as a feature-in-waiting. It may create a tactical exit because an incumbent wants the team, customers, integration, or time saved. That is different from owning a category. Validation proves that someone wants the function. It does not prove that the startup will remain the cheapest, most trusted, or most convenient place to buy it after the function is bundled into a larger platform.

This pressure is visible in enterprise software, where the buyer rarely wants another isolated intelligence surface. The buyer wants a process completed inside existing systems, with identity, permissions, data lineage, exceptions, and accountability. A thin application that summarizes a document or drafts a response can be reproduced by the document suite. A company that administers a regulated review, carries validated domain logic, and becomes responsible for the evidence trail occupies a different position. It is no longer selling an answer. It is becoming an operational institution.

Consumer AI faces an analogous constraint in the finite budget of interfaces. A person can discover an unlimited number of applications but cannot maintain an unlimited number of subscriptions, memories, identities, privacy settings, and habits. AI utilities resemble mobile applications more than games. In games, novelty is part of the consumption and many distinct worlds can coexist. In utilities, the user often prefers one trusted assistant to absorb adjacent tasks. A specialized consumer company must therefore own something the general assistant cannot cheaply reproduce: a community, identity, transaction network, corpus, longitudinal relationship, professional method, or social environment with accumulated meaning.

This argument does not imply that every narrow product disappears. Geographic markets, regulation, language, physical operations, and specialized communities can sustain independent firms. It implies that the burden of proof has shifted. The investor should not ask only whether the product is better today. The investor should ask why the company remains a company after model quality converges, inference prices fall, the incumbent bundles the feature, and the customer learns to build a basic version internally.

The strategic opportunity lies beyond the most legible implementation. AI writing is obvious; control over a regulated document process is strategic. A generic research assistant is obvious; a validated domain knowledge system that can execute and test claims is strategic. A companion chatbot is obvious; a

persistent relationship environment with memory, identity, safety, and community governance is strategic. The visible interface may be the wedge, but the company becomes defensible only when use transforms the wedge into rights, learning, authority, or infrastructure.

Table 2. The difference between a visible AI opportunity and a strategic one

TERM OR ASPECT	DESCRIPTION
A generic writing or analysis application	A governed workflow that owns evidence, approvals, exceptions, and responsibility for a consequential outcome.
A research chatbot over papers	An executable domain knowledge product linking claims, assumptions, protocols, evidence, and validation.
A generic agent framework	A control plane with identity, policy, observability, portability, and trusted deployment across environments.
A privacy wrapper around a model	Rights-aware data infrastructure that enables useful computation while preserving governance, consent, and provenance.
A provenance badge	A legitimacy institution combining identity, origin, performance evidence, governance, and remedies.
A new social feed or AI companion	A bounded social or relational environment with continuity, roles, rules, rituals, status, and consequences.

CHAPTER 4

The Industries Created by AI's Own Deficits

The most strategic markets may be created not by what AI can do, but by what widespread AI makes newly scarce: trustworthy action, executable knowledge, private context, credible identity, and meaningful social structure.

Every important technology solves scarcity and creates a deficit. The automobile solved distance while creating markets for roads, insurance, traffic law, safety engineering, and urban planning. The internet solved distribution while creating markets for search, cybersecurity, identity, moderation, and digital payments. AI solves the scarcity of cognitive production. It therefore creates deficits wherever society previously relied on the cost of producing an answer, image, plan, or persuasive signal as indirect evidence of competence and intent.

This dynamic defines the most important strategic markets of the next platform. They are institutional deficit markets: markets for the structures that make abundant intelligence safe to use, expensive to deceive, capable of accumulating knowledge, and compatible with durable human cooperation. Their products may contain software, but their value comes from reducing uncertainty and assigning consequence.

The infrastructure of consequence

The first deficit is the governed agency. A language model that proposes an action is an information tool. An agent that can deploy code, transfer data, send a message, modify a scientific protocol, or commit an organization enters the domain of authority. Natural-language instructions cannot serve as the sole security boundary because the same model that interprets the rule is exposed to adversarial or mistaken content. External controls must identify the agent, bind it to a principal, limit its capabilities, record its actions, and distinguish proposal from execution.

NIST's 2026 work on software and AI agent identity and authorization treats this as an emerging enterprise infrastructure problem. The relevant functions include identification, authentication, least privilege, delegation, audit, non-repudiation, and response to prompt injection. These are not merely compliance checkboxes. They are prerequisites for deploying autonomy at scale. A company that owns a validated control plane for consequential agents can sell the right to delegate without losing institutional memory or accountability (NIST, 2026a; 2026b).

The second deficit is executable knowledge. Scientific results remain primarily narrative documents. A paper states claims, describes methods, presents figures, and points to supplementary material, but the logic connecting

assumptions, data, protocol, analysis, and conclusion is often reconstructed manually. The format was optimized for human communication under conditions of scarce publication, not for machine inspection across millions of results. AI can summarize the paper, but summarization is not the same as knowing under which conditions a claim holds or reproducing the process that generated it.

The reproducibility problem makes this distinction commercially relevant. In the Open Science Collaboration's 2015 project, replication effects across 100 psychology studies were roughly half the magnitude of the originals, and 36 percent of replications were statistically significant compared with 97 percent of the original results. The study concerned a specific field and has itself been debated, so it should not be universalized. Its enduring lesson is that narrative publication and peer review do not automatically create a machine-verifiable evidence base (Open Science Collaboration, 2015).

Research infrastructures such as Whole Tale and RO-Crate already point toward a different object. They package data, code, environment, workflow, provenance, and narrative so that a result can be recreated and inspected. The Outfinity Executable Science direction extends this logic from reproducible computation to structured scientific reasoning. Claims become objects linked to assumptions and evidence; protocols become executable plans with expected outcomes; contradictions become queryable relations; and provenance traces how a conclusion depends on data and methods. This is not a replacement for scientific judgment. It is an attempt to give judgment a more inspectable substrate (Brinckman et al., 2019; Soiland-Reyes et al., 2022; Leo et al., 2024).

The market entry should be narrow. A universal ontology of science is too ambitious as a first product and likely to fail on domain variation. A credible venture begins with a high-value decision in a bounded field: assessing target evidence in drug discovery, transferring an experimental protocol, monitoring contradictions in a therapeutic area, or validating a regulated scientific review. The open representation can become a standard, while the defensible company accumulates a validated corpus, mappings, exception cases, buyer relationships, and integration into the decision process.

The third deficit is sovereign and rights-aware computation. AI increases the economic value of context because private data, organizational history, genomic information, and operational exceptions can improve decisions. It simultaneously increases the cost of losing control over that context. The strategic question is no longer whether data is stored in a database. It is who may use them, for which purpose, through which computation, under which jurisdiction, and with what evidence of compliance.

This creates room for local, edge, federated, and privacy-preserving architectures, but the product cannot be a vague promise that data remain private. It must combine an explicit threat model, cryptographic or statistical protections where appropriate, governed workflows, rights and consent, provenance, and operational assurance. OpenDSU and the proposed Omics Firewall are relevant research probes because they treat data as an object with identity, policy, version, and cross-organizational rights rather than as raw material to be copied into one central platform. Hyperdimensional or vector-symbolic representations may contribute to efficient or robust

computation, but they should not be described as privacy guarantees without adversarial evidence.

Europe is a plausible early market for this layer because policy increasingly treats technological sovereignty as continuity and strategic agency rather than simple data localization. The opportunity is not to reproduce US hyperscalers at a smaller scale. It is to build products that make mixed infrastructures governable: local where control matters, cloud where elasticity matters, and interoperable enough that organizations can change providers without reconstructing their institutional memory.

These three markets - governed agency, executable knowledge, and rights-aware computation - share a single economic structure. They make intelligence consequential. They transform a model from an impressive supplier of outputs into a component of a system that can be trusted with work whose failure has a cost. The moat is not a decorative governance layer. It is the accumulated ability to complete the process under real constraints.

The infrastructure of legitimacy and social life

The fourth deficit is legitimacy. For much of economic history, producing a polished argument, image, credential, or technical proposal required enough skill and effort that the presentation carried some information about the producer. Generative AI weakens that relationship. Cowgill, Hernandez-Lagos, and Wright found that access to ChatGPT increased the quality of pitches while compressing differences between senders; in their hiring and startup-investment experiments, screening accuracy fell by 4 to 9 percent on average. The effect was not universal, and AI improved screening in some non-English settings, but the general mechanism is important: when impressive talk becomes cheap, evaluators demand costlier signals of underlying quality (Cowgill, Hernandez-Lagos, and Wright, 2026).

This is the commercial foundation of legitimacy ventures. A legitimacy venture does not sell truth as a magical cryptographic property. It reduces the cost of deciding who or what deserves reliance. Content provenance standards such as C2PA can verify an asset's origin, editing history, and association with a signer, but the specification explicitly states that a valid manifest does not establish that the content is true. Provenance is one component of legitimacy, not its substitute (C2PA, 2026).

A credible legitimacy system combines identity, provenance, claims, performance evidence, governance, and remedies. Identity says who stands behind an action. Provenance says how the artifact arrived. Claims specify what is being asserted. Performance evidence records whether similar assertions or agents have succeeded. Governance establishes who may issue, challenge, or revoke status. Remedies determine what happens when trust is abused. The product becomes valuable when this combination reduces a measurable transaction cost: fraud, due-diligence time, marketplace risk, certification expense, or the cost of admitting participants to a high-trust network.

The decentralized brand is one possible institutional form. It sits between a fully centralized corporate identity and an ungoverned community label. A credible core defines standards and protects the meaning of the brand, while contributors, operators, or local communities can participate under transparent rules. Cryptographic identity and provenance may support the structure, but the value lies in incentive alignment and credible distinction. The brand must be capable of excluding bad behavior, recognizing legitimate variation, and preserving continuity without making every participant dependent on one opaque platform.

The fifth deficit is social architecture. The dominant social network reduced social life to an open feed because the feed was a highly efficient interface for distribution and advertising. It was never a complete model of belonging. It collapses professional identity, friendship, performance, outrage, entertainment, courtship, and political argument into a single public arena. AI-generated content and synthetic participants increase the supply of material entering that arena without increasing the human capacity to interpret it.

Evidence about social algorithms is mixed, which is itself instructive. Large experiments on Meta platforms found limited effects from switching to chronological feeds in some settings, while a 2026 randomized field experiment on X found that exposure to its algorithmic feed increased engagement and shifted several political attitudes and following behavior. Another 2026 study showed that feed-ranking changes could reduce exposure to polarizing content and improve perceptions of social norms without reducing enjoyment. The responsible conclusion is not that one algorithm determines society. It is that architecture changes incentives, exposure, and the formation of relationships, and that alternative designs are empirically testable (Guess et al., 2023; Gauthier et al., 2026; Brady et al., 2026).

The next social platform may therefore not be a better feed. It may be a system of bounded environments with a defined purpose, limited membership, persistent memory, roles, rituals, progression, and consequences. A group for scientific debate can turn rivalry into rule-governed performance. A rehearsal environment can let people practice leadership, conflict, interviews, grief, or parenting with AI characters and human peers. A deep-interest community can convert knowledge into identity and status. A productive social environment can reproduce the texture of guilds, workshops, and shared projects as routine employment becomes less central for some participants.

These environments do not require a technically maximal metaverse. A small room with ten humans, a persistent AI host, a memory of prior sessions, and a credible ritual can be more valuable than a vast but socially empty virtual city. The core product is the institution: who belongs, what behavior counts, how status is earned, what is remembered, and what consequences follow. AI supplies characters, moderation, personalization, and generated scenes, but it does not supply the social contract.

The commercial logic is retention based on meaning rather than exposure. A feed must continuously win the next second of attention. A bounded environment can accumulate identity, relationship, reputation, and history. Its moat is not simply engagement; it is the cost of leaving a social world in which

one's role and memory are recognized. This creates ethical risk because dependency and manipulation can also increase retention. Governance is therefore part of the product and part of the investment case. A system that cannot constrain its own incentives will eventually consume the trust on which its economics depend.

Legitimacy ventures and social technologies belong in the same chapter because both respond to the collapse of cheap signals. One rebuilds credible distinction in markets; the other rebuilds meaningful structure in social life. Investors need not treat these aims as philanthropy. When trust, identity, and belonging become scarce, organizations and users pay for institutions that produce them. The social usefulness of the product can be an economic asset when it lowers risk, improves selection, and supports relationships that survive the novelty of the interface.

Table 3. Institutional deficits created by AI abundance

TERM OR ASPECT	DESCRIPTION
Abundant but unbounded action	Agent identity, authorization, observability, assurance, and governed execution.
Abundant claims but weak evidence structure	Executable scientific knowledge, reproducible protocols, contradiction detection, and domain evidence products.
Valuable context under concentrated infrastructure	Local-first operating layers, privacy-preserving computation, rights-aware data systems, and sovereign deployment.
Cheap persuasive signals	Legitimacy infrastructure combining identity, provenance, performance evidence, governance, and remedies.
Infinite content and synthetic participation	Bounded social environments organized around roles, rituals, memory, progression, and purpose.

CHAPTER 5

Research Is Venture Infrastructure

When a category is not yet stable, research is not a preliminary expense before the company begins. It is the mechanism by which the company discovers what can become a durable asset and a new market.

The conventional startup method is optimized for uncertainty about demand. Build a small product, expose it to users, measure behavior, and iterate. This method remains valuable, but it is insufficient when the central uncertainty is architectural. A team can validate that customers want a research assistant without learning how scientific claims should be represented. It can validate interest in autonomous agents without resolving identity, delegation, and rollback. It can validate demand for privacy without establishing a threat model or a workable rights architecture. Rapid product feedback can select an interface while leaving the strategic problem untouched.

Research-centered venture building begins with a different question: what must be true for this market to exist in a durable form? The answer is decomposed into uncertainties that can be investigated through theory, prototypes, benchmarks, adversarial testing, standards work, and deployment with domain partners. The objective is not to remain in a laboratory. It is to convert ignorance into assets before the market has settled on a weak architecture.

This approach is especially important in a general-purpose technology transition. Bresnahan and Trajtenberg's "too little, too late" problem arises because core technology providers and downstream innovators cannot fully coordinate their investments. Model companies optimize general capability and infrastructure economics. Domain organizations optimize their current operations. The missing complements - executable knowledge, governed agency, portable workflows, privacy, legitimacy - can remain underbuilt because no incumbent initially owns the entire benefit. A venture studio with a shared research base can occupy that coordination gap.

The economic product of research is the right to outlearn. A company has this right when each experiment improves its ability to run the next experiment and when competitors cannot obtain the same learning merely by purchasing the same model or reading the same public paper. The right may come from a validated representation, a difficult benchmark, a corpus with lawful rights, a runtime that captures execution traces, a network of domain collaborators, or tacit knowledge accumulated through repeated failures. Research becomes defensible when it changes the slope of learning, not when it merely increases the volume of technical activity.

This criterion also disciplines intellectual property. Patents can matter when they protect a bottleneck, infringement is observable, enforcement is credible,

and alternatives are limited. They are not a substitute for an economic position. Open-source code can be strategically valuable when adoption, inspection, and standard formation matter more than secrecy. The commercial moat may then reside in certification, validated data, managed deployment, assurance, customer-specific rights, and the operational knowledge required to make an open architecture reliable. The question is not whether the artifact is open or closed. It is which layer must be open to enlarge the market and which scarce complement the company is positioned to own.

Axiologic Research provides a concrete example of portfolio-level compounding. AssistOS explores the workspace and operating layer; Ploinky explores controlled execution; MRP-VM explores explicit, inspectable processes derived from specifications; OpenDSU contributes experience with provenance, identity, and governed data; neuro-symbolic work investigates validation and composition; executable science applies the same discipline to claims and protocols. These components should not be forced into one commercial product. Their strategic value is that each venture can begin with deeper architectural continuity and a lower cost of testing adjacent hypotheses.

The organizational rule should be shared research, separate companies. Research objects, open standards, evaluation harnesses, and infrastructure modules can be reused across the studio. A venture should be formed only when a distinct buyer, problem, team, IP boundary, and financing path are visible. This prevents the studio from becoming a conglomerate of unresolved ideas and prevents each company from rebuilding the same foundational capabilities. It also preserves option value: a research result may support several use cases before the strongest commercial boundary becomes clear.

Capital must follow the uncertainty being retired. Grants, research partnerships, and patient studio capital are appropriate when the output is a method, benchmark, protocol, or technical proof. Validation capital is appropriate when the key question is whether a domain partner will integrate the result into consequential work. Company-formation capital becomes rational when there is a buyer, a wedge, an assignable asset base, and a team capable of owning the market. Large venture rounds should arrive when capital can accelerate learning, infrastructure, rights, and distribution - not when it merely subsidizes a generic product against a platform that can bundle it.

This sequencing is more demanding than ordinary SaaS because it refuses the emotional comfort of early metrics that do not prove the thesis. It is also more aligned with venture returns. A small application can reach revenue quickly and still have a low ceiling. A research-driven company may take longer to establish its architecture, but once it owns a consequential standard, evidence base, physical system, or institutional role, additional capital can widen a position that is difficult to reconstruct.

The model should not romanticize duration. Long-term work is not valuable because it is long, and difficult work is not valuable because it is difficult. Research earns time only by producing increasingly specific evidence. A failed hypothesis can be a valuable output if it removes an entire branch of the design space. A prototype that attracts praise but teaches nothing about reliability, rights, or adoption is not progress. The studio's advantage must be epistemic

discipline: the ability to distinguish activity from learning and learning from a defensible commercial asset.

Table 4. Research outputs as investment assets

TERM OR ASPECT	DESCRIPTION
Theory and representation	Defines the category, exposes assumptions, and turns a vague vision into testable architecture.
Benchmark and evaluation	Creates a credible basis for comparing systems in a market otherwise dominated by demos and claims.
Runtime, protocol, or reference implementation	Makes repeated deployment possible and converts tacit architecture into reusable infrastructure.
Validated corpus and lawful data rights	Creates a learning loop tied to outcomes that competitors cannot reproduce from public model access alone.
Open standard or open-source core	Accelerates adoption, partner contribution, and category formation while shifting monetization to scarce complements.
Tacit operational knowledge	Improves adaptation speed, exception handling, and the probability that later capital compounds rather than merely increases burn.




Figure 3. Research becomes venture infrastructure when reusable objects create focused wedges and, eventually, independent institutions.

CHAPTER 6

The Outfinity Investment Thesis

Outfinity is not searching for applications that AI makes easy to build. It is searching for institutions that AI makes necessary.

The investment thesis can now be stated without reference to a particular venture list. Artificial intelligence is renewing the computing platform. The visible application layer will produce rapid experimentation, genuine utility, and many tactical exits, but it will also be subject to model commoditization, platform envelopment, interface consolidation, and a finite budget of user attention. Venture-scale durability will concentrate in companies that own a non-generic core and become part of the infrastructure through which intelligence produces trusted consequences.

This means Outfinity should prefer opportunities that become more necessary as AI improves. A company whose value disappears when general models become cheaper is positioned against the direction of history. A company whose value increases because cheaper models produce more agents, more scientific claims, more sensitive context, more synthetic media, or more pressure on social coordination is positioned with it. The distinction is not between AI and non-AI companies. It is between ventures that rent the new abundance and ventures that organize the scarcities the abundance creates.

The second test is the non-generic core. Every candidate must answer what cannot be obtained by a competent competitor using the same models, cloud services, and public code. The answer may be a research trajectory, lawful and outcome-linked data, physical infrastructure, a regulated permission, a standard, a trusted brand, a distribution relationship, or authority over a workflow. "Better prompts" and temporary access to a model are not sufficient. Neither is complexity for its own sake. The core must improve bargaining power as the market matures.

The third test is compounding learning. Each deployment should leave behind a reusable asset: a better representation, evaluation set, exception taxonomy, integration, protocol, provenance graph, or relationship. Services can be an excellent wedge when they harvest complexity and turn it into product. They are a weak venture model when every customer requires a fresh project and the only accumulated asset is revenue history. The investor should be able to describe why the tenth implementation is cheaper, safer, or more informative than the first.

The fourth test is institutional position. Winning companies in the AI era will often control rights and responsibilities rather than merely interfaces. They may become the approved environment for an agent, the evidence system for a scientific decision, the trusted intermediary for private data, the certification layer for a marketplace, or the social institution that gives a community continuity. This is where brand, workflow authority, identity, and governance

become moats. They are difficult to acquire late because they depend on accumulated behavior and credible commitments.

The fifth test is capital compatibility. Large capital is justified when it can accelerate a defensible trajectory: deeper research, difficult infrastructure, regulatory access, physical deployment, data rights, or distribution that reinforces the core. It is not justified merely because competitors are raising. A company that must continuously outspend platforms to maintain a feature is fragile. A company that can use capital to extend a lead in learning or to occupy a standard-setting position may be capable of the asymmetric return that venture portfolios require.

The current Outfinity directions are best understood as probes against this framework. AssistOS Enterprise and related components explore governed agency, local-cloud portability, and the operating layer. Scientific Workbench and Executable Science explore structured knowledge and reproducible research. MRP-VM, the Neuro-Symbolic Systems Lab, and SOPTrace GxP explore explicit processes, validation, and assurance. OpenDSU, Ploinky Wormhole, the Cryptography and Privacy Lab, and Omics Firewall explore rights-aware coordination across organizational boundaries. Decentralized brands and structured social environments explore legitimacy, identity, belonging, and governance. These names are contingent; the deficits they address are the thesis.

Some probes will become features of another venture, some will remain open research, and some may support companies outside the present pipeline through partnership with domain founders and investors. This is not a weakness. A research studio should not force every technical result into the first commercial wrapper available. Its task is to maintain a coherent map, discover where a buyer and a defensible asset coincide, and form the company at that boundary.

The required investor relationship differs from conventional thematic exposure. Outfinity is looking for partners willing to finance learning in stages, contribute domain access, and judge progress by the retirement of strategic uncertainty. Family offices, venture funds, corporate partners, and public research programs can play different roles, but they must share a willingness to separate early evidence from premature scaling. The common objective is not patience as a virtue. It is a larger terminal position than the obvious application market can usually support.

There is also a societal belief beneath the financial thesis. A world in which machine intelligence is governed, inspectable, portable, privacy-aware, scientifically executable, and socially legitimate is more resilient than one in which a few opaque interfaces mediate cognition and action. Investors are not required to treat social benefit as the primary objective. They should recognize, however, that the absence of these institutions creates costs that markets eventually price: fraud, liability, failed adoption, reputational collapse, regulatory intervention, employee resistance, and social exhaustion. Building the necessary institutions can therefore align public usefulness with private return.

The thesis is deliberately selective. Outfinity will miss fast-growing products that do not fit it, and some of those products will become excellent companies.

A thesis is valuable precisely because it excludes. The studio's comparative advantage is not speed in every market. It is the combination of technical depth, research continuity, cryptographic and agentic infrastructure, institutional thinking, and willingness to work in categories that appear premature until the platform makes them inevitable.

Table 5. The Outfinity underwriting framework

TERM OR ASPECT	DESCRIPTION
Historical alignment	The opportunity becomes more necessary, not less, as models improve, costs fall, and agents spread.
Non-generic core	The company owns a scarcity that cannot be rented from the same infrastructure available to competitors.
Compounding learning	Every deployment produces reusable knowledge, rights, infrastructure, or evidence that improves the next one.
Institutional position	The venture accumulates authority, trust, standard-setting power, workflow control, or responsibility for outcomes.
Capital compatibility	Additional capital accelerates defensibility and category control rather than subsidizing a feature against a bundling platform.
Societal durability	The business does not consume the trust, autonomy, or community on which its long-run economics depend.

Table 6. Current research probes and the structural gaps they investigate

TERM OR ASPECT	DESCRIPTION
Governed agency and sovereign execution	AssistOS Enterprise, Ploinky, private deployment, agent identity, permissions, observability, and model portability.
Executable knowledge and scientific memory	Scientific Workbench and Executable Science representations for claims, protocols, evidence, and reproducibility.
Formal assurance and inspectable reasoning	MRP-VM, neuro-symbolic validation, benchmark systems, and SOPTrace GxP.
Private cross-organizational coordination	OpenDSU, Ploinky Wormhole, cryptography and privacy research, federated workflows, and Omics Firewall.
Legitimacy and social architecture	Decentralized brands, verifiable participation, bounded communities, persistent relationship systems, and structured social environments.

Conclusion. Invest in the Missing Institutions

History does not repeat products. It repeats migrations of control. The mainframe concentrated computation, the personal computer redistributed it, the web made information addressable, the cloud industrialized infrastructure, and mobile platforms made software continuous with everyday life. Artificial intelligence now makes interpretation and action programmable. The central investment question is who will organize that capability when it becomes ordinary.

The obvious answer is the model provider or the largest existing platform. They will capture enormous value and absorb many application categories. The less obvious answer is that no platform renewal is complete without institutions the core technology does not supply. AI can generate a scientific hypothesis but cannot, by generation alone, make the evidence executable. It can act but cannot assign itself legitimate authority. It can produce identity signals but cannot decide which identities deserve trust. It can fill a feed but cannot create belonging. It can process private context but cannot determine the rights under which that context should move.

These deficits are where long-horizon venture opportunity begins. They are difficult because the product category, architecture, standard, and buyer often emerge together. They reward research not as ornament but as a method of discovery. They require capital that can fund staged learning and then scale decisively when a non-generic core becomes visible. They also create the possibility of companies whose value to society is inseparable from their economic role: institutions that let intelligence be used without surrendering evidence, autonomy, privacy, or social meaning.

Outfinity's position is therefore simple. Do not invest only in what AI makes easy. Invest in what widespread AI makes necessary. Build ventures where the new platform is still missing its operating layer, scientific memory, rights architecture, legitimacy system, and social institutions. The tactical market will remain noisy. The strategic market is the reconstruction of the environment in which intelligence can be trusted to matter.

References

- Axiologic Research. 2026. AssistOS: The Operating Layer of Sovereign Intelligence. Working edition, July 2026.
- Bena, Jan, Bo Bian, and Mariassunta Giannetti. 2025. "Prompted to Start: How Generative AI Is Transforming Entrepreneurship." Working paper. DOI: 10.2139/ssrn.5749564.
- Brady, William J., et al. 2026. "Redesigning Algorithms to Intervene on Social Norm Misperceptions during a National Election." *Nature*. DOI: 10.1038/s41586-026-10536-1.
- Bresnahan, Timothy F., and Shane Greenstein. 1999. "Technological Competition and the Structure of the Computer Industry." *Journal of Industrial Economics* 47(1): 1-40.
- Bresnahan, Timothy F., and Manuel Trajtenberg. 1995. "General Purpose Technologies: Engines of Growth?" *Journal of Econometrics* 65(1): 83-108.
- Brinckman, Adam, et al. 2019. "Computing Environments for Reproducibility: Capturing the Whole Tale." *Future Generation Computer Systems* 94: 854-867.
- C2PA. 2026. C2PA Technical Specification, version 2.4, including Harms Modelling guidance on provenance and truth.
- Cowgill, Bo, Pablo Hernandez-Lagos, and Nataliya Langburd Wright. 2026. "Does AI Cheapen Talk? Theory and Evidence from Global Entrepreneurship and Hiring." *Management Science*. DOI: 10.1287/mnsc.2024.07027.
- Eisenmann, Thomas, Geoffrey Parker, and Marshall Van Alstyne. 2011. "Platform Envelopment." *Strategic Management Journal* 32(12): 1270-1285.
- Gauthier, Germain, Roland Hodler, Philine Widmer, and Ekaterina Zhuravskaya. 2026. "The Political Effects of X's Feed Algorithm." *Nature* 652: 416-423.
- Gawer, Annabelle, and Michael A. Cusumano. 2014. "Industry Platforms and Ecosystem Innovation." *Journal of Product Innovation Management* 31(3): 417-433.
- GitHub. 2025. Octoverse 2025: A New Developer Joins GitHub Every Second as AI Leads TypeScript to Number One.
- Graham, Paul. 2012a. "How to Get Startup Ideas." November 2012.
- Graham, Paul. 2012b. "Schlep Blindness." January 2012.
- Guess, Andrew M., et al. 2023. "How Do Social Media Feed Algorithms Affect Attitudes and Behavior in an Election Campaign?" *Science* 381: 398-404.
- Jacobides, Michael G., Carmelo Cennamo, and Annabelle Gawer. 2018. "Towards a Theory of Ecosystems." *Strategic Management Journal* 39(8): 2255-2276.
- Klepper, Steven. 1997. "Industry Life Cycles." *Industrial and Corporate Change* 6(1): 145-182.
- Klepper, Steven, and Kenneth L. Simons. 1997. "Technological Extinctions of Industrial Firms: An Inquiry into Their Nature and Causes." *Industrial and Corporate Change* 6(2): 379-460.
- Leo, Simone, et al. 2024. "Recording Provenance of Workflow Runs with RO-Crate." *PLOS ONE* 19(9): e0309210.
- National Institute of Standards and Technology. 2026a. AI Agent Standards Initiative.
- National Cybersecurity Center of Excellence. 2026b. Software and AI Agent Identity and Authorization: Draft Concept Paper.
- National Venture Capital Association. 2026. NVCA 2026 Yearbook: The Venture Industry in Transition.
- Open Science Collaboration. 2015. "Estimating the Reproducibility of Psychological Science." *Science* 349(6251): aac4716.

- Outfinity Venture Studio. 2026. Investing in an AI-Dominated Economy: An Outfinity Venture Studio Guide and Investment Thesis on Research-Driven Startups, Durable Moats, Market Consolidation, and Venture-Scale Returns. Volume I.
- Soiland-Reyes, Stian, et al. 2022. "Packaging Research Artefacts with RO-Crate." *Data Science* 5(2): 97-138.
- Stanford Institute for Human-Centered Artificial Intelligence. 2025. The 2025 AI Index Report.
- Stanford Institute for Human-Centered Artificial Intelligence. 2026a. The 2026 AI Index Report: Technical Performance.
- Stanford Institute for Human-Centered Artificial Intelligence. 2026b. The 2026 AI Index Report: Science.
- Stanford Institute for Human-Centered Artificial Intelligence. 2026c. The 2026 AI Index Report: Responsible AI.
- Teece, David J. 1986. "Profiting from Technological Innovation: Implications for Integration, Collaboration, Licensing and Public Policy." *Research Policy* 15(6): 285-305.
- Wolpert, David H., and William G. Macready. 1997. "No Free Lunch Theorems for Optimization." *IEEE Transactions on Evolutionary Computation* 1(1): 67-82.